

WELCOME TO SNOWFLAKE TRAINING



Name: Rajesh Hegde

Contact: rajeshh@useready.com


Prior Experience:

Goal: To make sure you walk out with better understanding about Snowflake

Hobby/Passion: Photography/Long Drive



*** Note: These sessions will be recorded for training & quality improvement purpose



WELCOME TO DAY 1

- ❑ Key Concepts & Architecture
- ❑ Snowflake Editions & Regions
- ❑ Overview of Key Features
- ❑ Overview of the Data Lifecycle
- ❑ Continuous Data Protection

- Snowflake is an analytic data warehouse provided as Software-as-a-Service (SaaS).
- Snowflake provides a data warehouse that is faster, easier to use, and far more flexible than traditional data warehouse offerings.

In this Topic:

Data Warehouse as a Cloud Service

Snowflake Architecture

-Database Storage


-Query Processing

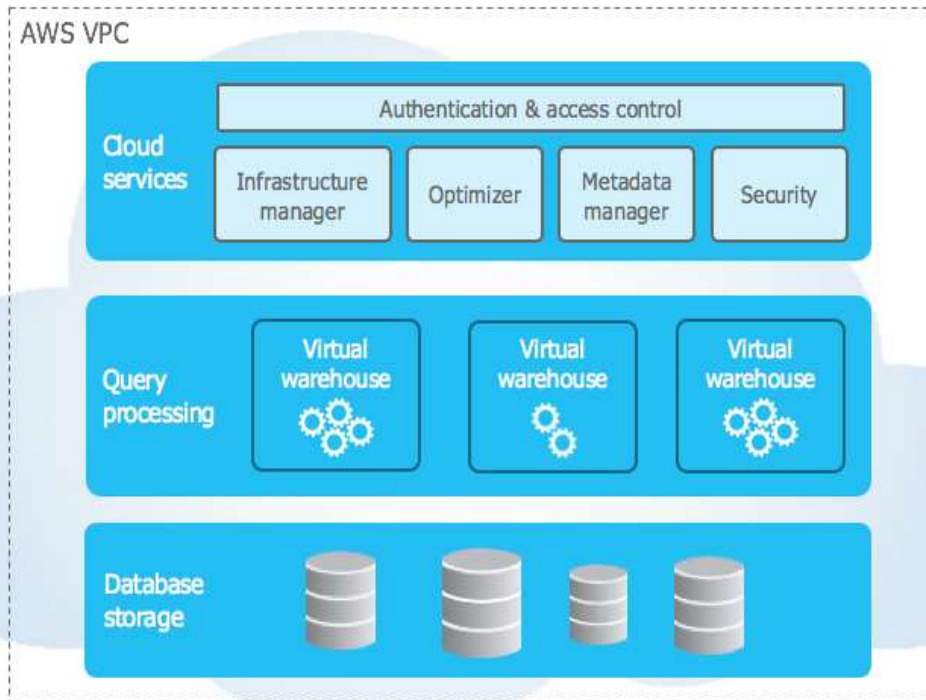
-Cloud Services

Connecting to Snowflake



- ❑ Snowflake runs completely on cloud infrastructure. All components of Snowflake's service (other than an optional command line client), run in a public cloud infrastructure.
- ❑ At this time, Snowflake runs exclusively in the Amazon Web Services (AWS) cloud infrastructure.
- ❑ Snowflake uses virtual compute instances provided by AWS EC2 (Elastic Compute Cloud) for its compute needs and AWS S3 (Simple Storage Service) storage for persistent storage of data.





- ❑ Hybrid of traditional shared-disk database architectures and shared-nothing database architectures
- ❑ Central data repository for persisted data that is accessible from all compute nodes in the data warehouse
- ❑ Processes queries using MPP (massively parallel processing) compute clusters where each node in the cluster stores a portion of the entire data set locally
- ❑ Snowflake's unique architecture consists of three key layers.

- Cloud Services - The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch. The cloud services layer also runs on compute instances provisioned by Snowflake from Amazon EC2.

Access Control

Authentication


Infrastructure Management

Metadata Management


Query Parsing and Optimization



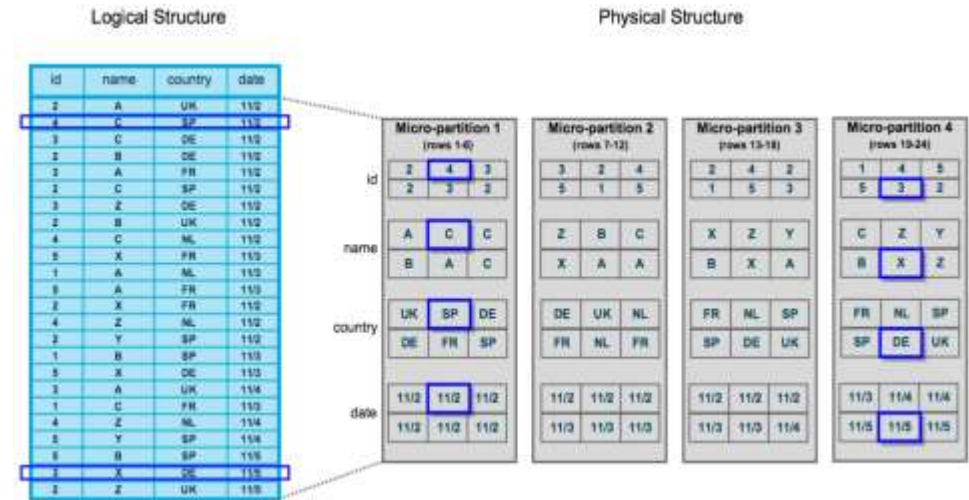
- ❑ Query Processing - Query execution is performed in the processing layer. Snowflake processes queries using “virtual warehouses”. Each virtual warehouse is an MPP compute cluster composed of multiple compute nodes allocated by Snowflake from Amazon EC2.



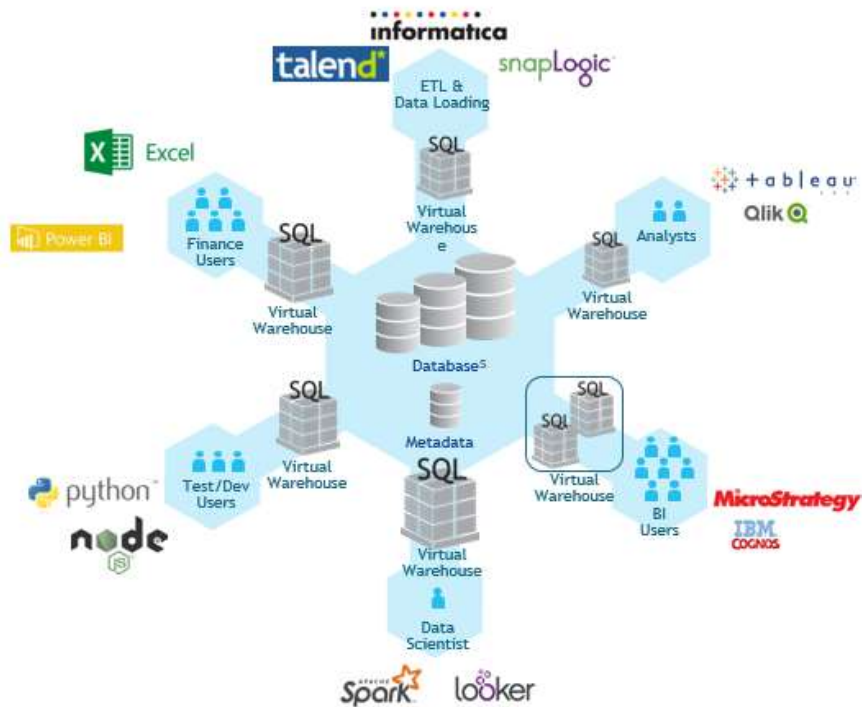
- Database Storage - When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format. Snowflake stores this optimized data using Amazon Web Service's S3 (Simple Storage Service) cloud storage.



- ▣ **Micro-partitions:** All data in Snowflake tables is automatically divided into micro-partitions
 - Contiguous units of storage
 - 50 MB and 500 MB of uncompressed data
 - enables extremely efficient DML and fine-grained pruning for faster queries
 - Cluster Key is used to co-locate data within same micro-partitions
 - User can create custom cluster key



<https://docs.snowflake.net/manuals/user-guide/tables-micro-partitions.html>



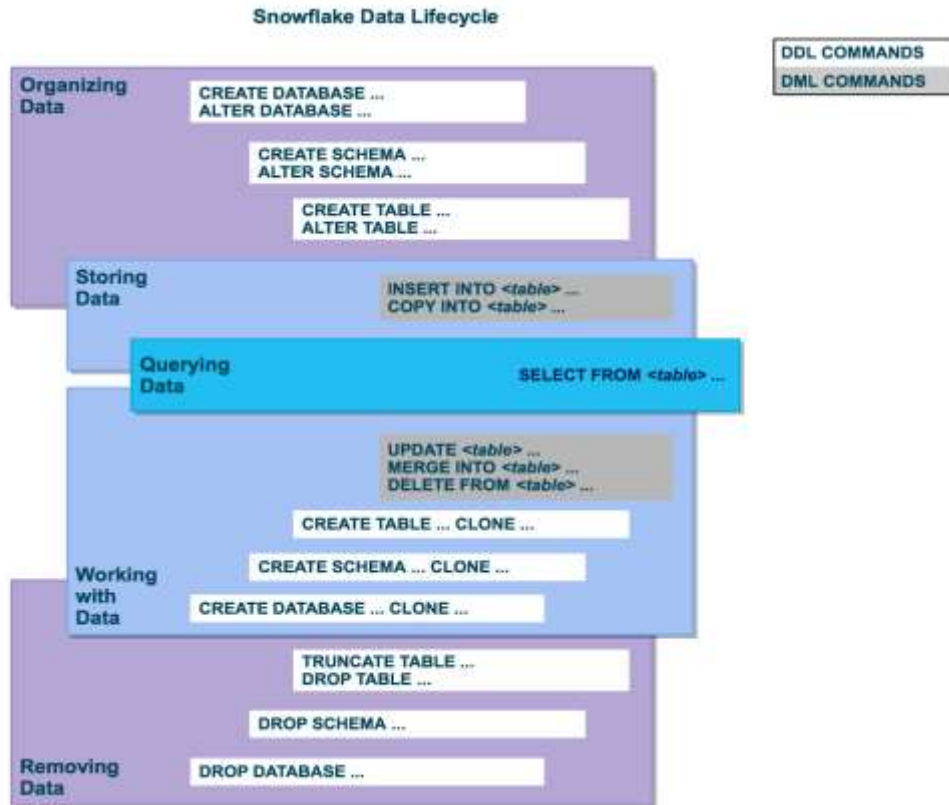
Snowflake supports multiple ways of connecting to the service:

- ❑ A web-based user interface from which all aspects of managing and using Snowflake can be accessed.
- ❑ Command line clients (e.g. SnowSQL) which can also access all aspects of managing and using Snowflake.
- ❑ ODBC and JDBC drivers that can be used by other applications (e.g. Tableau) to connect to Snowflake.
- ❑ Native connectors (e.g. Python) that can be used to develop applications for connecting to Snowflake.
- ❑ Third-party connectors that can be used to connect applications such as ETL tools (e.g. Informatica) and BI tools to Snowflake.

- ❑ The Snowflake Edition and Region that your organization choose determine the unit costs for the credits you consume and the data storage you use. Another factor that impacts unit costs is whether you have an *On Demand* or *Capacity* account:
- ❑ On Demand: Usage-based pricing with no long-term licensing requirements.
- ❑ Capacity: Discounted pricing based on an up-front Capacity commitment.
- ❑ For pricing details, go to the [pricing page](#) on the Snowflake website.

- ❑ Security and Data Protection
- ❑ SQL Support
- ❑ Tools and Interfaces
- ❑ Connectivity
- ❑ Data Import and Export

- ❑ Lifecycle Diagram
- ❑ Organizing Data
- ❑ Storing Data
- ❑ Querying Data
- ❑ Working with Data
- ❑ Removing Data



Continuous Data Protection encompasses a comprehensive set of features that help protect data stored in Snowflake against human error, malicious acts, and software or hardware failure. The features include:

- ❑ Network policies for granting or restricting users access to the site based on their IP address (i.e. IP whitelisting).
- ❑ User verification required for account access, including support for:
 - Multi-factor authentication (standard for all accounts; enabled per user by Snowflake on request).
 - Federated authentication (for Snowflake Enterprise Edition).
- ❑ Access to all objects in the system through security roles.
- ❑ Automatic encryption of data:
 - 256-bit AES encryption of data at rest and in transit.
 - 128-bit or 256-bit AES encryption of all files staged for bulk loading/unloading data.
- ❑ Maintenance of historical data (i.e. data that has been changed or deleted):
 - Querying and restoring historical data using Snowflake Time Travel.
 - Disaster recovery of historical data (by Snowflake) through Snowflake Fail-safe.

- ❑ Snowflake supports most SQL defined in ANSI SQL-99, as well as parts of the SQL-2003 analytic extensions.
- ❑ <https://docs.snowflake.net/manuals/user-guide/intro-summary-sql.html>

| Feature | Supported Type/Component | Notes |
|------------------|-----------------------------------|--|
| Sources | Local files | First upload into a Snowflake internal staging location or external location (i.e. S3 bucket), then load into table. |
| | Files in S3 | Can be any user-supplied bucket in Amazon S3. |
| | Delimited (CSV, TSV, etc.) | Any single-character delimiter is supported; default is comma (i.e. CSV). |
| File Formats | JSON | |
| | Avro | Includes support for Snappy compression in Avro files. |
| | Parquet | |
| | XML | Supported as a preview feature. For more information, see Preview Features . |
| File Compression | gzip | During upload, uncompressed files are automatically compressed using gzip, unless compression is explicitly set to NONE. |
| Compressed Files | gzip, bzip2, deflate, raw deflate | Explicitly set the compression type to one of these values or use AUTO (default) to automatically detect the type. |
| Encoding | UTF-8 | |

DRIVERS AND CONNECTORS

- ❑ Overview of the Ecosystem
- ❑ Snowflake CLI Client (SnowSQL)
- ❑ Snowflake Connector for Python
- ❑ Snowflake Connector for Spark
- ❑ Snowflake Node.js Driver
- ❑ Snowflake JDBC Driver
- ❑ Snowflake ODBC Driver

- Snowflake works with a broad array of industry-leading tools and technologies.
- From ETL to business intelligence and advanced analytics, our native solutions, partnerships, and other integrations enable you to leverage Snowflake to quickly and easily deliver insights into your data.

In this Topic:

JDBC or ODBC Drivers

Data Integration and Management

Business Intelligence

Advanced Analytics

Programmatic Interfaces

SQL Editing and Querying Tools

- ❑ If you need to connect to Snowflake using a tool or technology that is not listed here, we suggest first attempting to connect through our JDBC or ODBC drivers.
- ❑ These drivers provide general, multi-purpose connection functionality for most tools and technologies that support JDBC or ODBC.
- ❑ <https://docs.snowflake.net/manuals/user-guide/ecosystem.html>

- ❑ SnowSQL is the next-generation command line client for connecting to Snowflake to execute SQL queries and perform all DDL and DML operations, including loading data into and unloading data out of database tables.
- ❑ SnowSQL (snowsql executable) can be run as an interactive shell or in batch mode through stdin or using the -f option.
- ❑ <https://docs.snowflake.net/manuals/user-guide/snowsql.html>

- ❑ The Snowflake Connector for Python provides an interface for developing Python applications that can connect to Snowflake and perform all standard operations. It provides a programming alternative to developing applications in Java or C/C++ using the Snowflake JDBC or ODBC drivers.
- ❑ The connector is a native, pure Python package that has no dependencies on JDBC or ODBC. It can be installed using pip on Linux, Mac OSX, and Windows platforms where either Python 2.7.9 (or higher) or Python 3.4.3 (or higher) is installed.
- ❑ The connector supports developing applications using the Python Database API v2 specification (PEP-249), including using the following standard API objects:
 - ❑ Connection objects for connecting to Snowflake.
 - ❑ Cursor objects for executing DDL/DML statements and queries.
 - ❑ SnowSQL, the command line client provided by Snowflake, is an example of an application developed using the connector.
- ❑ <https://docs.snowflake.net/manuals/user-guide/python-connector.html>

- ❑ The Snowflake Connector for Spark, provided as a Spark package (spark-snowflake), brings Snowflake into the Spark ecosystem, enabling Snowflake to be used as a Spark data source, similar to other data sources (PostgreSQL, HDFS, S3, etc.).
- ❑ <https://docs.snowflake.net/manuals/user-guide/spark-connector.html>

- ❑ Written in pure JavaScript, the Snowflake Node.js driver provides a native asynchronous Node.js interface to Snowflake. The driver supports Node.js v4.0.0 and higher. For more information about Node.js, see nodejs.org.

The typical workflow for using the driver is:

- ❑ Establish a connection with Snowflake.
 - ❑ Execute statements, e.g. queries and DDL/DML commands.
 - ❑ Consume the results.
 - ❑ Terminate the connection.
-
- ❑ <https://docs.snowflake.net/manuals/user-guide/nodejs-driver.html>

- ❑ Snowflake provides a JDBC type 4 driver that supports core JDBC functionality. The JDBC driver must be installed in a 64-bit environment and requires Java 1.7 (or higher).
- ❑ The driver can be used with most client tools/applications that support JDBC for connecting to a database server. [sfsql](#), the now-deprecated command line client provided by Snowflake, is an example of a JDBC-based application.
- ❑ <https://docs.snowflake.net/manuals/user-guide/jdbc.html>

- ❑ Snowflake provides an ODBC driver for connecting to Snowflake using ODBC-based client applications.
- ❑ Snowflake provides specific versions of the driver for the following platforms:

Microsoft Windows (32-bit and 64-bit)

Mac OS X, v10.9 or higher

Linux (64-bit), developed and tested with:

- CentOS 6 or higher
- Ubuntu 14

- ❑ <https://docs.snowflake.net/manuals/user-guide/odbc.html>

- ❑ Download snowSQL client & ODBC driver
- ❑ Install these
- ❑ Check if PATH variable is updated with snowSQL path
- ❑ Login to SnowSQL with your credentials

WAREHOUSE ADMINISTRATION CAPABILITIES

PRE-REQ - Install Of Drivers And Snowsql

- ❑ Overview of Using the Web Interface
- ❑ Resource Monitors

- Creating and modifying all Snowflake objects, including virtual warehouses, databases, and all database objects
- Creating and managing users and other account-level objects, if you have the necessary administrator roles
- Loading data into tables
- Submitting and monitoring queries
- Admin Control

- ❑ **View the Snowflake Documentation**
- ❑ **Visit the Support Portal**
- ❑ **Download**
- ❑ **Help Panel**



- ❑ A virtual warehouse consumes Snowflake credits for each hour (or portion of an hour) that it runs. The number of credits consumed depends on the size of the warehouse and how long it runs.
- ❑ To help control costs and avoid unexpected credit usage related to using warehouses, Snowflake provides resource monitors.
- ❑ Resource Monitors can be used by account administrators to impose limits on the number of Snowflake credits that warehouses can consume within each monthly billing period.
- ❑ Resource monitors provide fine-grained control, allowing limits to be applied selectively to specific roles.

- Credit quota:

The number of Snowflake credits allocated to the monitor for each monthly billing period.

- Credit usage:

The number of Snowflake credits consumed during the current monthly billing period.

- Triggers:

A trigger specifies a credit quota threshold (as a percentage) and an action to perform when the threshold is reached. If credit usage reaches the threshold within the monthly billing period, the trigger fires and the action is performed. The following actions are supported:

- Suspend a warehouse once the statements being executed by the warehouse have completed.
- Suspend a warehouse immediately, which aborts any statements being executed by the warehouse at the moment.

Snowflake provides the following DDL commands for creating, managing, and using resource monitors:

- ❑ CREATE RESOURCE MONITOR
- ❑ ALTER RESOURCE MONITOR
- ❑ SHOW RESOURCE MONITORS
- ❑ DROP RESOURCE MONITOR

- Syntax

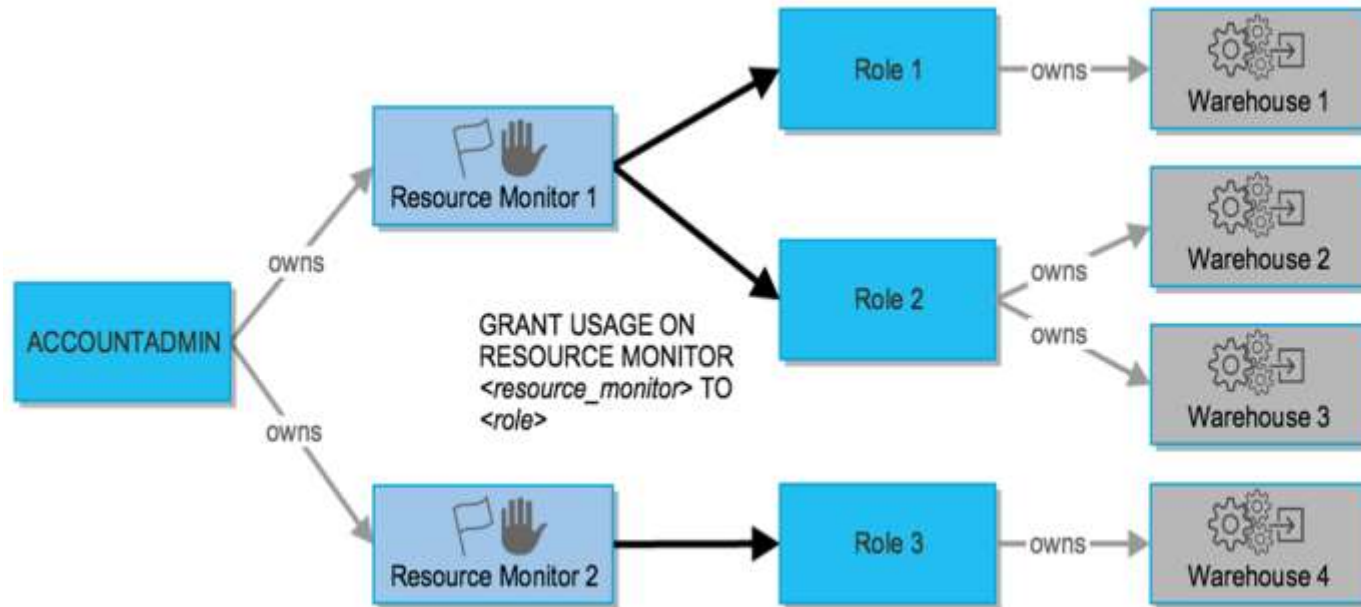
```
CREATE [ OR REPLACE ] RESOURCE MONITOR <identifier> WITH  
  [ CREDIT_QUOTA = <number> ]  
  [ TRIGGERS <trigger1_definition> <trigger2_definition> ... ]
```

- Where:

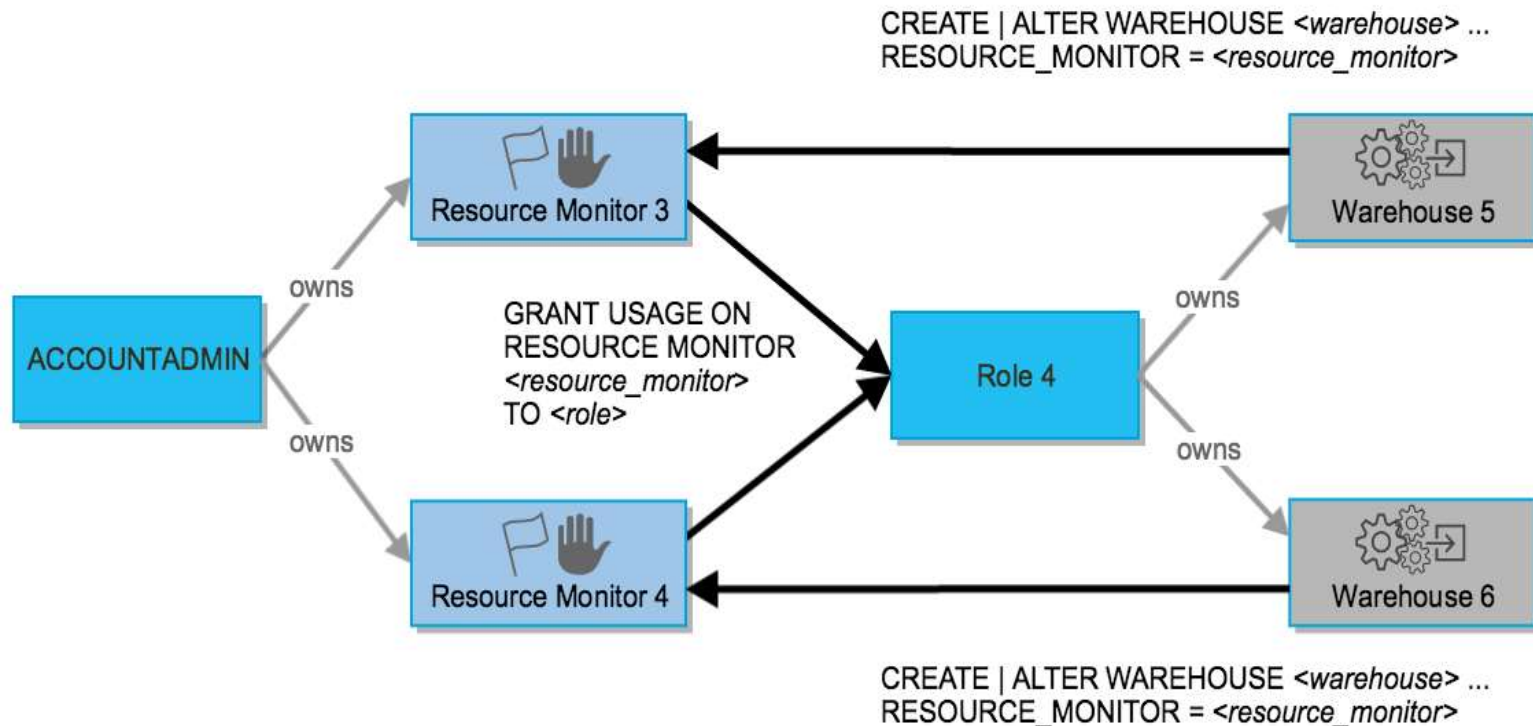
-- <trigger_definition> is:

```
ON <threshold> PERCENT DO <action>
```

GRANTING RESOURCE MONITOR USAGE TO ROLES (REQUIRED)



EXPLICITLY ASSIGNING A RESOURCE MONITOR TO A WAREHOUSE (OPTIONAL)



DATA DEFINITION AND MANIPULATION

- ❑ **PRE-REQ Creating Of Warehouse And DB , Applying Roles & Security**
- ❑ DDL Commands
- ❑ DML Commands

- Data Definition Language (DDL) commands are used to create, manipulate, and modify objects in Snowflake, such as users, virtual warehouses, databases, schemas, tables, views, columns, and functions.
- The following commands serve as the base for all DDL commands:
 - ALTER
 - COMMENT
 - CREATE
 - DESC
 - DROP
 - SHOW
 - USE

- **ALTER**

Modifies the metadata of an object (account/user or database) or the parameters for a session.

Syntax: ALTER <object_type> <identifier> <command> ...

- **Account/User Objects:**

ALTER ACCOUNT
ALTER NETWORK POLICY
ALTER USER
ALTER ROLE
ALTER WAREHOUSE
ALTER DATABASE

- **COMMENT**

Add a comment or overwrite an existing comment for an existing object.

Syntax: COMMENT [IF EXISTS] ON <object_type> <identifier> IS '<string_literal>';

COMMENT [IF EXISTS] ON COLUMN <table_identifier>.<column_identifier> IS '<string_literal>';

Comments can be added to all objects, e.g. :

- users
- roles
- warehouses
- databases
- schemas
- tables
- views

- **CREATE**

Creates a new object of the specified type.

*Syntax: CREATE [OR REPLACE] <object_type> [IF NOT EXISTS] <identifier>
[<object_properties>]
[<object_params>]
[COMMENT = '<string_literal>']*

- **Account/User/Database Objects:**

CREATE USER
CREATE ROLE
CREATE WAREHOUSE
CREATE DATABASE
CREATE SCHEMA
CREATE TABLE
CREATE VIEW

- **DESC**

Describes the details for the specified object.

Syntax: DESC <target_type> <identifier>

- **Account/User/Database:**

DESC USER
DESC WAREHOUSE
DESC DATABASE
DESC NETWORK POLICY
DESC SCHEMA
DESC TABLE
DESC VIEW

- **DROP**

Removes the specified object from the system.

Syntax: DROP <target_type> [IF EXISTS] <identifier> [CASCADE | RESTRICT]

- **Account/User/Database:**

DROP USER
DROP ROLE
DROP WAREHOUSE
DROP DATABASE
DROP SCHEMA
DROP TABLE
DROP VIEW
DROP FUNCTION

- **SHOW**

Lists all the existing objects for the specified object type.

Syntax: SHOW <object_types> [LIKE '<pattern>'] [IN <parent_object_name>]

SHOW <object_types> [LIKE '<pattern>'] [IN <parent_object_type> [<parent_object_name>]]

- **Account/User/Database:**

SHOW PARAMETERS

SHOW FUNCTIONS

SHOW USERS

SHOW TRANSACTIONS

SHOW SCHEMAS

SHOW OBJECTS

SHOW TABLES

- **USE**

- Specifies the database, schema, warehouse, or role to use for the current user session:
- Specifying a warehouse for a session enables executing DML statements and queries in the session.
- If a database or schema are not specified for a session, any objects referenced in SQL statements or queries executed in the session must be qualified with the database and schema, also known as a namespace, for the object.

Syntax: USE ROLE <role_identifier>

USE [DATABASE] <db_identifier>

USE [SCHEMA] [<db_identifier>.]<schema_identifier>

DML commands for inserting, deleting, updating, and merging data in Snowflake tables:

- INSERT
- INSERT (Multi-table)
- MERGE
- UPDATE
- DELETE

INSERT:

Insert data into a table.

Syntax:

```
INSERT INTO <target_table> (<column1>,<column3>,<column3>) VALUES  
(<value1>,<value2>,<value3>)
```

```
INSERT INTO <target_table> VALUES ( <value1>, <value2>, ... )
```

INSERT (Multi-table):

Inserts one or more rows with specified column values into multiple tables. Supports both conditional and unconditional inserts.

Syntax:

```
INSERT ALL  
  intoClause+  
<subquery>
```

```
INSERT { FIRST | ALL }  
  { WHEN <condition> THEN intoClause+ }+  
  [ ELSE intoClause ]  
<subquery>
```

MERGE:

Inserts, updates, and deletes data in a table based on data in a second table or a subquery. This can be useful if the second table is a change log that contains new rows to be added, modified rows to be updated, and marked rows to be deleted.

Syntax: MERGE INTO <target_table> USING <source> ON <join_expr> matchedClause | notMatchedClause ...

-- Where:

matchedClause ::=

*WHEN MATCHED [AND <case_predicate>] THEN
DELETE | UPDATE SET <column> = <update_expr> ...*

notMatchedClause ::=

*WHEN NOT MATCHED [AND <case_predicate>] THEN
INSERT [(<column_list>)] VALUES (<insert_expr> ...)*

UPDATE:

Updates specified rows in the target table with new values.

Syntax:

```
UPDATE <target_table>  
  SET <column1> = <value>, <column2> = <value>, ...  
  [FROM <additional_tables>]  
  [WHERE <condition_query>]
```

DELETE:

Remove data from a table.

Syntax:

```
DELETE FROM <target_table>  
[USING <additional_tables>]  
[WHERE <condition_query>];
```

DATA LOADING AND UNLOADING

- ❑ Databases & Tables
- ❑ Data Loading
- ❑ Data Unloading
- ❑ Queries
- ❑ Binary Data
- ❑ Dates & Timestamps
- ❑ Semi-structured Data
- ❑ **Snowflake Time Travel & Fail-safe**
- ❑ Sample Data Sets



- Snowflake supports bulk import (i.e. loading) of data from one or more files into tables in Snowflake databases
- Supported file formats for data loading:
 - Any flat, delimited plain text format (comma-separated values, tab-separated values, etc.)
 - Semi-structured data in JSON, Avro, Parquet, or XML format (XML is currently supported as a preview feature)
 - Supported character encoding format for data files: UTF-8

Internal Locations – use PUT while uploading data file to staging

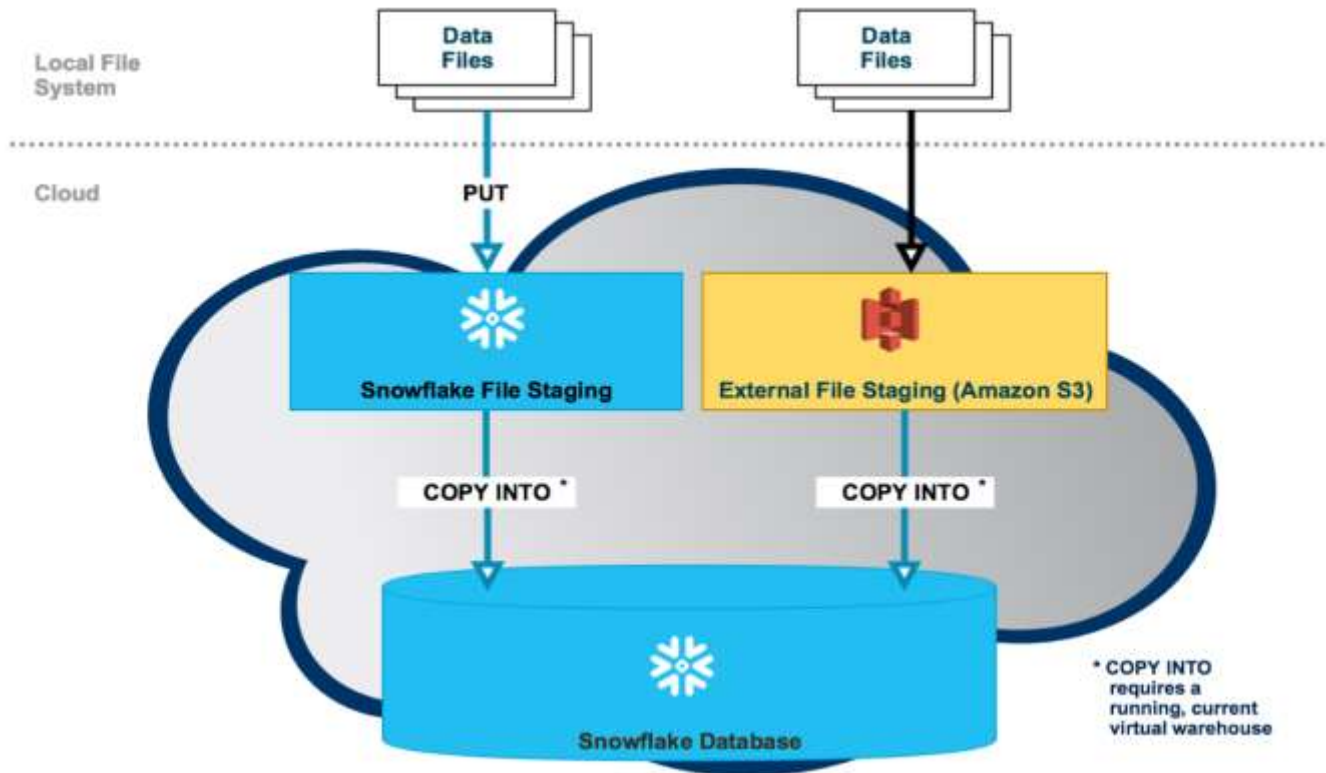
User Staging location (inside VWH)– accessed using “~”

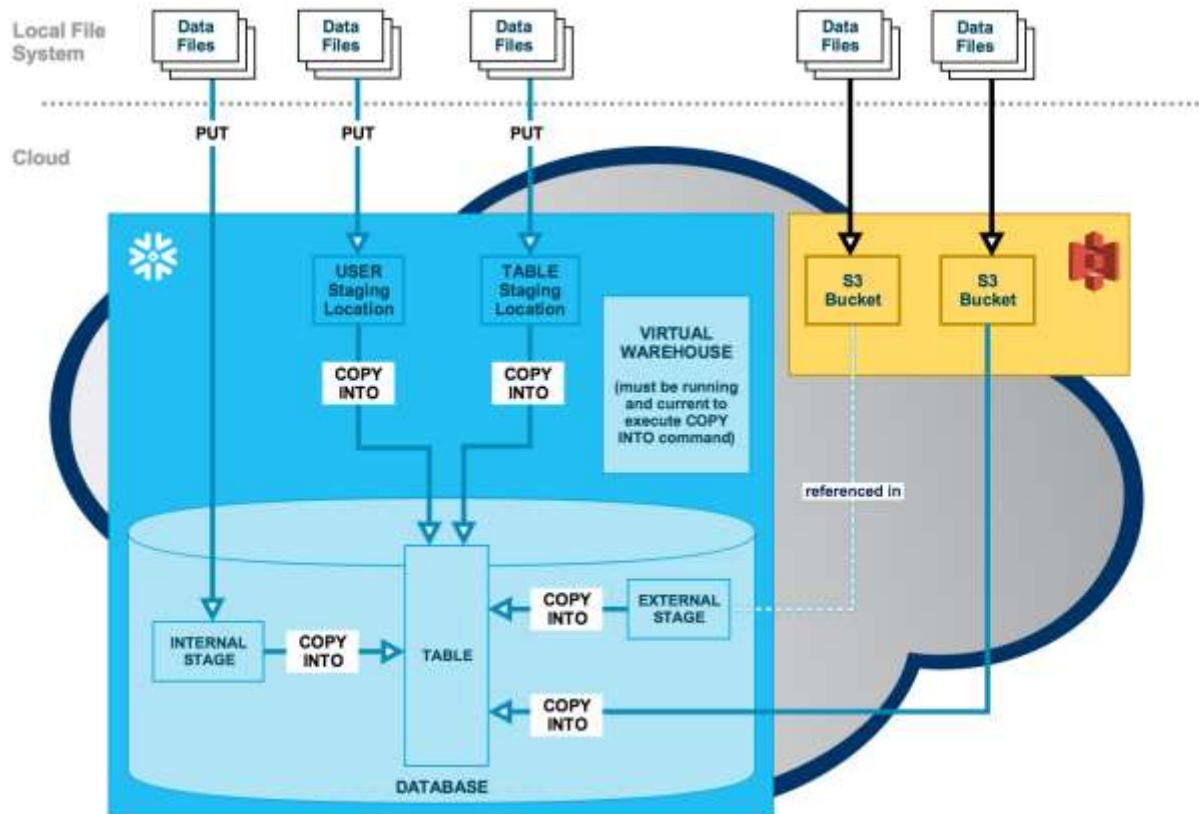
Table Staging location (inside VWH) – accessed using “%”

Internal Stage (inside database)

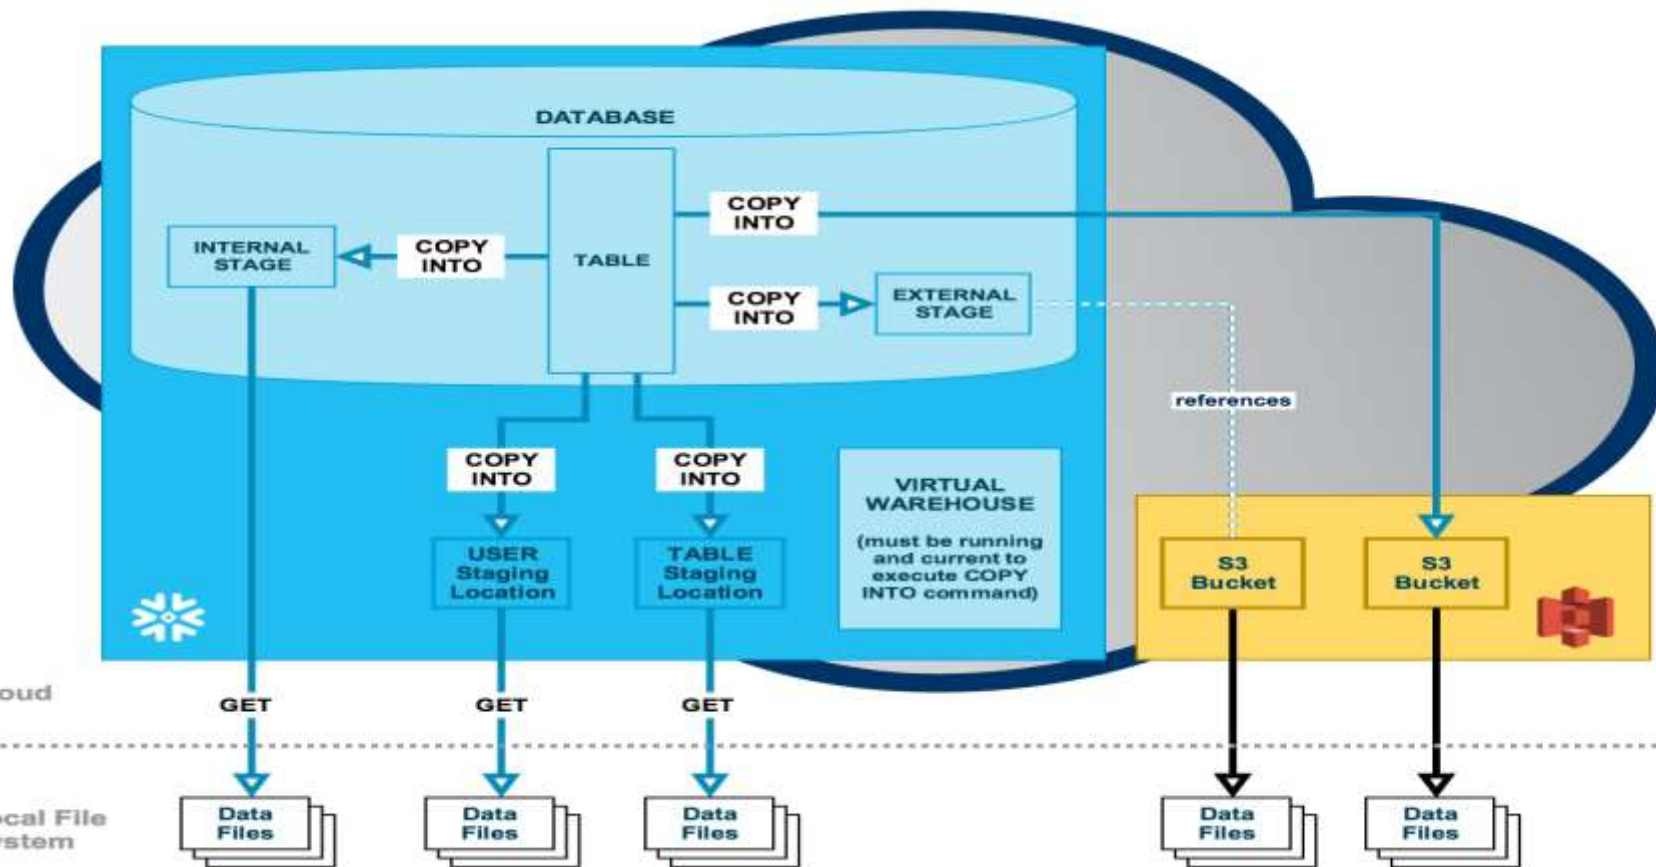
External Locations

S3 Buckets





- Similar to data loading, Snowflake supports bulk export (i.e. unloading) of data from database tables into flat, delimited text files.
- *Required Information and Metadata for Data Unloading*
- *Unloading into Snowflake Locations*
- *Unloading into External Locations*



- Snowflake supports specifying a `SELECT` statement instead of a table in the `COPY INTO` location command.
- `SELECT` statements in `COPY INTO` location support the full syntax and semantics of Snowflake SQL queries, including `JOIN` clauses, which enables downloading data from multiple tables.



- Snowflake supports standard SQL, including a subset of ANSI SQL-99 and the ANSI SQL-2003 analytic extensions

In this Topic:

Querying Semi-structured Data

Querying Metadata Columns

Using Sequences in Queries



- For each data file staged to an internal stage or staging location, Snowflake can generate additional columns which contain metadata about the data being copied.
- Currently, the following metadata columns can be queried or copied into tables:
- **METADATA\$FILENAME** - Name of the staged data file the current row belongs to
- **METADATA\$FILE_ROW_NUMBER** - Row number for each record in the container staged data file

In this Topic:

Querying Metadata Columns in Staged Files

Loading Metadata Columns into a Table

- You can query metadata columns just as you would other columns.
- Example:

```
-- Create an internal stage
create or replace stage mystage;

-- Stage a data file
put file:///tmp/store.csv @mystage;

-- Query the filename and row number metadata columns and the regular data columns in the staged file:
select metadata$filename, metadata$file_row_number, t.$1, t.$2, t.$3, t.$4 from
@mystage/store.csv.gz t;
```

- **NOTE:** NULL in t.\$4

- The COPY INTO table command can copy metadata from staged data files into a target table.
- Example

```
create or replace table rajtable (  
  Filename string, row_id integer, order_id string, order_date date, ship_date date);
```

If header is removed from file

```
copy into rajtable(Filename, row_id, order_id, order_date, ship_date)  
from (select metadata$filename, metadata$file_row_number, t.$1, t.$2, t.$3  
from @mystage/store.csv.gz t);
```

If header is included use a File Format

```
copy into rajtable(Filename, row_id, order_id, order_date, ship_date)  
from (select metadata$filename, metadata$file_row_number, t.$1, t.$2, t.$3  
from @mystage/store.csv.gz t) FILE_FORMAT = ( FORMAT_NAME = 'CSV_WITH_HEADER' TYPE = 'gzip');
```

```
select * from rajtable;
```

- Snowflake's BINARY data type allows you to store arbitrary binary data.

In this Topic:

*Binary Input and Output
Using Binary Data*

- Supported Binary Formats
 - Hex
 - Base64
 - UTF-8
- Session Parameters for Binary Values
- **BINARY_INPUT_FORMAT** - Used for loading data into Snowflake and for performing conversion to binary in the one-argument version of TO_CHAR , TO_VARCHAR
- **BINARY_OUTPUT_FORMAT** - Used for displaying binary result sets, unloading data from Snowflake, and performing conversion to string in the one-argument version of TO_BINARY.

Converting Between Hex and Base64:

The BINARY data type can be used as an intermediate step when converting between hex and base64 strings.

- Convert from hex to base64 using TO_CHAR , TO_VARCHAR:

```
select c1, to_char(to_binary(c1, 'hex'), 'base64') from hex_strings;
```

- Convert from base64 to hex:

```
select c1, to_char(to_binary(c1, 'base64'), 'hex') from base64_strings;
```

Converting Between Text and UTF-8 Bytes:

Strings in Snowflake are composed of Unicode characters, while binary values are composed of bytes. By converting a string to a binary value with the UTF-8 format, we can directly manipulate the bytes that make up the Unicode characters.

- Convert single-character strings to their UTF-8 representation in bytes using TO_BINARY:

```
select c1, to_binary(c1, 'utf-8') from characters;
```

- Convert a UTF-8 byte sequence to a string using TO_CHAR , TO_VARCHAR:

```
select to_char(X'41424320E29D84', 'utf-8');
```

- Snowflake's date, timestamp, and time zone support enables you to calculate and store consistent information about the time of events and transactions.
- Session Parameters for Dates, Times, and Timestamps:

Input Formats:

- `DATE_INPUT_FORMAT`
- `TIME_INPUT_FORMAT`
- `TIMESTAMP_INPUT_FORMAT`

Output Formats:

- `DATE_OUTPUT_FORMAT`
- `TIME_OUTPUT_FORMAT`
- `TIMESTAMP_OUTPUT_FORMAT`

- **DATE_INPUT_FORMAT**

Type:

Session — Can be set for Account ► User ► Session

Data Type:

String

Description:

Input format for dates or AUTO. AUTO specifies that Snowflake attempts to automatically detect the format of dates stored in the system during the session. For more information about date and time input and output formats, see Date and Time Input / Output.

Default Value:

AUTO

- **TIME_INPUT_FORMAT**

| | |
|----------------|--|
| Type: | Session — Can be set for Account ▶ User ▶ Session |
| Data Type: | String |
| Description: | Input format for times or AUTO. AUTO specifies that Snowflake attempts to automatically detect the format of times stored in the system during the session. For more information about date and time input and output formats, see Date and Time Input / Output. |
| Default Value: | AUTO |

- **TIMESTAMP_INPUT_FORMAT**

Type:

Session — Can be set for Account ► User ► Session

Data Type:

String

Description:

Input format for timestamps or AUTO. AUTO specifies that Snowflake attempts to automatically detect the format of timestamps stored in the system during the session. For more information about date and time input and output formats, see Date and Time Input / Output.

Default Value:

AUTO

- **DATE_OUTPUT_FORMAT**

Type:

Session — Can be set for Account
► User ► Session

Data Type:

String

Description:

Display format for date. For more information about date and time input and output formats, see Date and Time Input / Output.

Default Value:

YYYY-MM-DD

- **TIME_OUTPUT_FORMAT**

Type:

Session — Can be set for Account
► User ► Session

Data Type:

String

Description:

Display format for TIME. For more information about date and time input and output formats, see Date and Time Input / Output.

Default Value:

HH24:MI:SS

- **TIMESTAMP_OUTPUT_FORMAT**

Type:

Session — Can be set for Account
► User ► Session

Data Type:

String

Description:

Display format for TIMESTAMP.
For more information about date
and time input and output formats,
see Date and Time Input / Output.

Default Value:

YYYY-MM-DD HH24:MI:SS.FF3
TZHTZM

- To support these new types of data, semi-structured data formats, such as JSON, Avro, Parquet, and XML, with their support for flexible schemas, have become popular standards for transporting and storing data.

Snowflake provides native support for semi-structured data, including:

- Flexible-schema data types for loading semi-structured data without transformation
- Automatic conversion of data to optimized internal storage format
- Database optimization for fast and efficient SQL querying

Snowflake provides powerful CDP features for ensuring the maintenance and availability of your historical data (i.e. data that has been changed or deleted):

- Querying, cloning, and restoring historical data in tables, schemas, and databases for up to 90 days through Snowflake Time Travel.
- Disaster recovery of historical data (by Snowflake) through Snowflake Fail-safe.

Next Topics:

Understanding & Using Time Travel

Understanding & Viewing Fail-safe

Storage Costs for Time Travel and Fail-safe

- Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It provides powerful tools for restoring data that may have been accidentally or intentionally deleted
- What is Time Travel?



Data Retention Period:

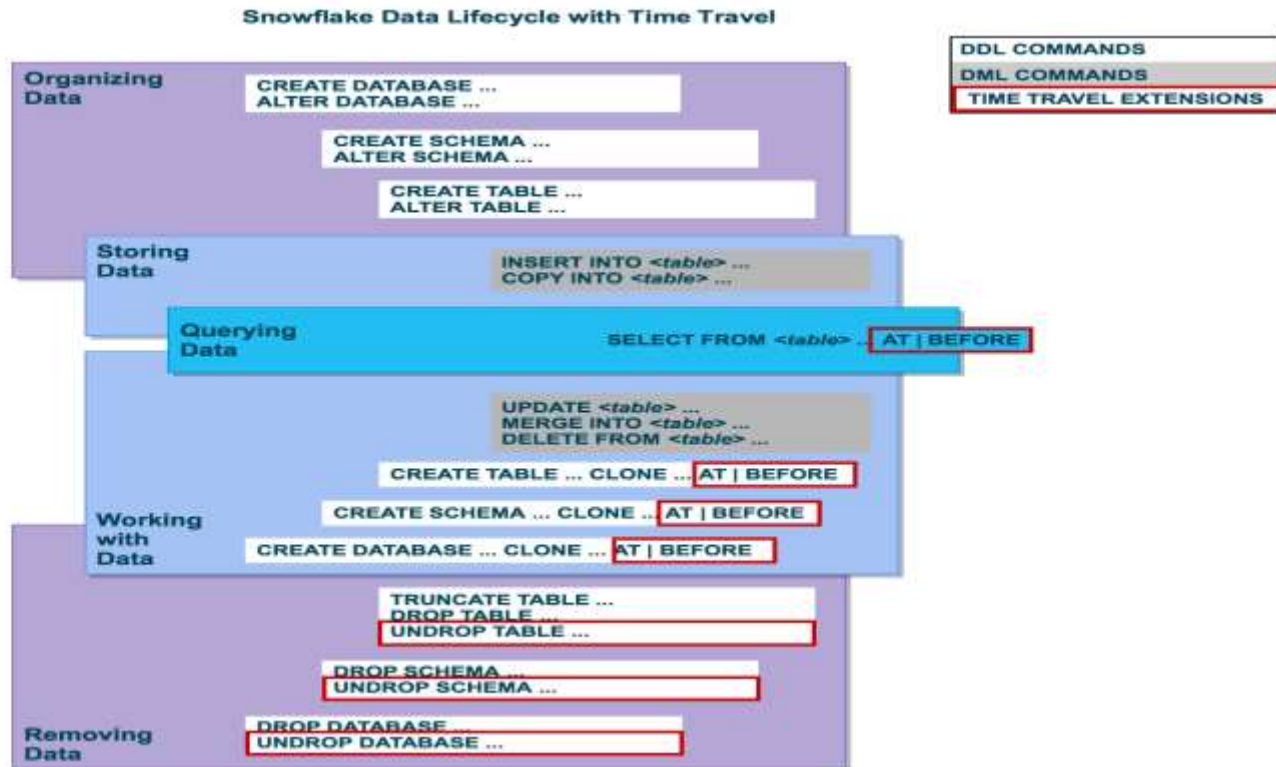
The standard retention period is 1 day (24 hours) and is automatically enabled for all Snowflake accounts:

- For Snowflake Standard Edition, the retention period can be changed to 0 at the account and object levels (for databases, schemas, and tables).
- For Snowflake Enterprise Edition, the retention period can be changed as follows:
 - For temporary or transient objects, 0 or 1 day.
 - For permanent objects, any value from 0 up to 90 days.

Time Travel SQL Extensions:

To support Time Travel actions, the following SQL extensions are provided:

- **AT | BEFORE** clause which can be specified in SELECT statements and **CREATE...CLONE** commands (immediately after the table name). The clause uses one of the following parameters to pinpoint the exact historical data you wish to access:
 - TIMESTAMP
 - OFFSET (time difference in seconds from the present time)
 - STATEMENT (identifier for statement, e.g. query ID)
- **UNDROP** command for tables, schemas, and databases.



Dropping and Restoring Objects:

To drop a table, schema, or database, use the following commands:

- DROP TABLE
- DROP SCHEMA
- DROP DATABASE

A dropped object that has not been purged from the system can be restored using the following commands:

- UNDROP TABLE
- UNDROP SCHEMA
- UNDROP DATABASE

- ❑ Data Types
- ❑ SQL Format Models
- ❑ Object Name Resolution
- ❑ Constraints
- ❑ Query Syntax
- ❑ Transactions
- ❑ Table Functions
- ❑ User-Defined Functions
- ❑ Information Schema



- Snowflake supports most basic SQL data types. Data types are automatically coerced when necessary and appropriate.
- Different Data Types:
 - **Numeric Data Types**
 - **Text and Binary Data Types**
 - **Logical Data Types**
 - **Date and Time Data Types**
 - **Semi-Structured Data Types**
 - **Unsupported Data Types**

- **NUMBER**

Fixed-point numbers of up to 38 digits, with precision and scale specified. By default, precision is 38 and scale is 0, i.e. `NUMBER(38, 0)`.

- **FLOAT, FLOAT4 , FLOAT8**

Floating point numbers. Snowflake uses double-precision (64 bit) IEEE 754 floating point numbers.


```
create or replace table test_number(num number,  
num10 number(10,1),  
dec decimal(20,2),  
numeric numeric(30,3),  
int int,  
integer integer,  
d double, f  
float,  
dp double precision,  
r real );
```

- **VARCHAR**

Text data type. The maximum length is 16MB (uncompressed). The maximum number of Unicode characters that can be stored in a VARCHAR column depends on whether the characters are single-byte or multi-byte:

Single-byte: 16,777,216

Multi-byte: Between 8,388,609 (2 bytes per character)
and 4,194,304 (4 bytes per character)

- **BINARY**

BINARY data type. The maximum length is 8MB. Unlike VARCHAR, the BINARY data type has no notion of Unicode characters, so the length is always measured in terms of bytes.

- **BOOLEAN**

BOOLEAN can have TRUE or FALSE values. BOOLEAN can also have an “unknown” value, which is represented by NULL. Boolean columns can be used in expressions (e.g. SELECT list), as well as predicates (e.g. WHERE clause).

BOOLEAN Example:

```
create or replace table test_Boolean  
( b boolean,  
  n number,  
  s string);
```

```
insert into test_boolean values (true, 1, 'yes'), (false, 0, 'no'), (null, null, null);
```

```
select * from test_boolean;
```

- **DATE**

Snowflake supports a single DATE data type for storing dates (without time). DATE accepts dates in the most common forms such as YYYY-MM-DD or DD-MON-YYYY etc. All accepted timestamps are valid inputs for dates as well

- **TIME**

Snowflake supports a single TIME data type for storing times in the form of HH:MI:SS. TIME supports an optional precision parameter for fractional seconds, e.g. TIME(3)

- **TIMESTAMP , TIMESTAMP_***

Snowflake supports three variations of the TIMESTAMP data type, as well as a special TIMESTAMP alias. All TIMESTAMP variations, including the alias, supports an optional precision parameter for fractional seconds, e.g. TIMESTAMP(3).

- Snowflake supports SQL queries that access semi-structured data using special operators and functions
- the examples in this topic refer to a table named **customers** that contains a single VARIANT column named **src**:

```
select src:name from customers;
```

```
+-----+  
| SRC:NAME |  
+-----+  
| "Patrick Jones" |  
| "Anna Glass" |  
+-----+
```

The following data types are used to represent arbitrary data structures which can be used to import and operate on semi-structured data (JSON, Avro, etc.)

- **VARIANT**

A tagged universal type, which can store values of any other type, including OBJECT and ARRAY, up to a maximum size of 16MB.

- **OBJECT**

Used to represent collections of key-value pairs, where the key is a non-empty string, and the value is a value of VARIANT type. Snowflake does not currently support explicitly-typed objects.

- **ARRAY**

Used to represent dense or sparse arrays of arbitrary size, where index is a non-negative integer (up to $2^{31}-1$), and values have VARIANT type.

- **LOB (Large Object)**

Snowflake does not support LOB data types (e.g. BLOB, CLOB).

- SQL format models (i.e. literals containing format strings) are used to specify conversion of numeric values to and from text strings. Format models can be used as arguments in the following functions:
- **TO_CHAR , TO_VARCHAR**
- **TO_DECIMAL , TO_NUMBER , TO_NUMERIC**

In this Topic:

Format Elements and Literals

Format Modifiers and Generic Space Handling

Using the Fill Mode Modifier and Printing

Parsing Input Strings

Numeric Format Models



A format model consists of a sequence of format elements and literals:

- A format element is a case-insensitive sequence of digits or letters, and, in some cases, signs (see tables below). Format elements can be directly concatenated to each other.
- Literals can be quoted (within double quotes) sequences of arbitrary characters, (a double quote is represented as two adjacent double quotes), or sequences of spaces or symbols from the following table:

| Symbol/Character | Notes |
|--------------------|--|
| . | In numeric models, treated as a format element when following 0, 9, or X; otherwise preserved as-is. |
| , | In numeric models, treated as a format element when following 0, 9, or X; otherwise preserved as-is. |
| ; | Always preserved as-is. |
| : | Always preserved as-is. |
| - | Always preserved as-is. |
| = | Always preserved as-is. |
| / | Always preserved as-is. |
| (| Always preserved as-is. |
|) | Always preserved as-is. |
| <i>blank space</i> | Always preserved as-is. |
| <i>tab</i> | Always preserved as-is. |

<https://docs.snowflake.net/manuals/sql-reference/sql-format-models.html>

- The following format elements control printing and matching input text, and are common to all format models:

| Element | Description |
|----------------|---|
| _ (underscore) | Nothing printed; optional space on input. |
| FM | Fill mode modifier; toggles fill mode for subsequent elements. |
| FX | Exact match modifier; toggles exact match mode for subsequent elements. |

- Numeric format models can be either fixed-position (with explicit placement of digits where 0, 9 or x elements are placed) or text-minimal (TM, TME and TM9 elements) ; these elements cannot be intermingled within the same model.

The following table lists all the supported elements for numeric format models:

- The **Usage** column indicates how each element can be used:
 - Fixed: Element can be used only for fixed-position numbers.
 - Minimal: Element can be used only for text-minimal models.

- Constraints define integrity and consistency rules for data stored in tables
- Snowflake provides the following constraint functionality:
 - Unique, primary, and foreign keys, and NOT NULL columns.
 - Named constraints.
 - Single-column and multi-column constraints.
 - Creation of constraints inline and out-of-line.
 - Support for creation, modification and deletion of constraints.

Snowflake supports the following constraint types from the ANSI SQL standard:

- UNIQUE
- PRIMARY KEY
- FOREIGN KEY
- NOT NULL

Query operators allow queries to be combined using set operators.

- **INTERSECT**

Returns rows from one query's result set which also appear in another query's result set, with duplicate elimination.

Syntax: SELECT ... FROM ... INTERSECT SELECT ... FROM

- **MINUS / EXCEPT**

Removes rows from one query's result set which appear in another query's result set, with duplicate elimination. The MINUS and EXCEPT keywords have the same meaning.

Syntax: SELECT ... FROM ... MINUS SELECT ... FROM
SELECT ... FROM ... EXCEPT SELECT ... FROM

- A transaction is a set of SQL statements, both reads and writes, that are processed as a unit.
- All the statements in the transaction are either applied (i.e. committed) or undone (i.e. rolled back).

In this Topic:

Scope of a Snowflake Transaction

Autocommit

Statement Rollback

Isolation Level

Read Committed Isolation

Resource Locking

Lock Wait Timeout

Deadlocks

Allowing Statement Errors to Abort Transactions

Transaction Commands and Functions

Aborting Transactions

- A transaction is associated with a single session. Multiple sessions cannot share the same transaction.
- A transaction can be started explicitly by executing a BEGIN statement. After a transaction has been started, it must be closed by executing either a COMMIT or ROLLBACK statement.
- If a session with an open transaction is closed, then the open transaction is rolled back.
- Every Snowflake transaction is assigned a unique start time (includes milliseconds), which serves as the ID for the transaction.

- **Autocommit**

A DML statement executed without explicitly starting a transaction is automatically committed on success or rolled back on failure at the end of the statement.

- **Statement Rollback**

If a DML statement executed in an explicitly-started transaction fails, the changes made by the DML will be rolled back. However, the transaction will be kept open, until the transaction is committed or rolled back.

- Transactional operations acquire locks on resources while they are being modified. This blocks other statements from modifying the resource concurrently until the lock holder completes the changes.

There are two types of locks:

- **Partition:** DML operations acquire partition locks within a table only for micro-partitions they might modify. This schema allows INSERT statements to add micro-partitions to the same table without locking any existing partitions, i.e. inserts can execute without being blocked by other DMLs.
- **Table:** A transaction only acquires a table lock during the commit phase, a length of time proportional to the amount of data inserted or modified in the transaction but overall very short. During the commit phase, INSERT and DML operations can continue. A table lock only blocks a separate transaction from performing a concurrent commit.

- If a transaction is running in a session and the session disconnects abruptly, preventing the transaction from committing or rolling back, the transaction will be left in a detached state, including any locks that the transaction is holding on resources. If this happens, you will likely want/need to abort the transaction.
- To abort a running transaction, the user who started the transaction or an account administrator can call the system function, `SYSTEM$ABORT_TRANSACTION`.

- Table functions return a set of rows instead of a single scalar value.
- **General Syntax**
- *SELECT * FROM <input_table> T, LATERAL <table_function>(T.<col1>, ..., <argN>) F;*

In this Topic:

General Syntax

List of Table Functions

| Sub-category | Function | Notes |
|--|---|--|
| Data Loading | •VALIDATE | For more information, see Data Loading. |
| Data Generation | •GENERATOR | |
| Semi-structured Queries | •FLATTEN | For more information, see Querying Semi-structured Data. |
| Query Results | •RESULT_SCAN | |
| Query History (Information Schema) | •QUERY_HISTORY , QUERY_HISTORY_BY_* | For more information, see Information Schema. |
| Warehouse & Storage Usage History (Information Schema) | •DATABASE_STORAGE_USAGE_HISTORY •WAREHOUSE_LOAD_HISTORY •WAREHOUSE_METERING_HISTORY •STAGE_STORAGE_USAGE_HISTORY | For more information, see Information Schema. |

User-defined functions (UDFs) let you extend the system to perform operations that are not available via the system-defined functions provided by Snowflake.

Snowflake currently supports two types of UDFs:

- *SQL UDFs* evaluate and return an arbitrary SQL expression.
- *JavaScript UDFs* allow you to manipulate data using the JavaScript programming language and runtime.

- **SQL**

A SQL UDF evaluates and returns an arbitrary SQL expression. The function definition must be a SQL expression that returns a single value. The expression may be a query expression, though for non-table functions, a query expression must be guaranteed to return at most one row, containing a single column.

- **JavaScript**

JavaScript user-defined functions (UDFs) allow you to manipulate data using the JavaScript programming language and runtime. Functions are created similarly to SQL UDFs.

- ❑ Operators
- ❑ All Functions
- ❑ Bitwise Expression Functions
- ❑ Conditional Expression Functions
- ❑ Context Functions
- ❑ Numeric Functions
- ❑ String & Binary Functions
- ❑ Regular Expression (String) Functions
- ❑ Date & Time Functions
- ❑ Semi-structured Data Functions
- ❑ Conversion Functions
- ❑ Aggregate Functions
- ❑ Window/Analytic Functions
- ❑ Miscellaneous Functions



- Operators for performing arithmetic, comparison, and logical operations on expressions.

In this Topic:

Arithmetic Operators

Comparison Operators

Logical Operators

| Operator | Description |
|----------|---|
| $a + b$ | Add two numeric expressions (number or real). |
| $+ a$ | Return a , which can cause implicit conversion of a to a numeric value. |
| $a - b$ | Subtract a numeric expression from another. |
| $- a$ | Negate the numeric input expression. |
| $a * b$ | Multiply two numeric expressions. |
| a / b | Divide two numeric expressions. |
| $a \% b$ | Compute the modulo of numeric expressions a per b . |

Operator**Description** $a = b$ a is equal to b . $a \neq b$ a is not equal to b . $a < > b$ a is less than or greater than b . $a > b$ a is greater than b . $a \geq b$ a is greater than or equal to b . $a < b$ a is less than b . $a \leq b$ a is less than or equal to b .

| Operator | Description |
|----------|---------------------------|
| a AND b | Matches both inputs. |
| a OR b | Matches either input. |
| NOT a | Does not match the input. |

- This topic provides a list of all system-defined Snowflake functions, including table functions, in alphabetical order.
- Find the entire list here <https://docs.snowflake.net/manuals/sql-reference/functions-all.html>

- ❑ BITAND
- ❑ BITAND_AGG
- ❑ BITNOT
- ❑ BITOR
- ❑ BITOR_AGG
- ❑ BITSHIFTLEFT
- ❑ BITSHIFTRIGHT
- ❑ BITXOR
- ❑ BITXOR_AGG

- ❑ `select BITAND(3, 4) from table1`
- ❑ `select BITOR(3, 4) from table1`

- ❑ BETWEEN
- ❑ CASE
- ❑ COALESCE
- ❑ DECODE
- ❑ EQUAL_NULL
- ❑ GREATEST
- ❑ IFF
- ❑ IFNULL
- ❑ IN / NOT IN
- ❑ IS NULL / IS NOT NULL
- ❑ IS_NULL_VALUE
- ❑ LEAST
- ❑ NULLIF
- ❑ NVL
- ❑ NVL2
- ❑ REGR_VALX
- ❑ REGR_VALY
- ❑ ZEROIFNULL

- ❑ `select column1, case
 when column1=1 then 'one'
 when column1=2 then 'two'
 else 'other'
end as result from (values(1),(2),(3)) v;`
- ❑ `select val, iff(val::int = val, 'integer', 'non-integer')`
`from (select column1 as val`
 `from values(1.0), (1.1), (-3.1415), (-5.000), (null));`
- ❑ `select a, b, nullif(a,b) from i;`

General Context

- ❑ CURRENT_CLIENT
- ❑ CURRENT_DATE
- ❑ CURRENT_TIME
- ❑ CURRENT_TIMESTAMP
- ❑ CURRENT_VERSION
- ❑ LOCALTIME
- ❑ LOCALTIMESTAMP

Session Context

- ❑ CURRENT_SESSION
- ❑ CURRENT_STATEMENT
- ❑ CURRENT_TRANSACTION
- ❑ CURRENT_USER
- ❑ LAST_QUERY_ID
- ❑ LAST_TRANSACTION

Session Object Context

- ❑ CURRENT_DATABASE
- ❑ CURRENT_ROLE
- ❑ CURRENT_SCHEMA
- ❑ CURRENT_SCHEMAS
- ❑ CURRENT_WAREHOUSE

- Display the current warehouse, database, and schema for the session:

```
select current_warehouse(), current_database(), current_schema();
```

- Display the current date, time, and timestamp (note that parentheses are not required to call these functions):

```
select current_date, current_time, current_timestamp;
```

Rounding and Truncation

- ❑ ABS
- ❑ CEIL
- ❑ FLOOR
- ❑ MOD
- ❑ ROUND
- ❑ SIGN
- ❑ TRUNC(for date DATE_TRUNC) , TRUNCATE

Exponent and Root

- ❑ CBRT
- ❑ EXP
- ❑ POW, POWER
- ❑ SQRT
- ❑ SQUARE

Logarithmic

- ❑ LN
- ❑ LOG

Trigonometric

- ❑ ACOS
- ❑ ACOSH
- ❑ ASIN
- ❑ ASINH
- ❑ ATAN
- ❑ ATAN2
- ❑ ATANH
- ❑ COS
- ❑ COSH

COT

- ❑ DEGREES
- ❑ HAVERSINE
- ❑ PI
- ❑ RADIANS
- ❑ SIN
- ❑ SINH
- ❑ TAN
- ❑ TANH

- ❑ `select n, scale, round(n, scale) from tab2;`
- ❑ `select column1, square(column1) from (values (0), (1), (-2), (3.15), (null)) v;`
- ❑ `select sin(0), sin(pi()/3), sin(radians(90));`

General

- ❑ ASCII
- ❑ BIT_LENGTH
- ❑ CHARINDEX
- ❑ CHR , CHAR
- ❑ CONCAT , ||
- ❑ CONTAINS
- ❑ EDITDISTANCE
- ❑ ENDSWITH
- ❑ ILIKE
- ❑ INITCAP
- ❑ INSERT
- ❑ LEFT
- ❑ LENGTH
- ❑ LIKE
- ❑ LOWER
- ❑ LPAD
- ❑ LTRIM
- ❑ OCTET_LENGTH
- ❑ PARSE_URL
- ❑ POSITION
- ❑ REPEAT
- ❑ REPLACE
- ❑ REVERSE
- ❑ RIGHT
- ❑ RPAD
- ❑ RTRIM
- ❑ RTRIMMED_LENGTH
- ❑ SPACE
- ❑ SPLIT
- ❑ SPLIT_PART
- ❑ STARTSWITH
- ❑ SUBSTR , SUBSTRING
- ❑ TRANSLATE
- ❑ TRIM
- ❑ UPPER
- ❑ UUID_STRING

Encode/Decode

- ❑ BASE64_DECODE_BINARY
- ❑ BASE64_DECODE_STRING
- ❑ BASE64_ENCODE
- ❑ HEX_DECODE_BINARY
- ❑ HEX_DECODE_STRING
- ❑ HEX_ENCODE
- ❑ TRY_BASE64_DECODE_BINARY
- ❑ TRY_BASE64_DECODE_STRING
- ❑ TRY_HEX_DECODE_BINARY
- ❑ TRY_HEX_DECODE_STRING

Hash/Cryptographic

- ❑ MD5 , MD5_HEX
- ❑ MD5_BINARY
- ❑ MD5_NUMBER
- ❑ SHA1 , SHA1_HEX
- ❑ SHA1_BINARY
- ❑ SHA2 , SHA2_HEX
- ❑ SHA2_BINARY

- ❑ `select v, lower(v) from lu;`
- ❑ `select v, upper(v) from lu;`
- ❑ `select '123456', pos, len, substr('123456', pos, len) from o;`
- ❑ `select`
 `md5_binary('Snowflake') as md5_binary,`
 `base64_encode(md5_binary('Snowflake'),0,$') as base64_encode,`

 `base64_decode_binary(base64_encode(md5_binary('Snowflake'),0,$'),'`
 `$') as base64_decode_binary;`
- ❑ `select sha2_binary('Snowflake', 384);`

- ❑ REGEXP
- ❑ REGEXP_COUNT
- ❑ REGEXP_INSTR
- ❑ REGEXP_LIKE
- ❑ REGEXP_REPLACE
- ❑ REGEXP_SUBSTR
- ❑ RLIKE

- ❑ `select * from rlike_ex where name regexp 'J.h.* [dD]+o.*';`
- ❑ `select body, regexp_count(body, '\\b\\S*o\\S*\\b') from message;`
- ❑ `select regexp_substr('Customers - (NY)', '\\([[:alnum:]]-\\+\\)') as customers;`

Construction

- ❑ DATE_FROM_PARTS
- ❑ TIME_FROM_PARTS
- ❑ TIMESTAMP_FROM_PARTS

Extraction

- ❑ DATE_PART
- ❑ DAYNAME
- ❑ EXTRACT
- ❑ HOUR / MINUTE / SECOND
- ❑ LAST_DAY
- ❑ MONTHNAME
- ❑ YEAR / QUARTER / MONTH / WEEK / DAY / ...



Addition/Subtraction

- ❑ ADD_MONTHS
- ❑ DATEADD
- ❑ DATEDIFF
- ❑ TIMEADD
- ❑ TIMEDIFF
- ❑ TIMESTAMPADD
- ❑ TIMESTAMPDIFF

Truncation

- ❑ DATE_TRUNC
- ❑ TRUNC

Conversion

- ❑ TO_DATE
- ❑ TO_TIME
- ❑ TO_TIMESTAMP /
TO_TIMESTAMP_

Time Zone

- ❑ CONVERT_TIMEZONE
- ❑

- ❑ `select dayname(to_date('2015-05-01')) as day;`
- ❑ `select datediff(year, '2010-04-09 14:39:20'::timestamp, '2013-05-08 23:39:20'::timestamp) as diff_years;`
- ❑ `select to_time('12:34:56') from (values(1)) v;`

JSON and XML Parsing

- ❑ CHECK_JSON
- ❑ CHECK_XML
- ❑ PARSE_JSON
- ❑ PARSE_XML
- ❑ STRIP_NULL_VALUE

Array Creation/Manipulation

- ❑ ARRAY_AGG
- ❑ ARRAY_APPEND
- ❑ ARRAY_CAT
- ❑ ARRAY_COMPACT
- ❑ ARRAY_CONSTRUCT
- ❑ ARRAY_CONSTRUCT_COMPACT
- ❑ ARRAY_INSERT
- ❑ ARRAY_PREPEND
- ❑ ARRAY_SIZE
- ❑ ARRAY_SLICE
- ❑ ARRAY_TO_STRING

Object Creation/Manipulation

- ❑ OBJECT_AGG
- ❑ OBJECT_CONSTRUCT
- ❑ OBJECT_DELETE
- ❑ OBJECT_INSERT

Extraction

- ❑ FLATTEN
- ❑ GET
- ❑ GET_PATH , :
- ❑ XMLGET

Conversion/Casting

- ❑ AS_type
- ❑ TO_ARRAY
- ❑ TO_JSON
- ❑ TO_OBJECT
- ❑ TO_VARIANT
- ❑ TO_XML

Type Predicates

- ❑ IS_type
- ❑ TYPEOF

- ❑ `select *, check_xml(xml_str) from my_table where check_xml(xml_str) is not null;`
- ❑ `select array_size(v) from colors;`
- ❑ `select n, as_real(v), typeof(v) from vartab;`
- ❑ `select * from vartab where is_decimal(v);`

Any Data Type

- ❑ CAST , ::
- ❑ TRY_CAST

Text/Character/Binary Data Types

- ❑ TO_CHAR , TO_VARCHAR
- ❑ TO_BINARY
- ❑ TRY_TO_BINARY

Numeric Data Types

- ❑ TO_DECIMAL , TO_NUMBER , TO_NUMERIC
- ❑ TO_DOUBLE
- ❑ TRY_TO_DECIMAL, TRY_TO_NUMBER, TRY_TO_NUMERIC
- ❑ TRY_TO_DOUBLE

Boolean Data Type

- ❑ TO_BOOLEAN
- ❑ TRY_TO_BOOLEAN

Date and Time Data Types

- ❑ TO_DATE
- ❑ TO_TIME
- ❑ TO_TIMESTAMP / TO_TIMESTAMP_*
- ❑ TRY_TO_DATE
- ❑ TRY_TO_TIME
- ❑ TRY_TO_TIMESTAMP / TRY_TO_TIMESTAMP_*

Semi-structured Data Types

TO_ARRAY

- ❑ TO_OBJECT
- ❑ TO_VARIANT

- ❑ `select to_date('2013-05-17') from (values(1)) v;`
- ❑ `select to_varchar('2013-04-05 01:02:03'::timestamp, 'mm/dd/yyyy, hh24:mi hours');`

General Aggregation

- ❑ ANY_VALUE
- ❑ AVG
- ❑ CORR
- ❑ COUNT
- ❑ COVAR_POP
- ❑ COVAR_SAMP
- ❑ GROUPING
- ❑ GROUPING_ID
- ❑ HASH_AGG
- ❑ LISTAGG
- ❑ MEDIAN

- ❑ MIN / MAX
- ❑ PERCENTILE_CONT
- ❑ PERCENTILE_DISC
- ❑ STDDEV
- ❑ STDDEV_POP
- ❑ STDDEV_SAMP
- ❑ SUM
- ❑ VAR_POP
- ❑ VAR_SAMP
- ❑ VARIANCE_POP
- ❑ VARIANCE , VARIANCE_SAMP

Bitwise Aggregation

- ❑ BITAND_AGG
- ❑ BITOR_AGG
- ❑ BITXOR_AGG

Linear Regression

- ❑ REGR_AVGX
- ❑ REGR_AVGY
- ❑ REGR_COUNT
- ❑ REGR_INTERCEPT
- ❑ REGR_R2
- ❑ REGR_SLOPE
- ❑ REGR_SXX
- ❑ REGR_SXY
- ❑ REGR_SYY

Semi-structured Data Aggregation

- ❑ ARRAY_AGG
- ❑ OBJECT_AGG

Cardinality Estimation (Using HyperLogLog)

- ❑ APPROX_COUNT_DISTINCT
- ❑ HLL
- ❑ HLL_ACCUMULATE
- ❑ HLL_COMBINE
- ❑ HLL_ESTIMATE
- ❑ HLL_EXPORT
- ❑ HLL_IMPORT

Similarity Estimation (Using MinHash)

- ❑ APPROXIMATE_JACCARD_INDEX
- ❑ APPROXIMATE_SIMILARITY
- ❑ MINHASH
- ❑ MINHASH_COMBINE

- ❑ `select k, min(d), max(d) from minmax_example group by k;`
- ❑ `select array_agg(distinct o_orderstatus) from orders where o_totalprice > 450000;`
- ❑ `select hll(o_orderdate),
hll_estimate(hll_import(hll_export(hll_accumulate(o_orderdate)))) from
orders;`
- ❑ `select approximate_similarity(mh) from
((select minhash(100, c5) mh from orders where c2 <= 50000)
union
(select minhash(100, c5) mh from orders where c2 > 50000));`

System Functions

- ❑ SYSTEM\$ABORT_SESSION
- ❑ SYSTEM\$ABORT_TRANSACTION
- ❑ SYSTEM\$CANCEL_ALL_QUERIES
- ❑ SYSTEM\$CANCEL_QUERY
- ❑ SYSTEM\$CLUSTERING_DEPTH
- ❑ SYSTEM\$CLUSTERING_INFORMATION
- ❑ SYSTEM\$CLUSTERING_RATIO
- ❑ SYSTEM\$TYPEOF
- ❑ SYSTEM\$WAIT

- CUME_DIST Finds the cumulative distribution of a value with regard to other values within the same window partition.
- DENSE_RANK Finds the rank of a value within a group of values, without gaps.
- FIRST_VALUE Returns the first value within an ordered group of values.
- LAG Returns the value of *expr* at negative row offset *offset*, i.e. the value of *expr* at *offset* rows before with respect to the current input row.
- LAST_VALUE
- LEAD
- NTH_VALUE Returns the nth value (up to 1000) within an ordered group of values.
- NTILE Divides an ordered data set equally into the number of buckets specified by *constant_value*. Buckets are sequentially numbered 1 through *constant_value*.
- PERCENT_RANK Returns the relative rank of a value within a group of values.
- RANK Returns the rank of a value within a group of values.
- ROW_NUMBER Returns a unique row number for each row within a window partition starting with 1 and increasing sequentially.
- WIDTH_BUCKET Constructs equi-width histograms, in which the histogram range is divided into intervals of identical size.

Utility Functions

- ❑ GET_DDL
- ❑ HASH

Data Generation Functions

- ❑ RANDOM
- ❑ SEQ1 / SEQ2 / SEQ4 / SEQ8
- ❑ UUID_STRING

Random Distribution Functions

NORMAL


RANDSTR

UNIFORM

ZIPF



- ❑ `select system$abort_session(1065153868222);`
- ❑ `select get_ddl('view', 'tpch_customer');`
- ❑ `select get_ddl('function', 'multiply(number, number)');`
- ❑ `select random(4711) from table(generator(rowcount => 5));`
- ❑ `select normal(0, 1, 1234) from table(generator(rowcount => 5));`



WELCOME TO DAY 2

- ❑ Snowflake Administration
- ❑ End To End Usage Of Snowflake
- ❑ Feedback Form
- ❑ JSON Data Upload And Query – completed on Day 2
- ❑ ETL Access
- ❑ Tableau Desktop/Server connectivity
- ❑ Q & A

In this Topic:

- ❑ *Snowflake Password Policy*
- ❑ *Creating Users*
- ❑ *Resetting a User Password (Administrators)*
- ❑ *Disabling / Enabling a User*
- ❑ *Unlocking a User*
- ❑ *Altering Session Parameters for a User*
- ❑ *Modifying Other User Properties*
- ❑ *Viewing Users*
- ❑ *Dropping a User*

Snowflake applies the following policy to passwords:

- ❑ Must be at least 8 characters long.
- ❑ Must contain at least 1 digit.
- ❑ Must contain at least 1 uppercase letter and 1 lowercase letter.

- Snowflake enables administrators to create and manage user credentials and default values using SQL. In addition, the Snowflake web interface provides a convenient wizard for creating individual users, as well as the ability to complete common activities such as resetting a user's password.
- Only users with the SECURITYADMIN role (i.e. security administrators) or higher can create, alter, or drop users.

□ Using Web Interface

Click on Account > Users> Create Button

In the User Name field, enter a unique identifier for the user and in the Password and Confirm Password fields, enter the login password for the user.

Leave the Force Password Change checkbox checked to force the user to change their password on their next login; otherwise, clear the checkbox.

Click the Next button. The Advanced screen opens.

Optionally enter the Login Name, Display Name, and personal information for the user.

Click the Next button. The Preferences screen opens.

Optionally enter defaults for the user: Virtual warehouse, Namespace in the form of <db_name> or <db_name>.<schema_name>, Role

Click the Finish button. Snowflake displays a success message.

- Using SQL

- Syntax

```
CREATE [ OR REPLACE ] USER [ IF NOT EXISTS ] <name>  
  [ <properties> ... ]  
  [ <session_params> ... ]  
  [ COMMENT = '<string_literal>' ]
```

- Example

```
create user user1 password='abc123' must_change_password = true;
```

- Using the Web Interface

Click on Account > Users.

Click on a user row to select it, then click the Reset Password button.
The Reset Password dialog opens.

Enter the new login password for the user, and confirm the password.

Leave the Force Password Change checkbox checked to force the user to change their password on their next login; otherwise, clear the checkbox.

Click the Finish button.

- Using SQL
- Syntax

```
ALTER USER [ IF EXISTS ] <name> RENAME TO <new_name>
```

```
ALTER USER [ IF EXISTS ] <name> ABORT ALL QUERIES
```

```
ALTER USER [ IF EXISTS ] <name> SET { [ <property> [ <property> ... ] ]  
    [ <session_param> [ <session_param> ... ]  
    }
```

```
ALTER USER [ IF EXISTS ] <name> UNSET { [ <property_name> [, <property_name> ... ] ] ,  
    [ <session_param_name> [, <session_param_name> ... ] ]  
    }
```

- Example

```
alter user user1 set password = 'welcome1' must_change_password = true;
```

- Using the Web Interface

Click on Account > Users.

Click on a user row to select it, then click the Disable User button. A confirmation dialog opens.

Click Yes to disable the user.

- ❑ Using SQL
- ❑ Example
- ❑ Disable a user:

```
alter user jane.smith set disabled = true;
```

- ❑ Enable a user:

```
alter user jane.smith set disabled = false;
```

- ❑ If a user login fails after five consecutive attempts, the user is locked out of their account for a period of time (currently 15 minutes). Once the period of time elapses, the system automatically clears the lock and the user can attempt to log in again.
 - ❑ Using SQL
 - ❑ Example
- ```
alter user jane.smith set mins_to_unlock= 0;
```



- ❑ Using SQL
- ❑ To show the session parameters for a user, use the following SQL syntax:  
`SHOW PARAMETERS [ LIKE '<pattern>' ] FOR USER <name>`
- ❑ To alter the session parameters for a user, use the following syntax:  
`ALTER USER <name> SET <session_param> = <value>`
- ❑ For example, allow a user to remain connected to Snowflake indefinitely without timing out:  
`alter user `jane.smith` set client_session_keep_alive = true;`
- ❑ To reset a session parameter for a user to the default value, use the following syntax:  
`ALTER USER <name> UNSET <session_param>`

- Using SQL
- Example

```
alter user jane.smith set last_name = 'Jones';
```

- Using SQL

- Syntax

DESC USER <name>

SHOW USERS [LIKE '<pattern>']

- Example

```
desc user janeksmith;
```

- Using Web Interface

Click on **Account > Users**.

Click on a user row to select it, then click the **Drop** button. A confirmation dialog opens.

Click **Yes** to drop the user.



- Using SQL

- Syntax

`DROP USER [ IF EXISTS ] <name>`

- Example

`drop user user1;`

- The Snowflake cloud architecture separates data warehousing into two distinct functions:

**virtual warehouse processing and data storage.**

- The costs associated with using Snowflake are based on your usage of the service for each of these functions.

### **In this Topic:**

*Virtual Warehouse Credit*

*Data Storage Usage*

*Managing Resource Monitors*

*Common Use Cases for Enforcing Credit Quotas*




- Warehouses come in different sizes: X-Small, Small, Medium, Large, X-Large, 2X-Large, and 4X-Large. The number of credits consumed by a warehouse per second while it is running is based directly on its size, but for the first time when it is launched credit for 60Sec is charged:

| X-Small | Small | Medium | Large | X-Large | 2X-Large | 3X-Large | 4X-Large |
|---------|-------|--------|-------|---------|----------|----------|----------|
| 1       | 2     | 4      | 8     | 16      | 32       | 64       | 128      |

- Users with the ACCOUNTADMIN role can view warehouse usage and data storage for your account:

Web Interface: Click on Account > Billing & Usage

SQL: `WAREHOUSE_METERING_HISTORY` (in the Information Schema)



The screenshot shows the Snowflake Account Billing & Usage interface. The top navigation bar includes links for Databases, Warehouses, Worksheets, History, and Account. The main section is titled "Account" and has tabs for Users, Roles, Sessions, and Billing & Usage. The Billing & Usage tab is active, showing a toggle between "Credit Usage" and "Storage Usage". The current view is "Credit Usage" for "October, 2015 (local)".

Annotations highlight key features:

- Toggle between credit usage and storage usage:** Points to the "Credit Usage" and "Storage Usage" toggle buttons.
- Select a month to view:** Points to the "October, 2015 (local)" dropdown menu.
- View details for a day:** Points to the "View details" link for a specific day in the "Usage by Day" table.
- View details for a warehouse:** Points to the "View details" link for a specific warehouse in the "Warehouse Name" table.

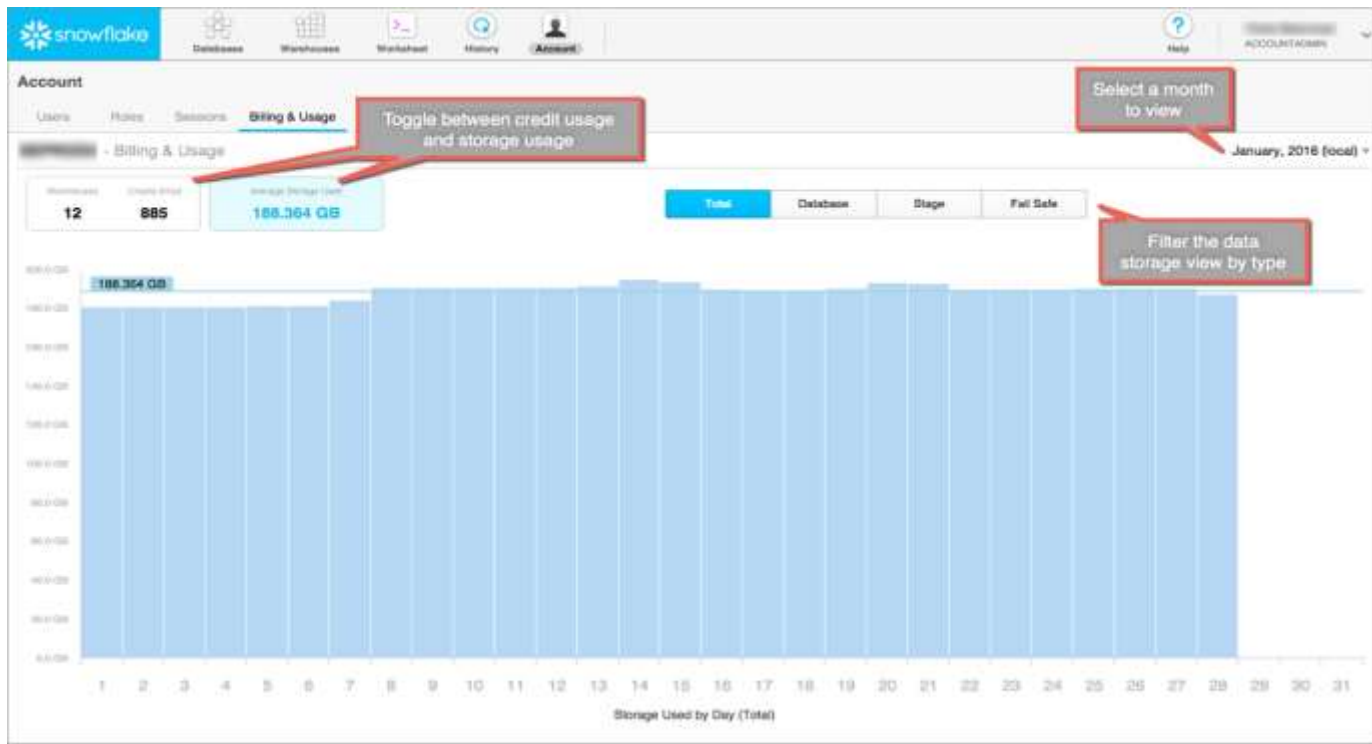
The "Usage by Day (All Warehouses)" table shows the following data:

| Date   | Credits Billed |
|--------|----------------|
| Oct 1  | 24             |
| Oct 2  | 26             |
| Oct 3  | 26             |
| Oct 4  | 24             |
| Oct 5  | 24             |
| Oct 6  | 24             |
| Oct 7  | 27             |
| Oct 8  | 27             |
| Oct 9  | 24             |
| Oct 10 | 24             |

The "Warehouse Name" table shows the following data:

| Warehouse Name | Credits Billed |
|----------------|----------------|
| WAREHOUSE1     | 745            |
| WAREHOUSE2     | 49             |
| WAREHOUSE3     | 49             |
| WAREHOUSE4     | 45             |
| WAREHOUSE5     | 10             |
| WAREHOUSE6     | 7              |
| WAREHOUSE7     | 7              |

- ❑ The monthly costs for storing data in Snowflake is based on a flat rate per terabyte (TB). The amount charged per TB depends on your type of account (Capacity or On Demand) and region (US or EU). For data storage pricing, see the pricing page on the Snowflake website.
- ❑ Users with the ACCOUNTADMIN role can view warehouse credit usage for your account:
  - ❑ Web Interface: Click on Account > Billing & Usage
  - ❑ SQL: Information Schema functions: - DATABASE\_STORAGE\_USAGE\_HISTORY - STAGE\_STORAGE\_USAGE\_HISTORY
- ❑ Users with the appropriate access privileges can use either the web interface or SQL to view the size (in bytes) of individual tables in a schema/database:
  - ❑ Web Interface: Click on Databases > db\_name > Tables
  - ❑ SQL: SHOW TABLES
- ❑ In addition, users with the ACCOUNTADMIN role can use SQL to view table size information:
  - ❑ SQL: TABLE\_STORAGE\_METRICS View (in the Information Schema)



- Defining a Global Quota for Your Account

A global credit quota for an account can be implemented by ensuring that all warehouses are associated with a single resource monitor. To accomplish this, grant the USAGE privilege for the resource monitor to the predefined PUBLIC role.

- Creating a Quota for a Group of Users

To limit the credit usage for a group of users without restricting other users, you can use multiple roles.

Snowflake provides three types of parameters that can be set for your account:

- Account parameters that affect your entire account.
- Session parameters that default to users and their sessions.
- Object parameters that default to objects (warehouses, databases, schemas, and tables).

**In this Topic:**

*Parameters*



- ❑ Snowflake provides industry-leading features that ensure the highest levels of security for your account, users, and all the data you store in Snowflake:

- ❑ Network/site access

Network policies for controlling site access, i.e. IP whitelisting

- ❑ Account/user authentication

Multi-factor authentication (MFA)

Federated authentication and SSO

- ❑ Object security

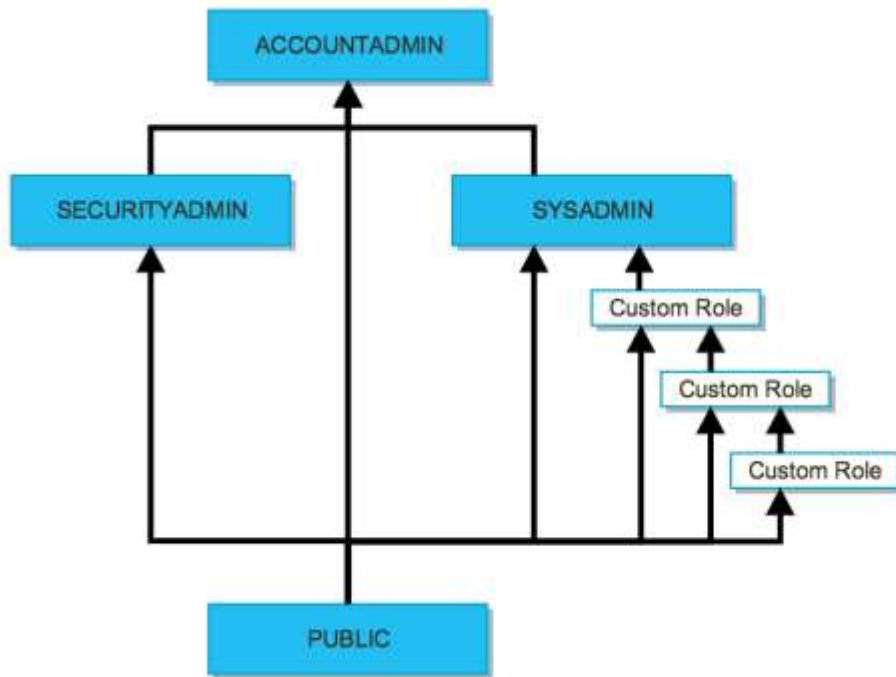
Hybrid model for controlling access to any object in the account (users, warehouses, databases, tables, etc.):

-DAC (discretionary access control)

-RBAC (role-based access control)

- ❑ Data security

Automatic data encryption using AES 256 strong encryption

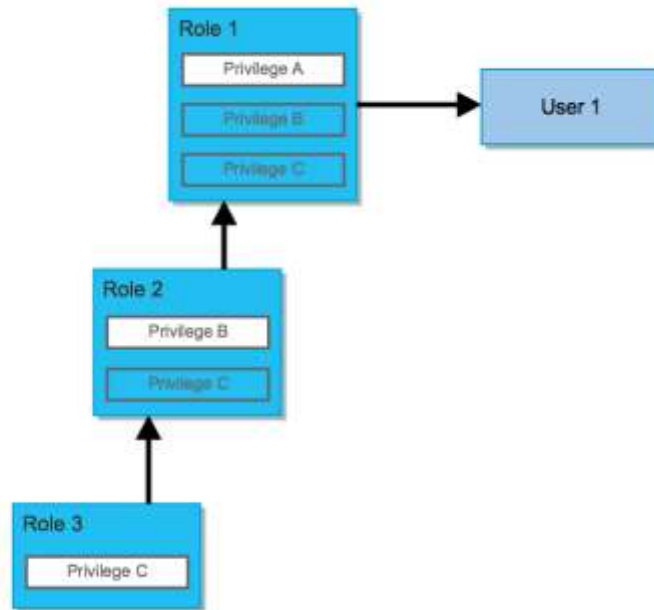


For a more specific example of role hierarchy and privilege inheritance, consider the following scenario:

- Role 3 has been granted to Role 2.
- Role 2 has been granted to Role 1.
- Role 1 has been granted to User 1.

• In this scenario:

- Role 2 inherits Privilege C
- Role 1 inherits Privileges B and C
- User 1 has all three privileges



During the course of a session, the user can use the USE ROLE command to change the current role. If the role is granted other roles, the user's session has privileges equal to the sum of the privileges of the current role and all the privileges of the roles granted to the current role within the role hierarchy.

When a user attempts to execute an action on an object, Snowflake compares the privileges available in the user's session against the privileges required on the object for that action. If the session has the required privileges on the object, the action is allowed.

The Snowflake Information Schema consists of a set of system-defined views and functions that provide metadata information about the objects in your account:

The views describe database objects, as well as non-database objects, defined for your account, including:

- Views from the Information Schema SQL standard that are relevant to Snowflake (tables, columns, views, etc).

- Views for the non-standard database objects that Snowflake supports (e.g. stages and file formats).

- Views for other, non-database objects (e.g. roles).

The functions provide historical information, including:

- Query history.

- Data storage usage for databases and stages.

- Warehouse usage.

The views in INFORMATION\_SCHEMA display metadata about all objects defined in the database, as well as Snowflake objects that are common across all databases, e.g. roles.

Some of the important views INFORMATION\_SCHEMA contains are:

- APPLICABLE\_ROLES View
- COLUMNS View
- DATABASES View
- FUNCTIONS View
- OBJECT\_PRIVILEGES View
- REFERENTIAL\_CONSTRAINTS View
- SCHEMATA View
- TABLE\_PRIVILEGES View
- TABLES View
- VIEWS View

The Snowflake Information Schema consists of a set of system-defined views and functions that provide metadata information about the objects in your account:

The views describe database objects, as well as non-database objects, defined for your account, including:

- Views from the Information Schema SQL standard that are relevant to Snowflake (tables, columns, views, etc).

- Views for the non-standard database objects that Snowflake supports (e.g. stages and file formats).

- Views for other, non-database objects (e.g. roles).

The functions provide historical information, including:

- Query history.

- Data storage usage for databases and stages.

- Warehouse usage.

- ❑ Given a batch of 15 SQL files what is the best method to insert a PAUSE between each?
- ❑ Can Snowflake support EXCLUDE list while cloning a database or schema? Some tables may not be needed in the clone.
- ❑ Snowflake document states that SEQ.CURRVAL is not supported. What is the workaround that should be suggested for entrenched Oracle users who prefer to maintain the current value across multiple SELECT statements?



## END TO END USAGE OF SNOWFLAKE

### PRE-REQ Hands-On Loading Data and Unloading

- ❑ Download data shared by trainer
- ❑ Create Multiple Warehouses For Upload And Analytics
- ❑ Create DB & table in Snowflake
- ❑ Load data to Table
- ❑ Connect to Tableau and Analyze

- ❑ The data is in csv format and you can download it in the link mentioned earlier.
- ❑ The name of the file is Crimes.csv
- ❑ The data is about the crimes that have happened across the Chicago, USA from 2016-2017.

- ❑ Go to cmd and install SNOWSQL  
*snowsql;*
- ❑ Login to your account  
*snowsql -a useready -u rajeshh;*
- ❑ Load file to staging  
*put file://C:\SnowflakeStuffs\OtherFiles\Crimes.csv @raj\_stage;*

- ❑ Go to <https://useready.snowflakecomputing.com/console/login#/> and login using username and password
- ❑ Go to Snowflake Worksheet  
Use Warehouse TRAINING\_TEST\_WAREHOUSE  
Use Database TRNG\_ANALYTICS\_DB  
Use ETL schema

□ Create a table CRIME\_SF

*create or replace table CRIME\_SF*

*(ID int not null, Case\_Number varchar(50), Date Timestamp, Block varchar(255), IUCR varchar(50), Primary\_Type varchar(255), Description varchar(255), Location\_Description varchar(255),*

*Arrest varchar(50), Domestic varchar(50), Beat int, District int, Ward int, Community\_Area int, FBI\_Code varchar(50), X\_Coordinate int, Y\_Coordinate int, Year int, Updated\_On timestamp, Latitude float, Longitude float, Location varchar(150));*

```
CREATE FILE FORMAT "TRNG_ANALYTICS_DB"."ANALYTICS".csv TYPE = 'CSV'
COMPRESSION = 'AUTO' FIELD_DELIMITER = ',' RECORD_DELIMITER = '\n'
SKIP_HEADER = 1 FIELD_OPTIONALLY_ENCLOSED_BY = '\042' TRIM_SPACE =
TRUE ERROR_ON_COLUMN_COUNT_MISMATCH = TRUE ESCAPE = 'NONE'
ESCAPE_UNENCLOSED_FIELD = '\134' DATE_FORMAT = 'AUTO'
TIMESTAMP_FORMAT = 'AUTO' NULL_IF = ('');
```

copy into CRIME\_SF

```
from @raj_stage/Crimes.csv FILE_FORMAT = (FORMAT_NAME =
"csv", compression="gzip");
```

```
copy into @raj_stage/Crimes_unload.csv
 from CRIME_SF
 FILE_FORMAT = (FORMAT_NAME = "csv", compression="gzip");
```

```
Get @raj_stage/Crimes_unload.csv file:///tmp;
```

```
list @raj_stage PATTERN='.*1.csv';
```

```
Remove @raj_stage pattern='.*.csv';
```

## JSON DATA UPLOAD AND QUERY

- ❑ Create warehouse, database, schema, file format
- ❑ Create table to store raw JSON
- ❑ Load data to Snowflake
- ❑ Load data to Snowflake tables
- ❑ Sample query



```
CREATE TABLE people_data.people_json_src
(
 src_json variant
);
```

```
create TABLE people_data.people
as
SELECT src_json:"_id"::varchar as id
 ,src_json:guid::varchar as guid ,src_json:isActive::varchar as isActive
 ,regexp_replace(src_json:balance,"\\$|\\'|,")::decimal(15,2) as balance ,src_json:picture::varchar as picture
 ,src_json:age::number as age ,src_json:eyeColor::varchar as eyeColor
 ,src_json:name.first::varchar as first_name ,src_json:name.last::varchar as last_name
 ,src_json:company::varchar as company ,src_json:email::varchar as email
 ,src_json:phone::varchar as phone ,src_json:address::varchar as address
 ,src_json:about::varchar as about ,src_json:registered::varchar as registered
 ,src_json:latitude::decimal(15,6) as latitude ,src_json:longitude::decimal(15,6) as longitude
 ,src_json:greeting::varchar as greeting ,src_json:favoriteFruit::varchar as favoriteFruit
FROM people_data.people_json_src
WHERE check_json(src_json) is null;
```

```
CREATE TABLE people_data.tags
as
SELECT
 src_json:_id::varchar as id
 ,tags.value::varchar as tag
FROM
 people_data.people_json_src p
 ,TABLE(FLATTEN(p.src_json:tags)) tags;
```

```
CREATE TABLE people_data.friends
as
SELECT
 src_json:_id::varchar as id
 ,friends.value:id::number as friend_id
 ,friends.value:name::varchar as friend_name
FROM
 people_data.people_json_src p
 ,TABLE(FLATTEN(p.src_json:friends)) friends;
```

```
select
 *
from
 people_data.people
where
 id in (select id from people_data.friends where
 friend_name='Parks Ray');
```

## ETL ACCESS

- ❑ Use of ODBC –System DSN
- ❑ Sample workflow
- ❑ Parameters

- View

SHOW PARAMETERS [ LIKE '<pattern>' ] IN ACCOUNT

- Alter

ALTER ACCOUNT SET <param> = <value>

- Reset

ALTER ACCOUNT UNSET <param>

- Snowflake provides parameters that let you control the behaviour of your account, individual user sessions, and objects.
- All the parameters have default values, which can be set and then overridden at different levels depending on the parameter type (Account, Session, or Object).

## In this Topic:

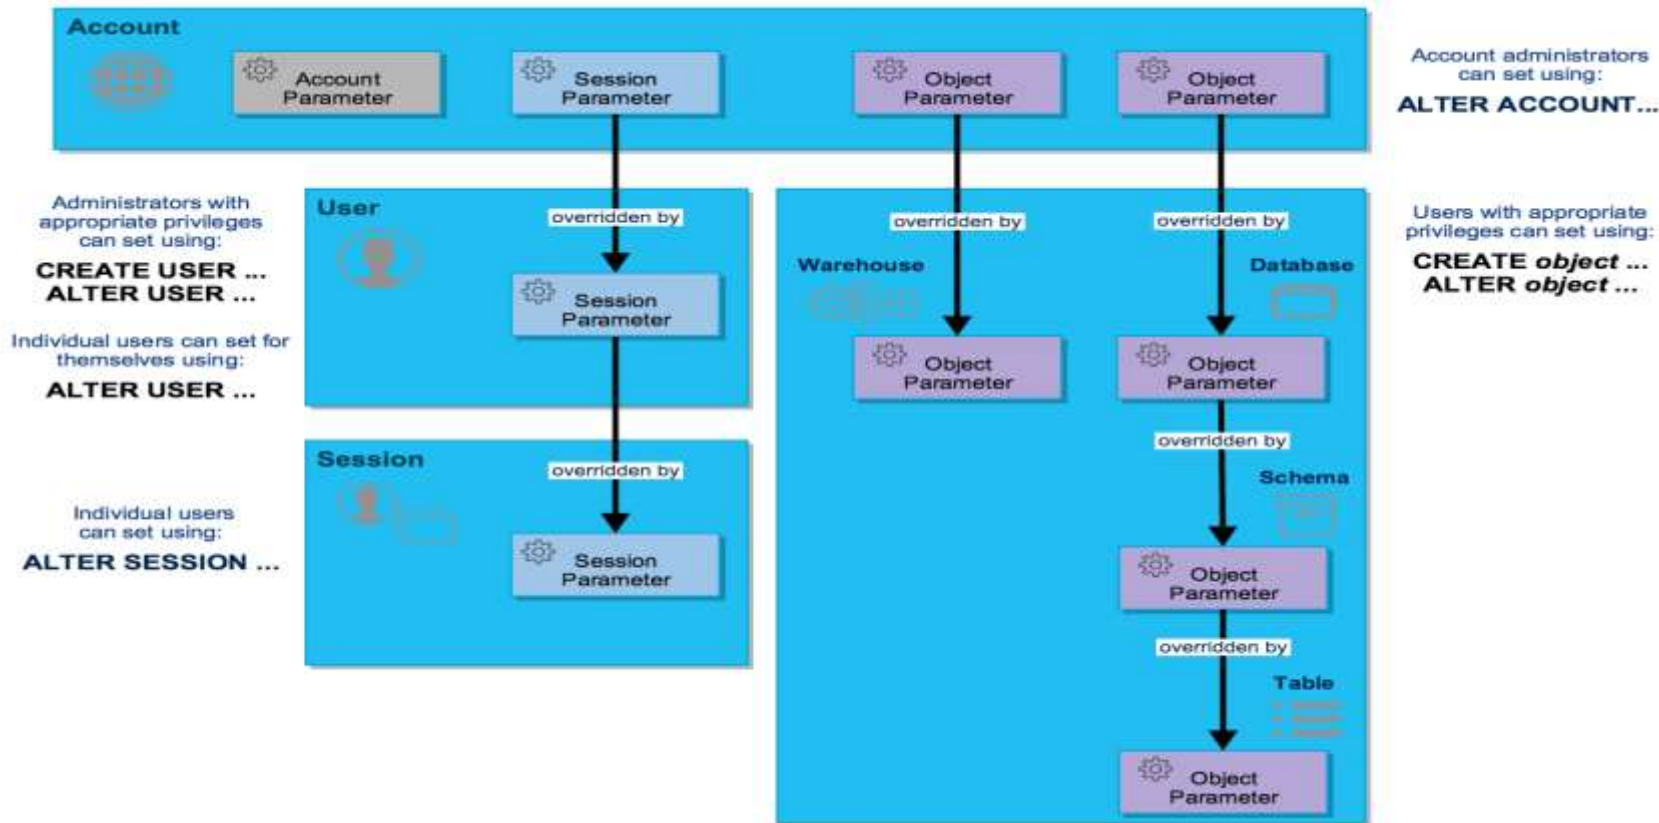
*Parameter Types and Hierarchy*

*Account Parameters*

*Session Parameters*

*Object Parameters*





- Account parameters can be set only at the account level by users with the ACCOUNTADMIN role (i.e. account administrators). Account parameters are set using the ALTER ACCOUNT command.
- Snowflake provides the following account parameters:
  - CLIENT\_ENCRYPTION\_KEY\_SIZE
  - NETWORK\_POLICY
  - PERIODIC\_DATA\_REKEYING
  - SAML\_IDENTITY\_PROVIDER
  - SSO\_LOGIN\_PAGE

- **CLIENT\_ENCRYPTION\_KEY\_SIZE:**

|                |                                                                                                                                                                                                                                                                  |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Type:          | Account — Can be set only for Account                                                                                                                                                                                                                            |
| Data Type:     | Integer                                                                                                                                                                                                                                                          |
|                | Specifies the AES encryption key size, in bits (128 or 256), used to encrypt/decrypt files stored in internal staging locations (for loading/unloading data).                                                                                                    |
| Description:   | Note that, if set to 256 (i.e. strong encryption), additional JCE policy files must be installed on each client machine from which data is loaded/unloaded. For more information about installing the required files, see Java Requirements for the JDBC Driver. |
| Default Value: | 128                                                                                                                                                                                                                                                              |

- **NETWORK\_POLICY:**

|                |                                                                                                                                                                                                                                                                                                                                                       |
|----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Type:          | Account — Can be set only for Account                                                                                                                                                                                                                                                                                                                 |
| Data Type:     | String                                                                                                                                                                                                                                                                                                                                                |
| Description:   | Supports restricting access to your account based on user IP address. References a network policy created using CREATE NETWORK POLICY. Note that NETWORK_POLICY is the only account parameter that can be set by both account administrators and security administrators. For more information about network policies, see Managing Network Policies. |
| Default Value: | None                                                                                                                                                                                                                                                                                                                                                  |

- **PERIODIC\_DATA\_REKEYING:**

Type:

Account — Can be set only for Account

Data Type:

Boolean

Description:

Enables/disables re-encryption of table data with new keys on a yearly basis to provide additional data protection:

- TRUE specifies that data is rekeyed after one year has passed since the data was last encrypted. Rekeying occurs in the background so no down-time is experienced and the affected data/table is always available.

- FALSE specifies that data is not rekeyed.

You can enable and disable rekeying at any time. Enabling/disabling rekeying does not result in gaps in your encrypted data:

- If rekeying is enabled for a period of time and then disabled, all data already tagged for rekeying is rekeyed, but no further data is rekeyed until you reenable it again.

- If rekeying is re-enabled, Snowflake automatically rekeys all data that has keys which meet the criteria (i.e. key is older than one year).

Default Value:

FALSE

- **SAML\_IDENTITY\_PROVIDER:**

Type:

Data Type:

Description:

Default Value:

Account — Can be set only for Account

JSON

Enables federated authentication. The parameter accepts a JSON object, enclosed in single quotes, with the following fields:

```
{ "certificate": "", "ssoUrl": "", "type" : "",
 "label" : "" }
```

Where:

**certificate:** Specifies the certificate (generated by the IdP) that verifies communication between the IdP and Snowflake.

**ssoUrl:** Specifies the URL endpoint (provided by the IdP) where Snowflake sends the SAML requests.

**type:** Specifies the type of IdP used for federated authentication ("OKTA", "ADFS", "Custom").

**label:** Specifies the button text for the IdP in the Snowflake login page. The default label is Single Sign On. If you change the default label, the label you specify can only contain alphanumeric characters (i.e. special characters and blank spaces are not currently supported).

Note that, if the "type" field is "Okta", a value for the label field does not need to be specified because Snowflake displays the Okta logo in the button.

None

- **SSO\_LOGIN\_PAGE**

Type:

Account — Can be set only for Account

Data Type:

Boolean

Description:

Disables preview mode for testing SSO (after enabling federated authentication) before rolling it out to users:

- If TRUE, preview mode is disabled and users will see the button for Snowflake-initiated SSO for your identity provider (as specified in SAML\_IDENTITY\_PROVIDER) in the Snowflake main login page.

- If FALSE, preview mode is enabled and SSO can be tested using the following URL:

- [https://<account\\_name>.snowflakecomputing.com/console/login?fedpreview=true](https://<account_name>.snowflakecomputing.com/console/login?fedpreview=true)

For more information, see [Configuring and Using Federated Authentication in Snowflake](#).

Default Value:

FALSE

Most parameters are session parameters, which can be set at the following levels:

**Account ► User ► Session**

### **Account**

Account administrators can use the ALTER ACCOUNT command to set session parameters for the account. The values set for the account default to individual users and their sessions.

### **User**

Administrators with the appropriate privileges can use the ALTER USER command to override session parameters for individual users. The values set for a user default to any sessions started by the user. In addition, users can override sessions parameters for themselves using ALTER USER.

### **Session**

Users can use the ALTER SESSION to override session parameters within their sessions.



Object parameters can be set at the following levels:

**Account ► Object**

**Account**

Account administrators can use the ALTER ACCOUNT command to set object parameters for the account. The values set for the account default to the objects created in the account.

**Object**

Users with the appropriate privileges can use the corresponding CREATE or ALTER commands to override object parameters for an individual object.

Currently, Snowflake provides the following object parameters:

- DATA\_RETENTION\_TIME\_IN\_DAYS (for databases, schemas, and tables)
- MAX\_CONCURRENCY\_LEVEL (for warehouses)
- STATEMENT\_QUEUED\_TIMEOUT\_IN\_SECONDS (for warehouses); session parameter also
- STATEMENT\_TIMEOUT\_IN\_SECONDS (for warehouses); session parameter also

- **DATA\_RETENTION\_TIME\_IN\_DAYS :**

|                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Type:          | Object (for databases, schemas, and tables)<br>Can be set for Account ► Database ► Schema ► Table                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| Data Type:     | Number<br><br>Number of days for which Snowflake retains historical data for performing Time Travel actions (SELECT, CLONE, UNDROP) on the object: <ul style="list-style-type: none"><li>•For Standard Edition accounts, this parameter can be set to either 1 or 0.</li><li>•For Enterprise Edition accounts, this parameter can be changed to any value from 0 up to 90 days. The value specified defaults to databases, schemas, and permanent tables, and can be overridden for each.</li></ul> A value of 0 for a database, schema, or table effectively disables Time Travel for the specified object.<br>For more information, see Understanding & Using Time Travel. |
| Description:   |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| Default Value: | 1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |

- **MAX\_CONCURRENCY\_LEVEL**

Type: Object (for warehouses) — Can be set for Account ► Warehouse

Data Type: Number

Description: Specifies the maximum number of SQL statements (queries, DDL, DML, etc.) a warehouse cluster can execute concurrently. When the max level is reached:

- For a single-cluster warehouse or a multi-cluster warehouse (in Maximized mode), additional statements are queued until resources are available.
- For a multi-cluster warehouse (in Auto-scale mode), additional clusters are started.
- With large, complex statements, fewer statements than the max level may execute concurrently. As each statement is submitted to a warehouse, Snowflake ensures that resources are not over-allocated; if there aren't enough resources available to execute a statement, it is automatically queued.

Default Value: 8

## STATEMENT\_QUEUED\_TIMEOUT\_IN\_SECONDS:


|                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Type:          | Session <b>and</b> Object (for warehouses)<br>Can be set for for Account ► User ► Session; can also be set for individual warehouses                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Data Type:     | Number                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Description:   | Maximum amount of time, in seconds, a SQL statement (query, DDL, DML, etc.) can be queued on a warehouse before it is canceled by the system. A value of 0 specifies that no timeout is enforced. This parameter can be used in conjunction with the MAX_CONCURRENCY_LEVEL parameter to ensure a warehouse is never backlogged. The parameter can be set within the session hierarchy. It can also be set for a warehouse to control the queue timeout for all SQL statements processed by the warehouse. When the parameter is set for both a warehouse and a session, the lowest non-zero value is enforced. |
| Default Value: | 0 (i.e. no timeout)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |

- **STATEMENT\_TIMEOUT\_IN\_SECONDS:**

|                |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Type:          | Session <b>and</b> Object (for warehouses)<br>Can be set for for Account ► User ► Session; can also be set for individual warehouses                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| Data Type:     | Number                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Description:   | <p>Time, in seconds, after which a running SQL statement (query, DDL, DML, etc.) is cancelled by the system. A value of 0 specifies that no timeout is enforced. The parameter can be set within the session hierarchy. It can also be set for a warehouse to control the runtime for all SQL statements processed by the warehouse. When the parameter is set for both a warehouse and a session, the lowest non-zero value is enforced.</p> <p>For example, a warehouse has a timeout of 1000 seconds and the current user session has a timeout of 500 seconds. Any statement that executes for longer than 500 seconds in the session is cancelled.</p> |
| Default Value: | 0 (i.e. no timeout)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |

- ❑ Feedback URL:
- ❑ <https://docs.google.com/forms/d/e/1FAIpQLSeJLPDihApnTep3XTvtmGNg4Bvcg4iCpq2Eja8aPp3rcV4YQw/viewform>
- ❑ Required to improve training quality

- ❑ What are the total no of cases?
- ❑ What are the total no of crimes?
- ❑ Percentage of arrested and not arrested
- ❑ Where have most crimes happened?
- ❑ What type of crime has been committed the most?
- ❑ At what time of the day do most of the crimes happen?



The image is a collage of four screenshots from the Microsoft Power BI Desktop application, illustrating the steps to connect to a Snowflake data source.

- Top Left:** The 'Get Data' window is open. In the search bar, 'snow' is entered, and 'All' is selected. The 'Snowflake' data source is listed in the results.
- Top Right:** The 'Snowflake' connection dialog box is shown. The 'Server' field contains 'userready.snowflakecomputing.com' and the 'Warehouse' field contains 'SF\_TRAINING\_VWH'. The 'Data Connectivity mode' is set to 'Import'.
- Bottom Left:** The 'Snowflake' login dialog box is shown. The 'User name' field contains 'rajeshh' and the 'Password' field is masked with dots. There are 'Back', 'Connect', and 'Cancel' buttons.
- Bottom Right:** The 'Navigator' window is open, showing a list of available data sources. The source 'userready.snowflakecomputing.com:SF\_TRAINING\_V...' is selected, and its contents are listed below.



## Connect

### To a File

Excel  
Text file  
Access  
JSON file  
Spatial file  
Statistical file  
More...

### To a Server

Tableau Server  
Amazon Redshift  
Google Analytics  
Google Sheets  
Microsoft SQL Server  
More...

|                          |                             |                                  |
|--------------------------|-----------------------------|----------------------------------|
| Tableau Server           | IBM BigInsights             | Progress OpenEdge                |
| Action Matrix            | IBM DB2                     | QuickBooks Online                |
| Action Vector            | IBM PDA (Netezza)           | Salesforce                       |
| Amazon Aurora            | Kognitio                    | SAP HANA                         |
| Amazon EMR               | MapR Hadoop Hive            | SAP NetWeaver Business Warehouse |
| Amazon Redshift          | Marketo                     | SAP Sybase ASE                   |
| Anaplan                  | MarkLogic                   | SAP Sybase IQ                    |
| Apache Drill             | MemSQL                      | SharePoint Lists                 |
| Aster Database           | Microsoft Analysis Services | <b>Snowflake</b>                 |
| Cisco Information Server | Microsoft PowerPivot        | Spark SQL                        |
| Cloudera Hadoop          | Microsoft SQL Server        | Splunk                           |
| DataStax Enterprise      | MonetDB                     | Teradata                         |
| EXASQL                   | MySQL                       | Teradata OLAP Connect            |
| Firebird                 | OData                       | Web Data Connector               |
| Google Analytics         | Oracle                      | Other Databases (ODBC)           |
| Google BigQuery          | Oracle Eloqua               |                                  |
| Google Cloud SQL         | Oracle Essbase              |                                  |

## Snowflake

Server:

Role:

Enter information to sign in to the server:

Username:

Password:

Initial SQL...

Sign In



USERREADY

[www.userready.com](http://www.userready.com)  
[contact@userready.com](mailto:contact@userready.com)

+1-212-370-0000

805 Third Avenue  
11<sup>th</sup> floor, New York, NY 10022



[www.snowflake.net](http://www.snowflake.net)

+1-844-766-9355

101 South Ellsworth Avenue,  
San Mateo, CA, United States