



Tera-Tom Coffing

Tera-Tom's 1000 page e-Book on Teradata Architecture and SQL

Imagine a book that is in color with over 1000 slides explaining every aspect of Teradata in easy to understand terms. Then imagine every single SQL command at your fingertips filled with tips and tricks for performance tuning. I think this book will be the only book you ever need on Teradata.

Here are three sample chapters on:

Teradata Space

OLAP Ordered Analytics

Date Functions

Contact Leslie.Nolander@CoffingDW.com **to purchase.**

Pricing

You can purchase a single-license of the electronic version for yourself or a corporate license for everyone in your organization.

An individual license is \$399 and a corporate license is \$5,000.00

Order before August 1st and get a 50% discount!

Contact Leslie.Nolander@CoffingDW.com **to purchase.**

To see the entire Table of Contents use this link:

<http://www.coffingdw.com/TbasicsV12/1000PageTOC.docx>

Teradata*, NCR, BYNET and SQL Assistant, are registered trademarks of Teradata Corporation, Dayton, Ohio, U.S.A., IBM* and DB2* are registered trademarks of IBM Corporation, ANSI* is a registered trademark of the American National Standards Institute. Microsoft Windows, Windows 2003 Server, .NET are trademarks of Microsoft. Ethernet is a trademark of Xerox. UNIX is a trademark of The Open Group. Linux is a trademark of Linus Torvalds. Java is a trademark of Oracle.

Coffing Data Warehousing shall have neither liability nor responsibility to any person or entity with respect to any loss or damages arising from the information contained in this book or from the use of programs or program segments that are included. The manual is not a publication of Teradata Corporation, nor was it produced in conjunction with Teradata Corporation.

Copyright © January 2012 by Coffing Publishing



All rights reserved. No part of this book shall be reproduced, stored in a retrieval system, or transmitted by any means, electronic, mechanical, photocopying, recording, or otherwise, without written permission from the publisher. No patent liability is assumed with respect to the use of information contained herein. Although every precaution has been taken in the preparation of this book, the publisher and author assume no responsibility for errors or omissions, neither is any liability assumed for damages resulting from the use of information contained herein.

Space

“For the wise man looks into space and he knows there is no limited dimensions.”

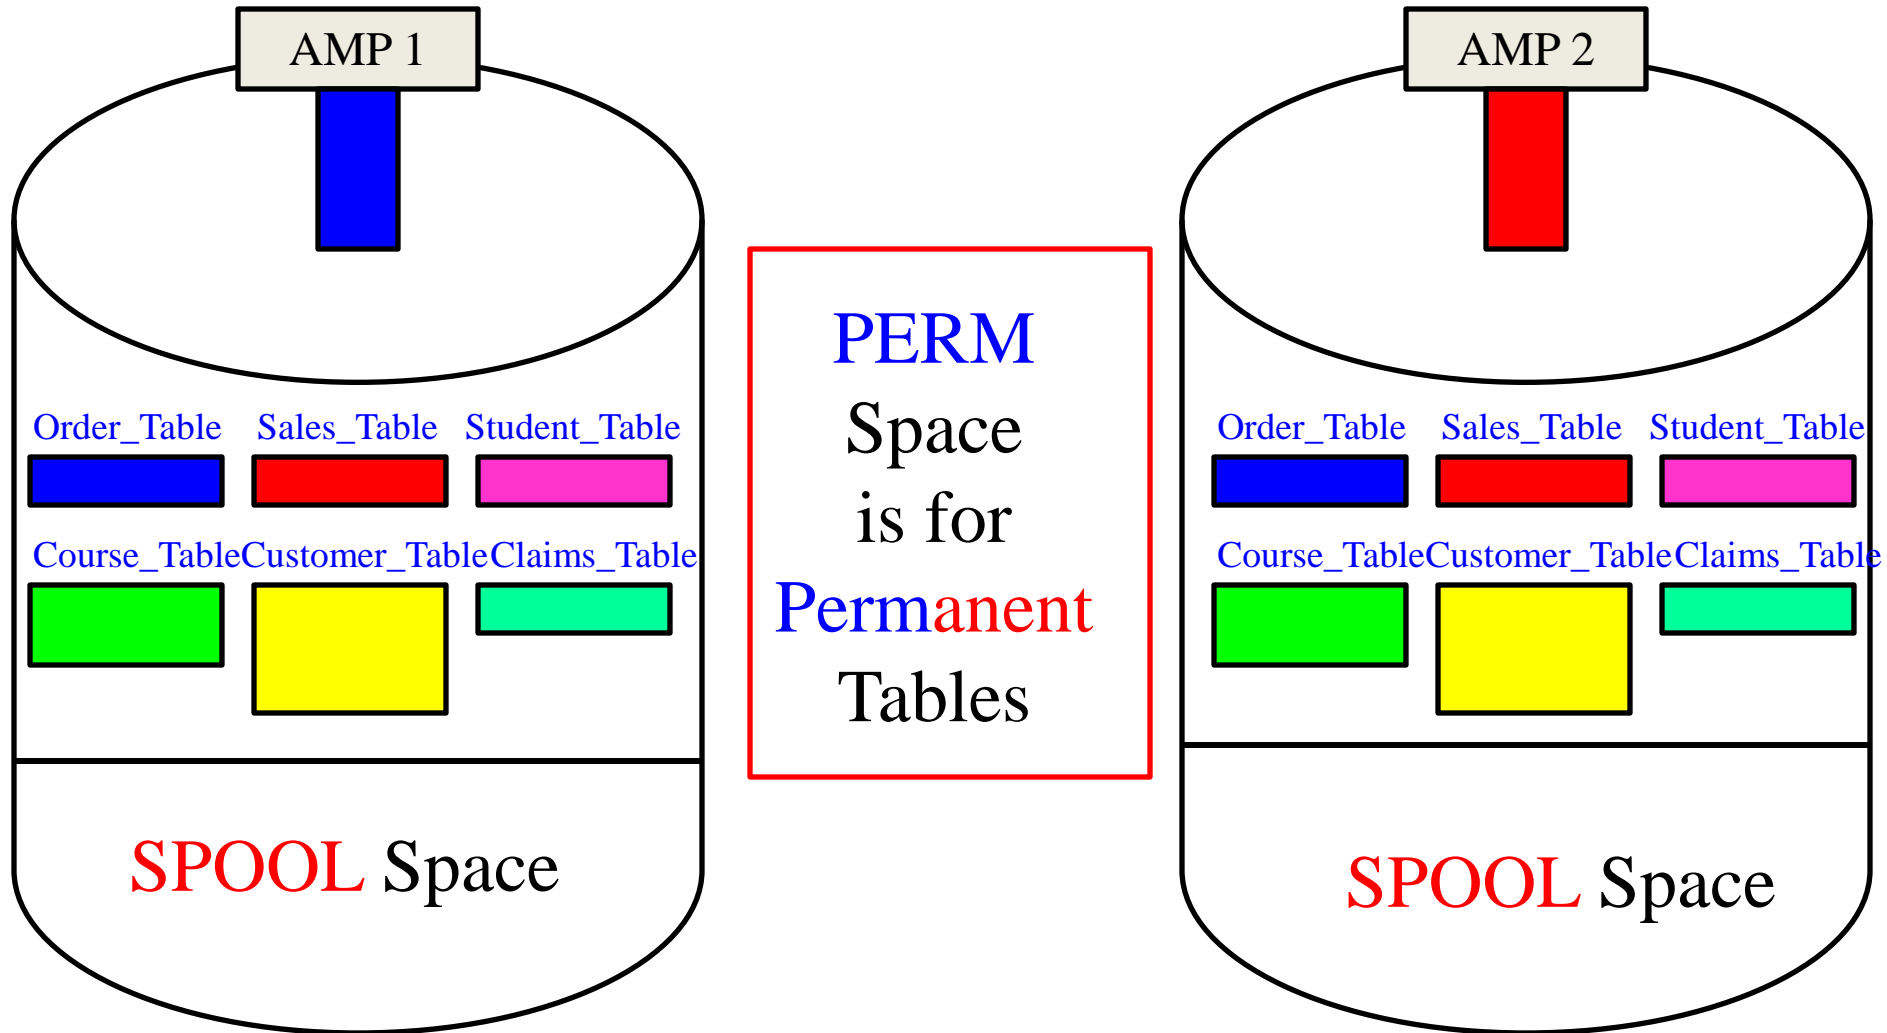
- Lao-tzu

Perm and Spool Space



Space has only to do with space on the data warehouse disks. Each AMP controls their own disk farm and about 60% of each disk will be used for tables and that is called PERM space. The other 40% (Spool) is work space for user queries and answer sets.

Perm Space is for Permanent Tables



PERM Space is where an AMP keeps its tables. That is what you need to understand. You will also find out that PERM spaces also houses Secondary Indexes, Join Indexes and Permanent Journals. Just remember that PERM is for the tables and indexes!

Spool Space is work space that builds a User's Answer Sets



AMP 1

Order_Table Sales_Table Student_Table
Course_Table Customer_Table Claims_Table

Dept_No	AVG(Salary)	Sum(Salary)
200	44944.44	89888.88
400	48316.67	144950.00

AMP 2

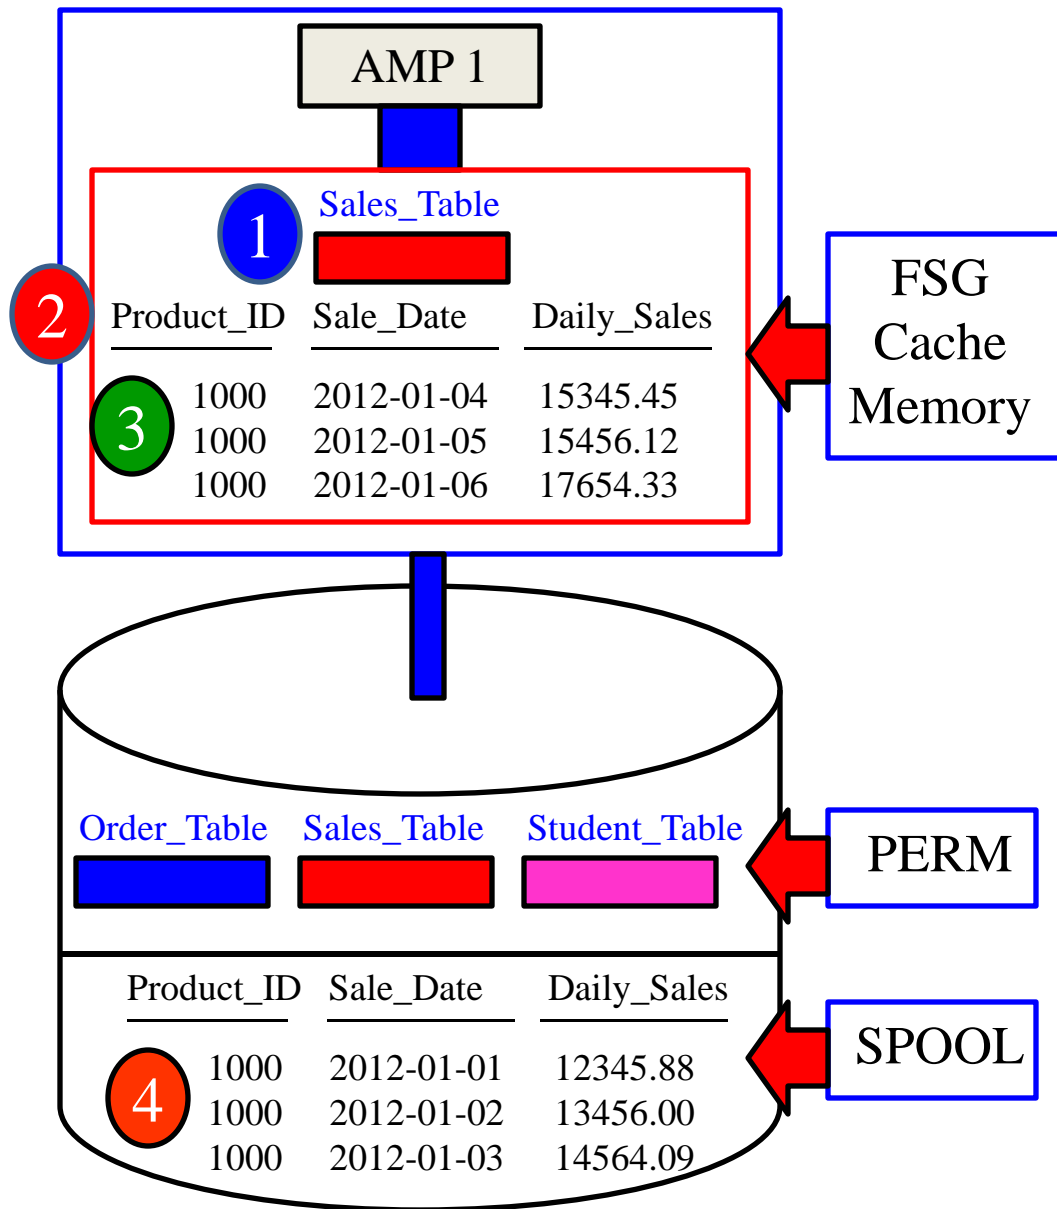
Order_Table Sales_Table Student_Table
Course_Table Customer_Table Claims_Table

Dept_No	AVG(Salary)	Sum(Salary)
100	23966.42	56868.78
300	51354.55	143450.03

Spool is
used by the
AMPS as
workspace
to build a
User's
Answer
Sets

Spool space is used by each AMP in order to build the answer set for the user.

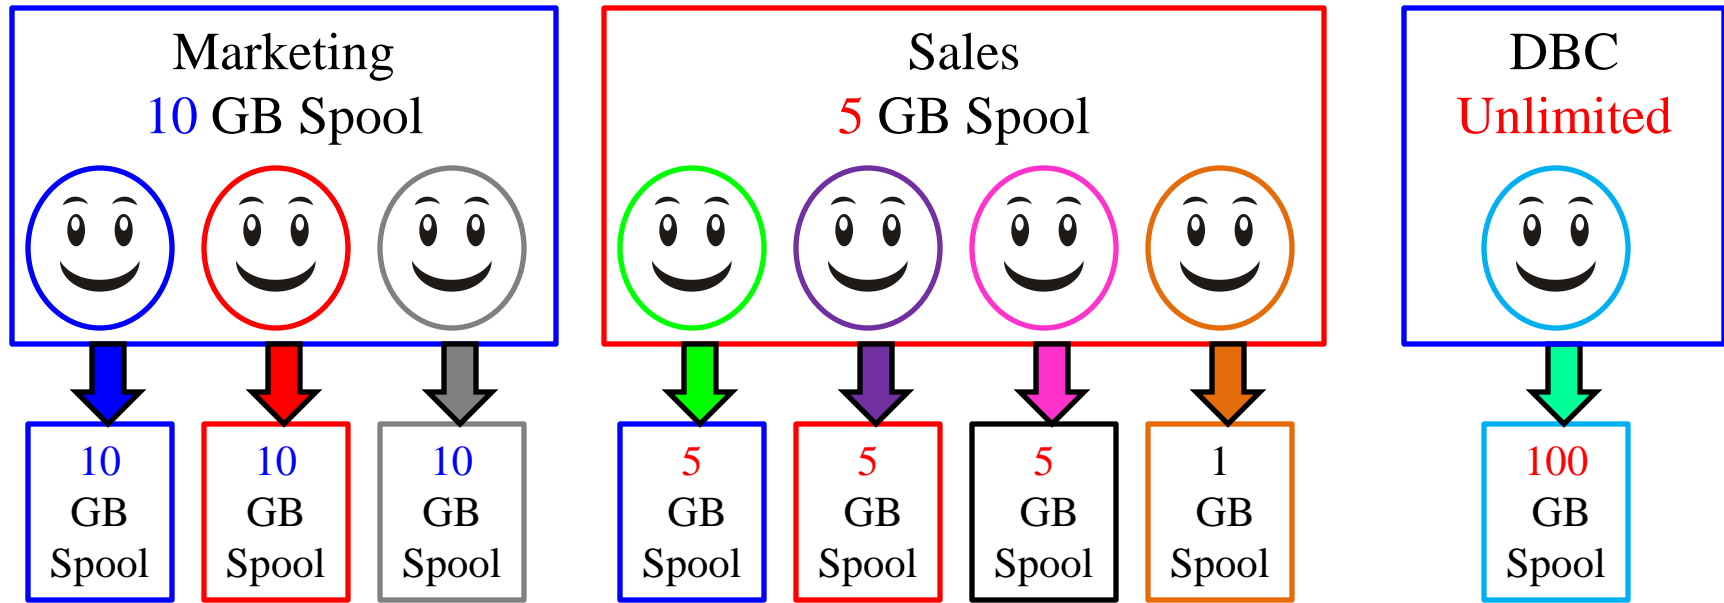
Spool Space is in an AMPs memory and on its Disk



- 1 Transfer the Sales_Table from the disk (Perm) to FSG Memory.
- 2 Get the Product_ID, Sale_Date and the Daily_Sales columns.
- 3 Build the Report in FSG Cache.
- 4 If there is no more room in FSG Cache than transfer the report to Spool on Disk.
- 5 Keep checking if the USER has gone over their SPOOL Limit.
- 6 The Report is done so transfer the report to the Parsing Engine over the BYNET.
- 7 DELETE the Spool Files.

AMPs have memory called File System Generating Cache (FSG) used for processing.

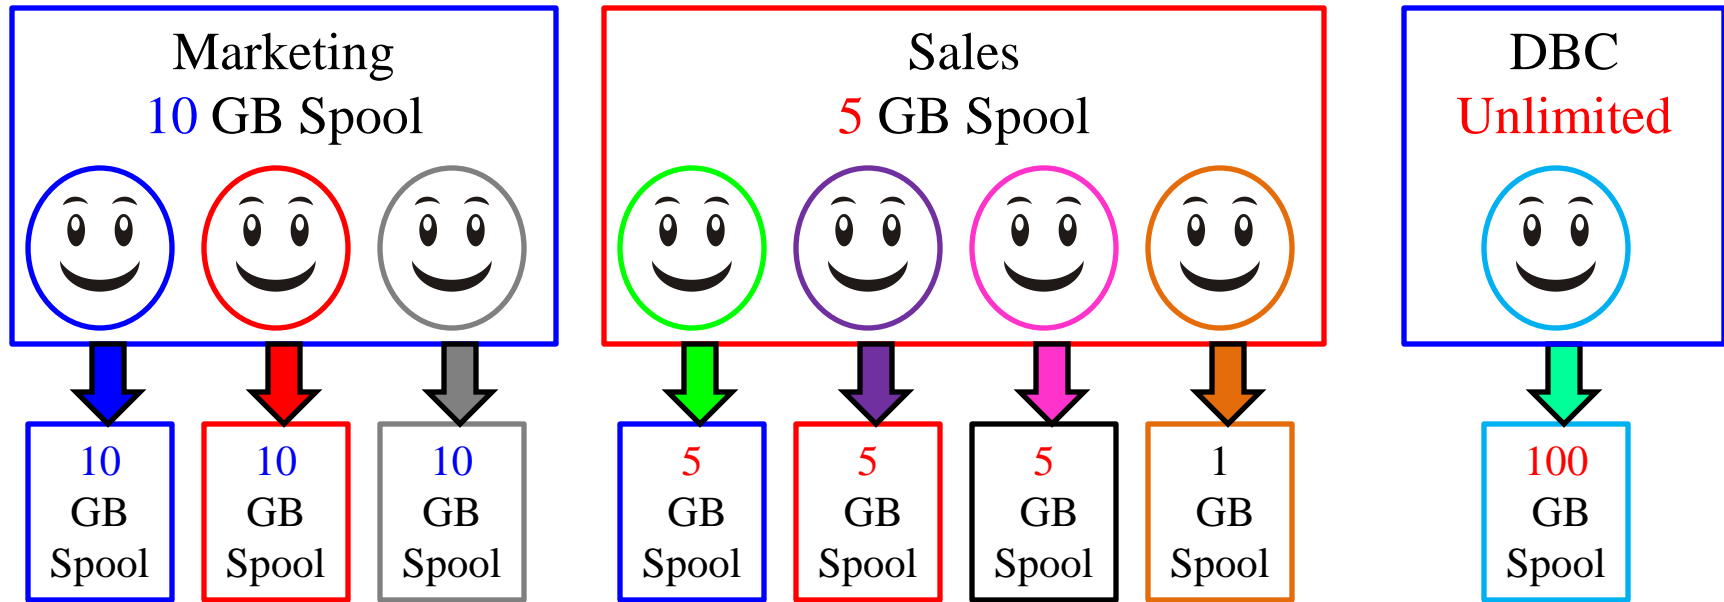
USERS are Assigned Spool Space Limits



Every User is assigned **Spool** Space so they can submit SQL and retrieve answer sets. The **Spool** in the database Marketing is 10 GB so each user **defaults** to 10 GB of Spool. Any User in Marketing can run queries, but are aborted if they go over the 10 GB limit. All 3 users in Marketing can query simultaneously and use 30 GB of Spool in total . Three users in Sales defaulted to the max (5 GB) but the **intern** was assigned **less**. The final user in DBC was given 100 GB of Spool because their brilliant.

Spool is assigned to users and the only way you are aborted is if you go over your spool limit. **Marketing** has **unlimited spool**, but the max for each user in marketing is 10 GB!

What is the Purpose of Spool Limits?




There are two reasons for Spool Limits:

If a user makes a mistake and runs a query that could take weeks to run it will abort the second the user goes over their allotted spool limit.

It keeps users from hogging the system.

Spool is assigned to users and the only way a user is aborted is if they go over their spool limit. **Marketing, Sales, and DBC** have **unlimited spool**, but the max for each individual user is 10 GB in Marketing, 5 GB in Sales, and our power user is at 100 GB.

Why did my query Abort and say “Out of Spool”?



How is it possible that I ran out of spool?


You ran out of spool because your query used over **your limit** of **10 GB** of spool.

It is also possible that you have logged onto multiple machines or ran multiple queries and the **combination** went over 10 GB of spool.


It is also very likely that the data you were working with was **NOT** evenly distributed (**skewed**) and this is a major cause of Spool errors.

Spool is assigned to users and the only way a user is aborted is if they go over their spool limit. No user has ever failed because they are in Marketing and Marketing has only 10 GB of spool. It doesn't work that way. Thousands of users in Marketing could run queries simultaneously because Marketing has unlimited amounts of spool, but each user in Marketing can't go over the default Max of 10 GBs for an individual user.

How can Skewed Data cause me to run “Out of Spool”?



Each User's Spool limit is actually done per AMP so if you are assigned 10 GBs of spool and the system has 10 AMPs you are really assigned 1 GB of Spool per AMP!



If data is skewed and you exceed your 1 GB limit on any AMP you are “out of spool”.

Spool is assigned to every user, but since Teradata is a parallel processing system each AMP is only concerned with themselves. Each AMP processes their portion of the data in parallel. Because of this philosophy your Spool Space (10 GB) is divided among the total AMPs in the system. If you have 10 GBs of Spool and there are 10 AMPs you get 1 GB per AMP. If you go over 1 GB on any AMP you are aborted and “Out of Spool”.

How come my Join caused me to run “Out of Spool”?

1

You might not have put in a Join Condition.

```
SELECT First_Name, Last_Name, Department_Name  
FROM   Employee_Table  as E  
       INNER JOIN  
       Department_Table as D
```


2

You might have Aliased the table and then fully qualified with the real table name.

```
SELECT First_Name, Last_Name, Department_Name  
FROM   Employee_Table  as E  
       INNER JOIN  
       Department_Table as D  
ON     Employee_Table.Dept_No = D.Dept_No ;
```

3

There might be skewed data on one of the tables.

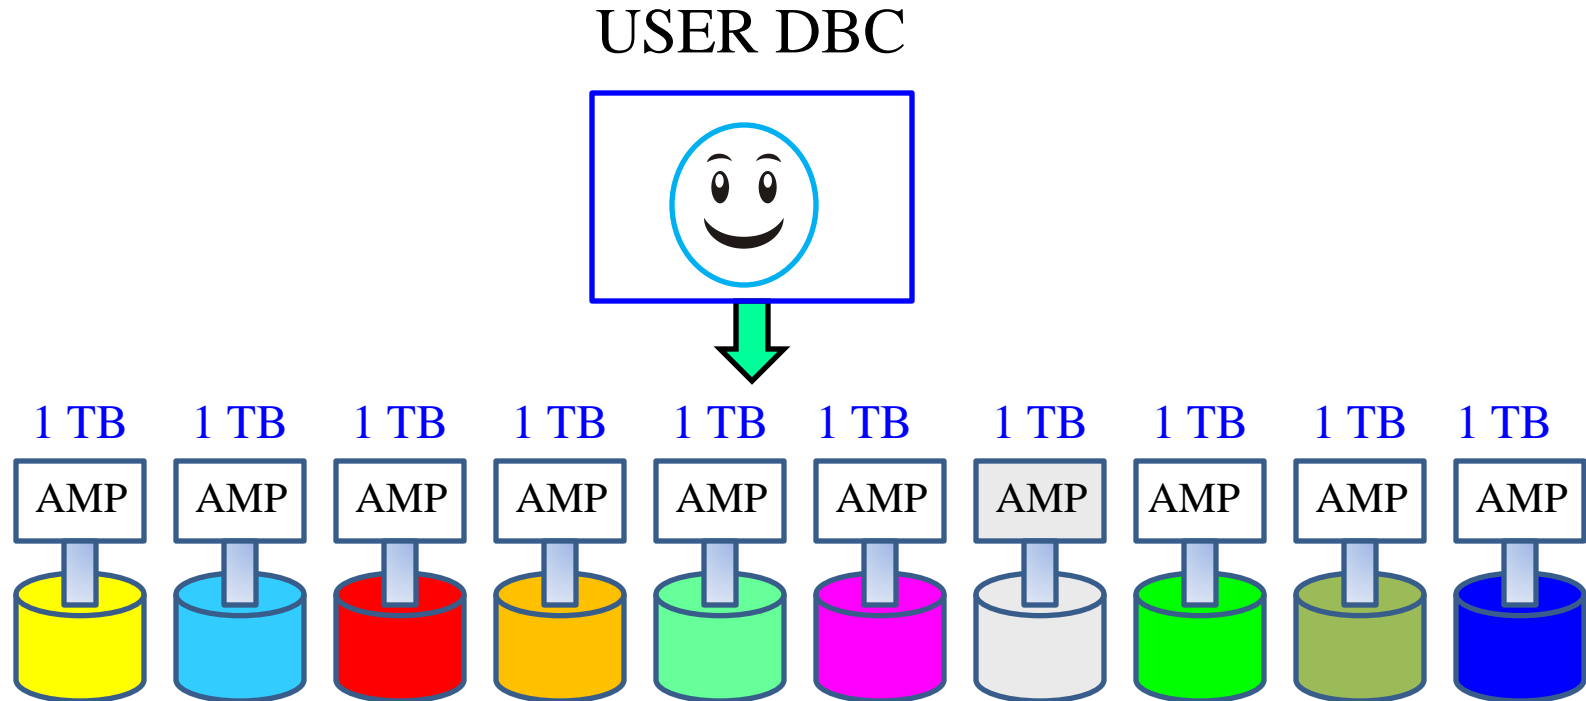


4

A Lot of NULLs on a table on an Outer Join.

```
SELECT e.*, d.* from Employee_Table as E  
LEFT OUTER JOIN Department_Table as D  
ON E.Dept_No= D.Dept_No ;
```

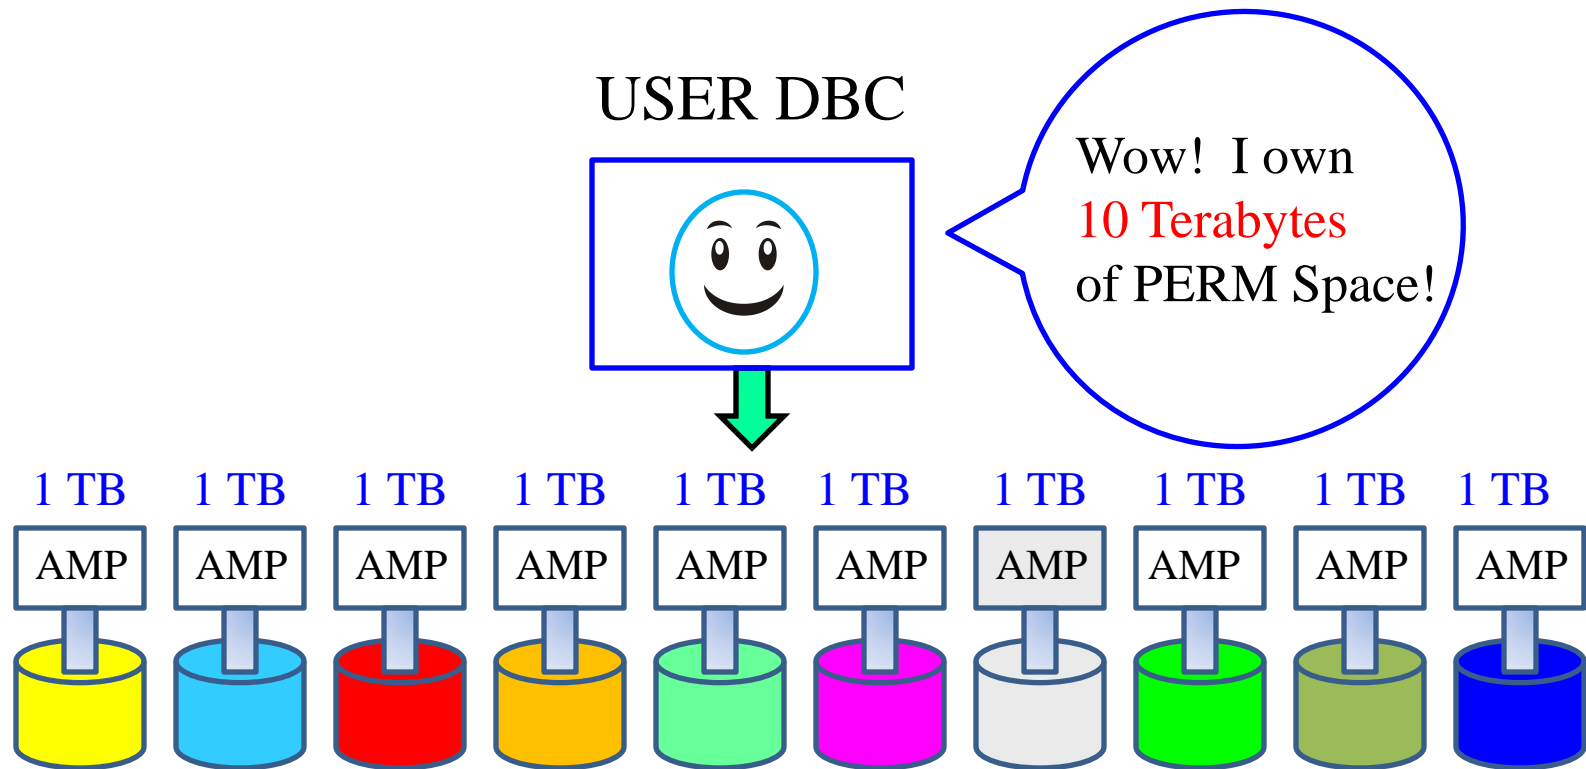
What does my system look like when it first arrives?



All Teradata systems start with one USER called **DBC**.

The first Teradata machine ever built came out in 1988 and it was called the DBC 1012. The DBC portion stood for Database Computer. The 1012 was named because 10 to the 12th power is equal to a Terabyte. So, the DBC 1012 was a Database Computer designed to process Terabytes of data. So, every system starts with one USER called DBC and DBC owns all the PERM Space in the system.

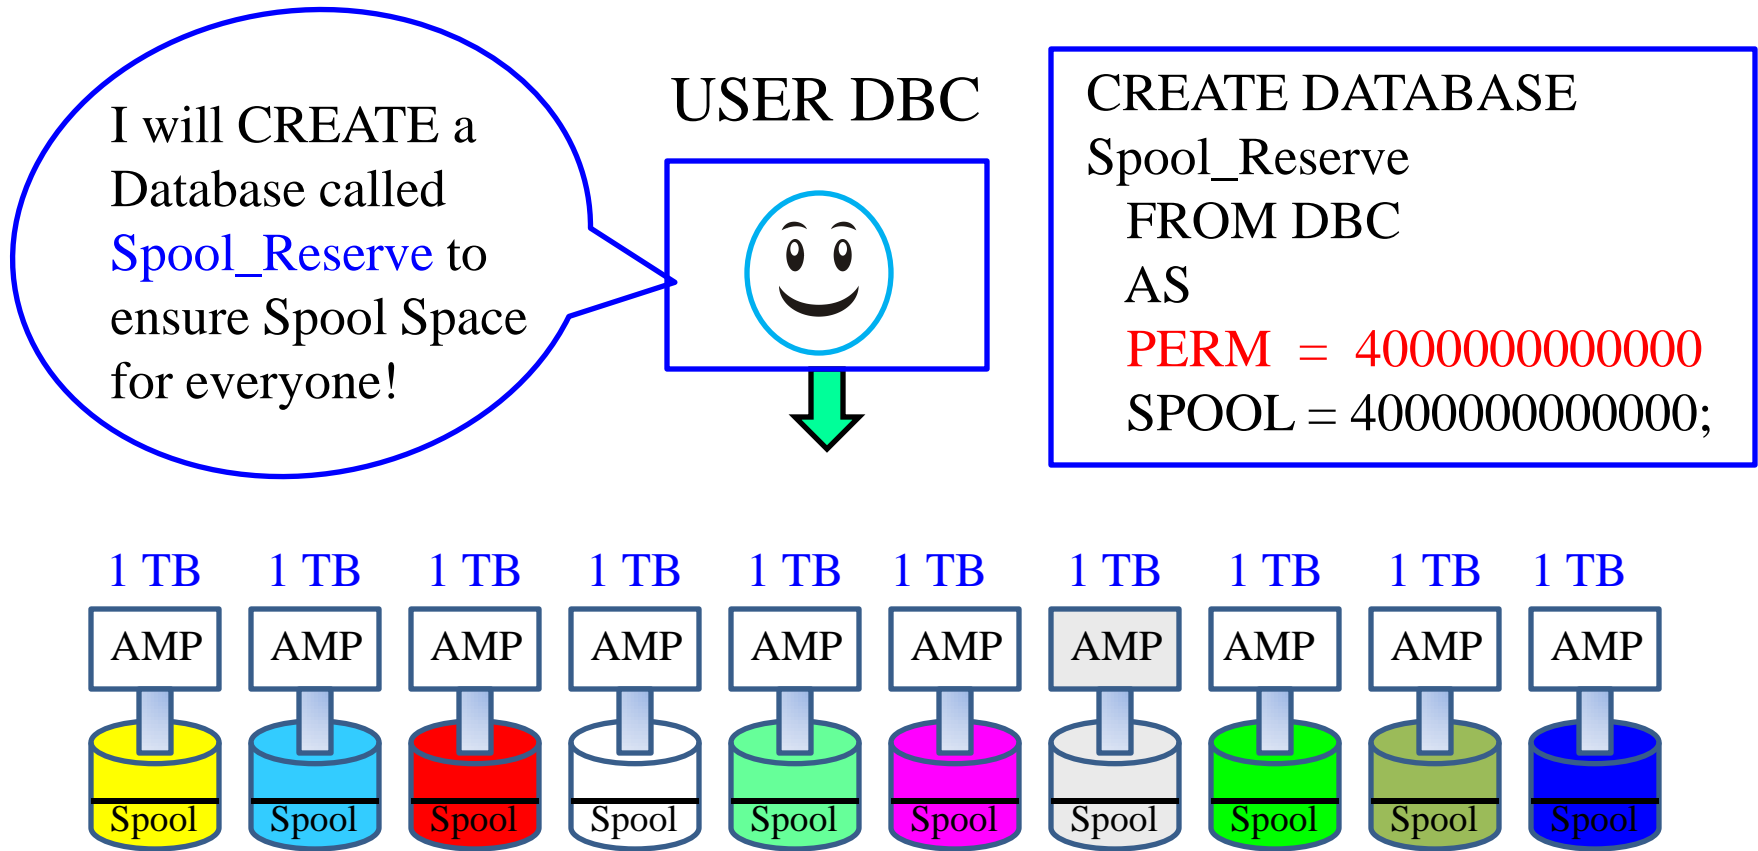
DBC owns all the PERM Space in the system on day one



When the system starts out new and arrives at your company **DBC** is the only **USER**.
DBC counts up all the **disk space** attached to each **AMP** and considers that **PERM** Space owned by DBC.

DBC owns all the disk space on day one of your systems arrival. DBC will then begin to put out space to other databases or users.

DBC's First Assignment is Spool Space



DBC will create a database called **Spool_Reserve** (any name will do), but it will reserve between 20% to 40% for Spool. What really happens is that DBC creates **Spool_Reserve** to claim **PERM** Space, but never places a table in the database.

When a database is given **PERM** Space and no object is created in that database it is used for **Spool**. **Spool** is unused **PERM**!

DBC's 2nd Assignment is to CREATE Users and Databases

I will now
CREATE USERs
and/or
DATABASEs

USER DBC

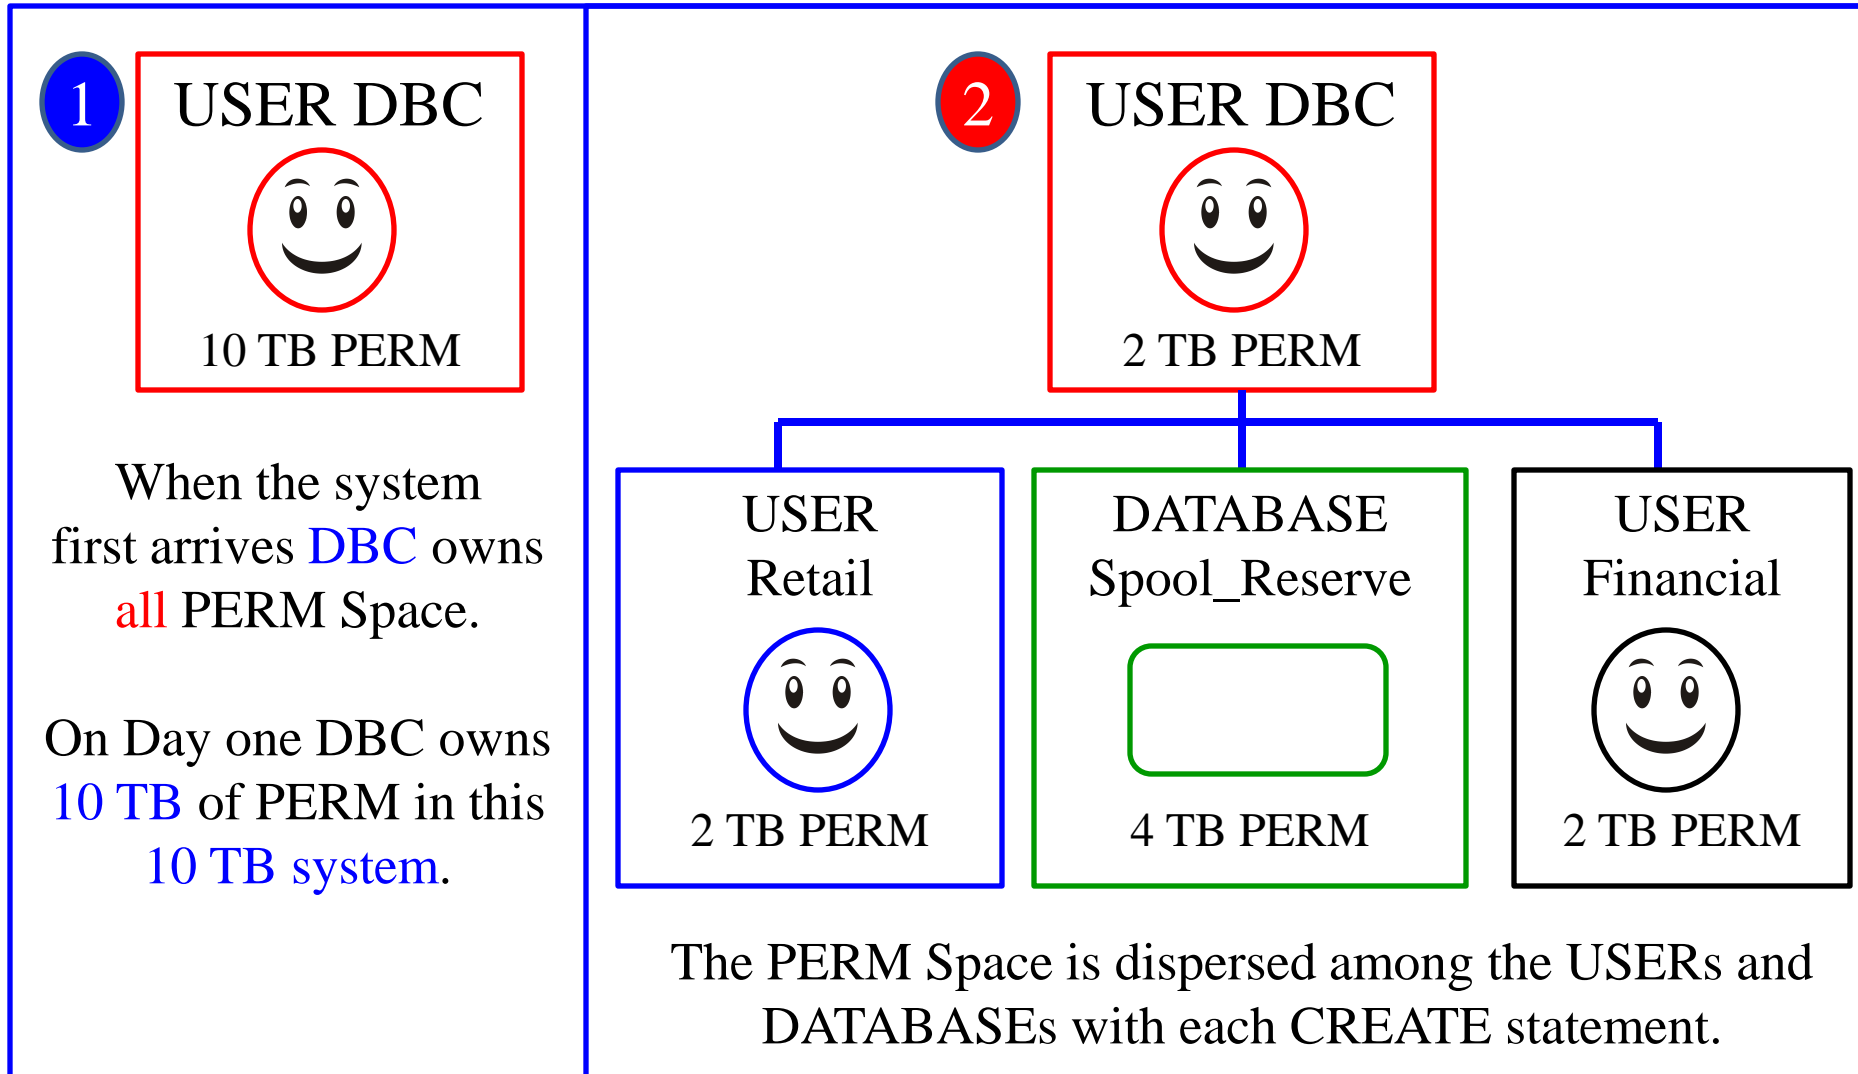


```
CREATE USER Retail
FROM DBC
AS
PASSWORD=abc123
PERM=2000000000000
SPOOL=10000000000
TEMPORARY = 10000000000
ACCOUNT='$Med'
DEFAULT DATABASE = DBC ;
```

```
CREATE USER Financial
FROM DBC
AS
PASSWORD=abc123
PERM=2000000000000
SPOOL=50000000000
TEMPORARY = 10000000000
ACCOUNT='$Med'
DEFAULT DATABASE = DBC ;
```

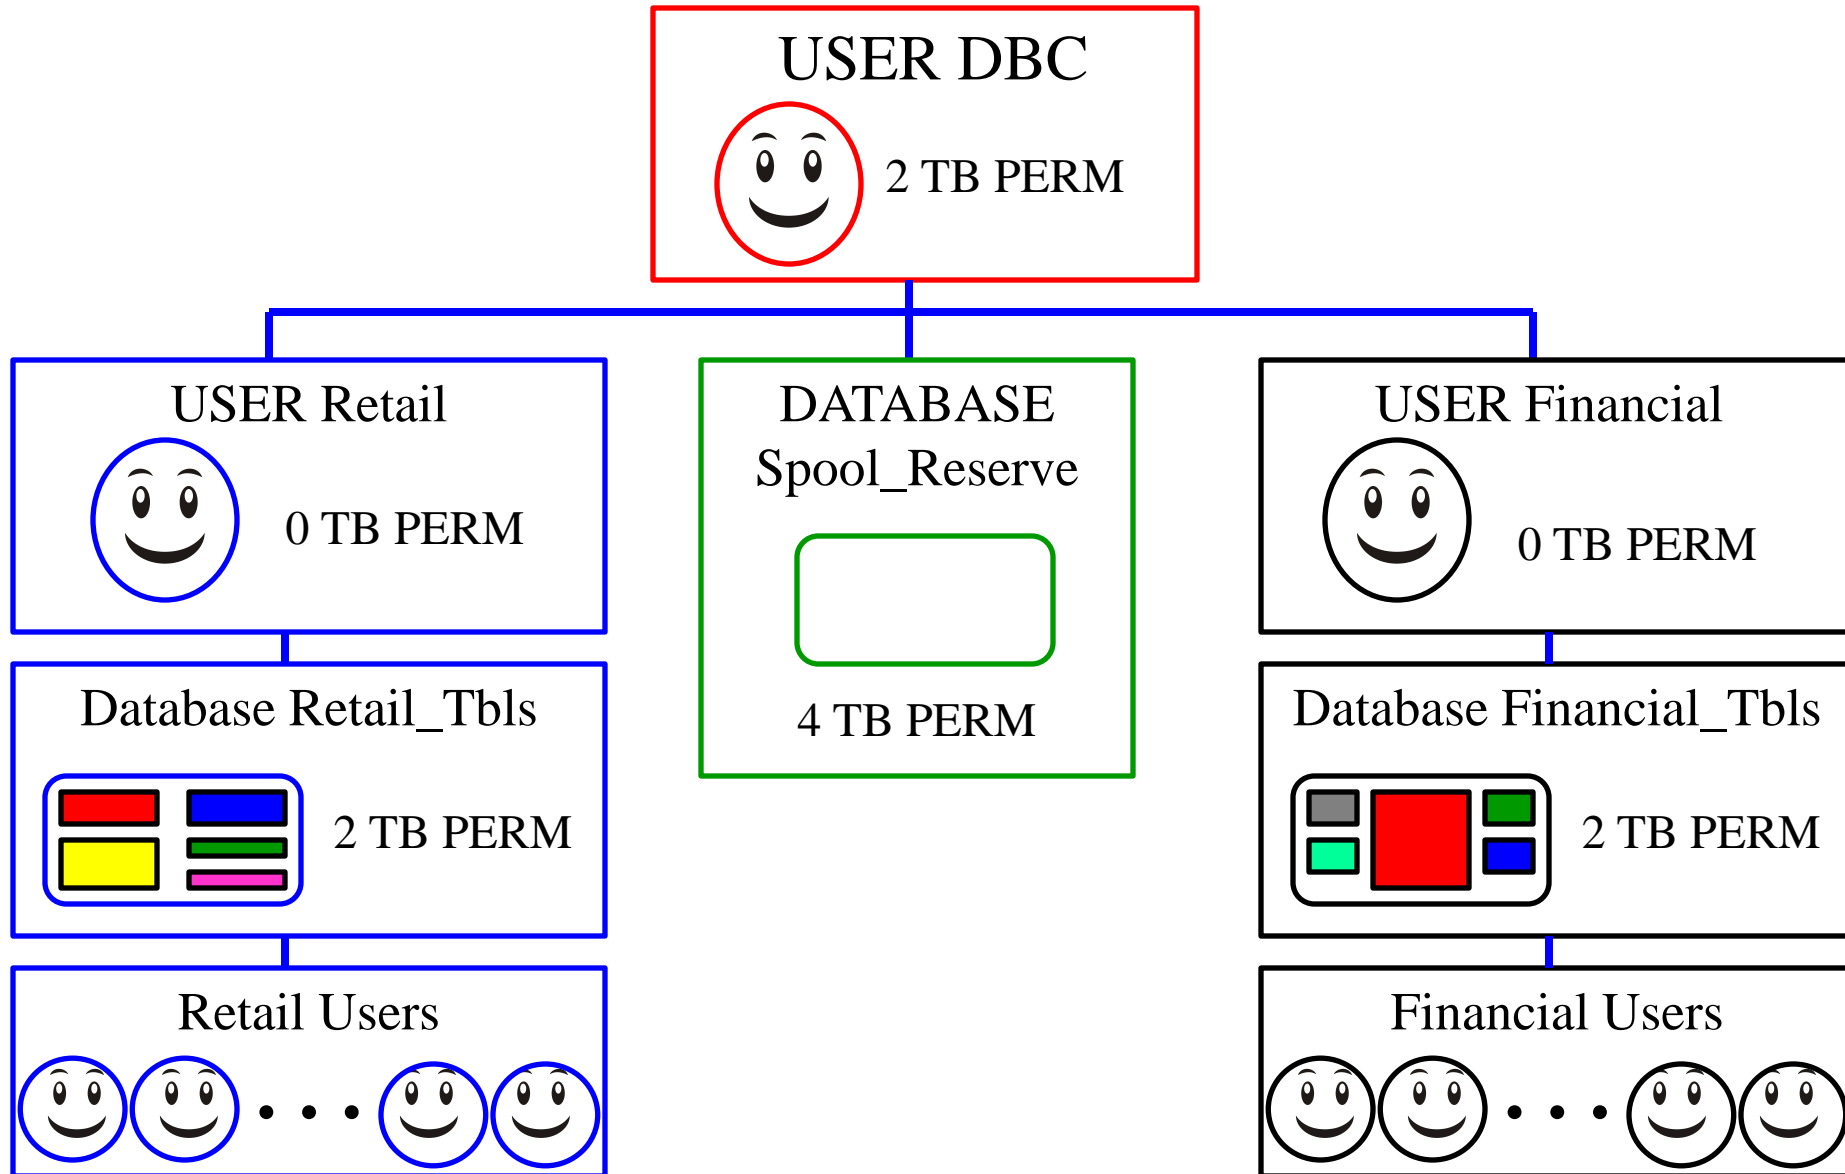
DBC's 2nd assignment will be to create some USERs or DATABASEs and the hierarchy begins. If a USER or DATABASE is assigned PERM space it can CREATE tables.

The Teradata Hierarchy Begins



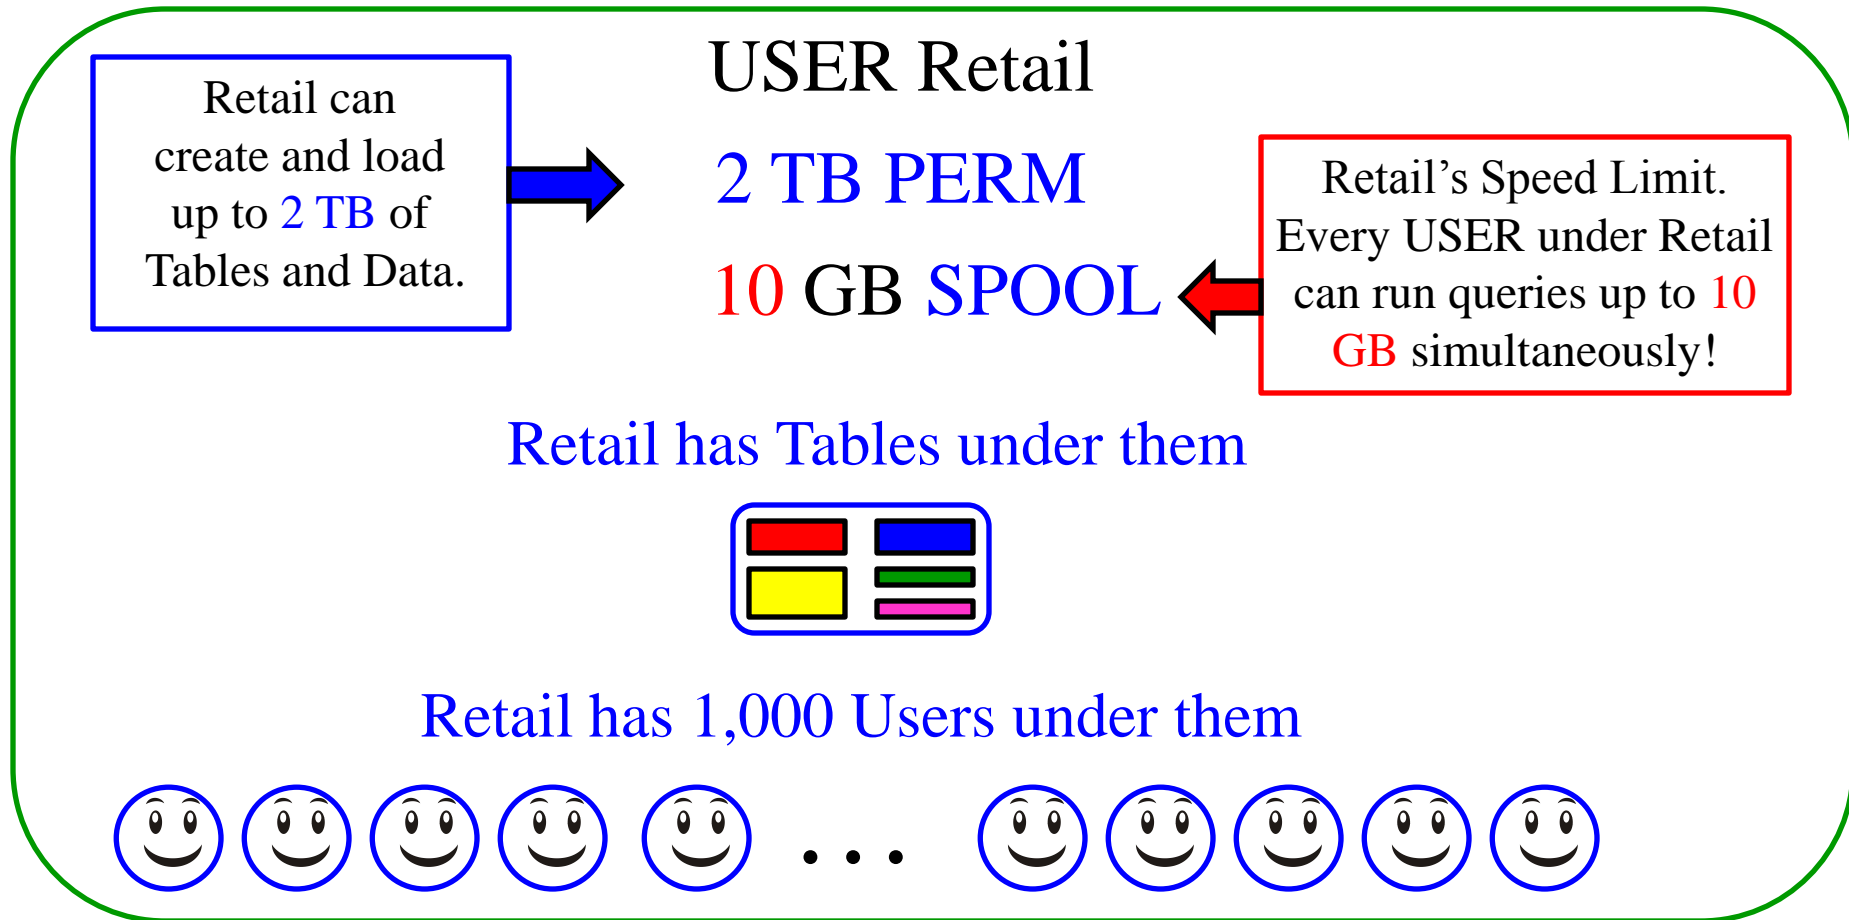
Notice in example 1 that DBC owns 10 TB of PERM space. Notice that after DBC created Spool_Reserve (4 TB), USER Retail (2 TB) and USER Financial (2 TB) that DBC now only owns only 2 TB of PERM space.

The Teradata Hierarchy Continues



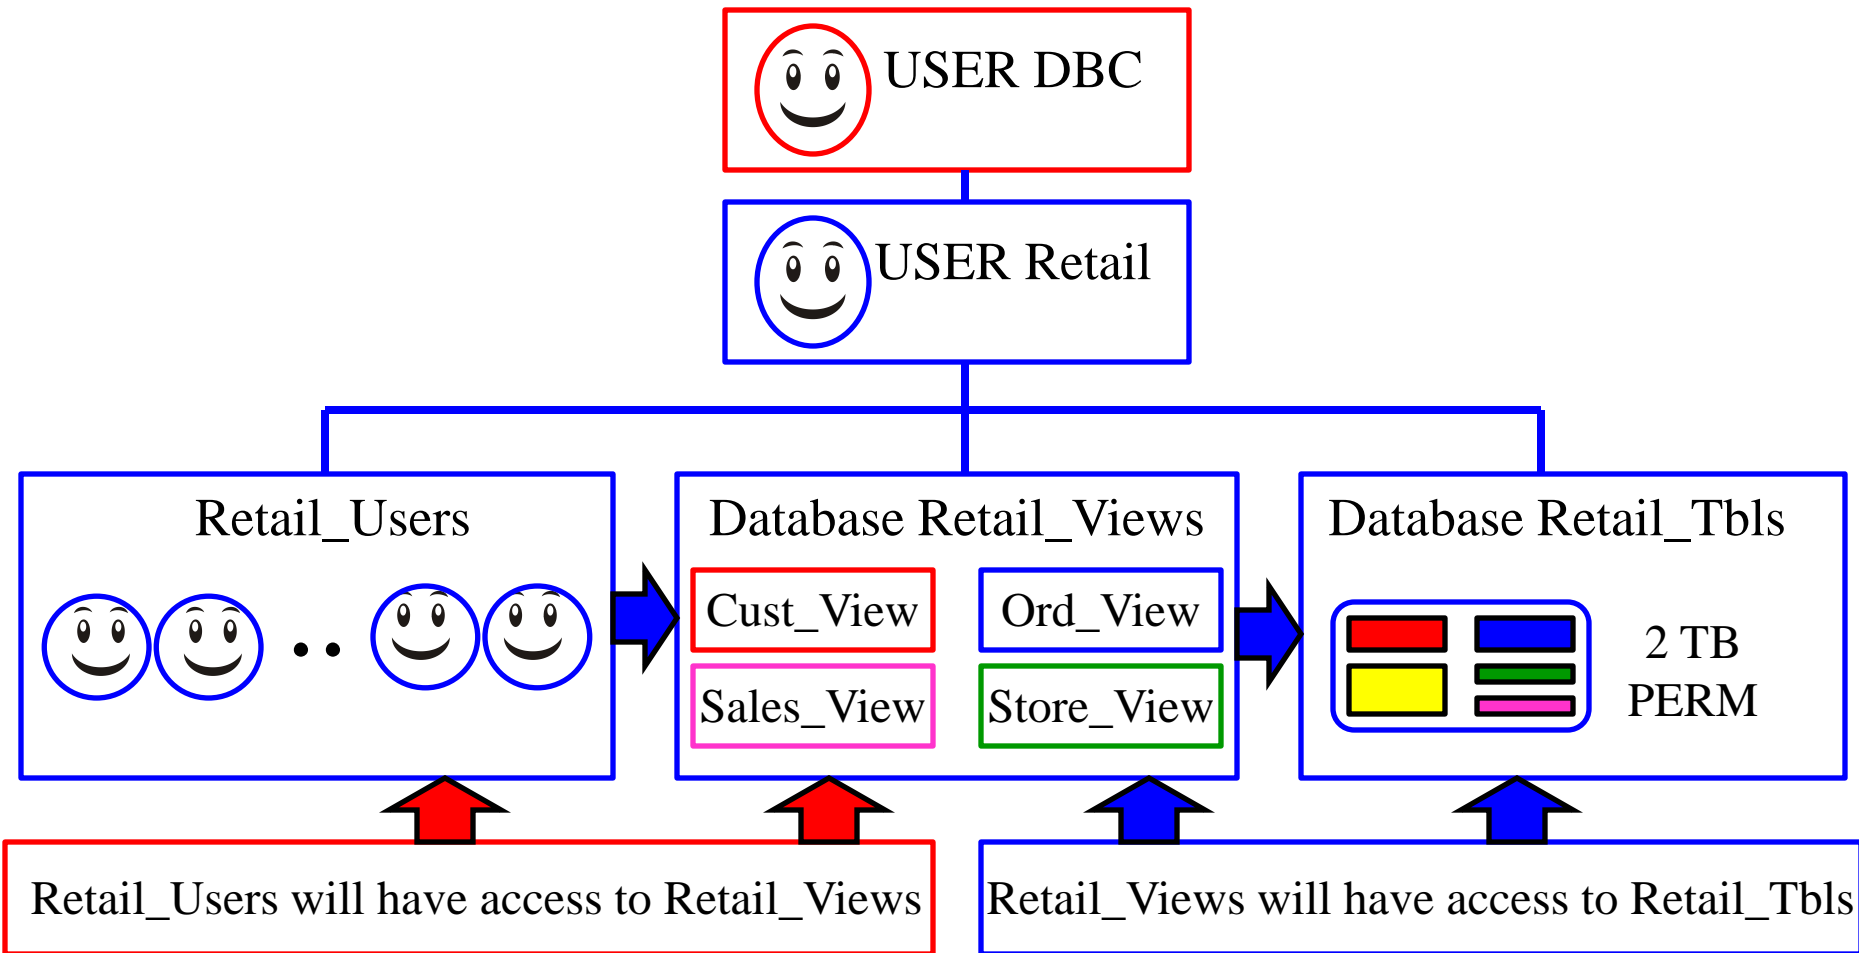
USER Retail and USER Financial now create the databases and users desired.

Differences between PERM and SPOOL



There are 1,000 users in Retail. Since Retail has 10 GB of spool that means that every user gets 10 GB of spool. That is the maximum limit for Retail. What it does **NOT** mean is that Retail is limited to only 10 GB of spool in total. Every user could logon and run a 9 GB query taking up Terabytes of Spool and nobody would run out of spool. Spool is system wide and calculated on an individual level only.

Databases, Users, and Views



For security purposes the Retail tables will be kept in their own database called **Retail_Tbls** (in this example). The general Retail User Population will NOT have access directly to these tables. A Database called **Retail_Views** houses the views that access the tables. So, the DBA will create Access Rights that allow the views to read the tables and the Users to **SELECT** from the views.

What are Similarities between a DATABASE and a USER?



1 A Database or a User can be assigned PERM Space

A If the Database **Marketing** is assigned 10 GB of PERM that means it can hold up to 10 GB of Permanent Tables.

B If the User **Maria** is assigned 10 GB of PERM that means she can hold up to 10 GB of Permanent Tables.

2 A Database or a User can be assigned Spool Space

A If the Database **Marketing** is assigned 10 GB of Spool that means all users under marketing can each run 10 GB queries.

B If the User **Maria** is assigned 10 GB of Spool that means she can run up to 10 GB queries and any user created under Maria will default to 10 GB queries.

What is the Difference between a DATABASE and a USER?



A USER has a login and password and therefore can
run queries

Objects that take up PERM Space

Permanent Space (Perm space) is the maximum amount of storage assigned to a user or database for holding:

- Table Rows
- Fallback Tables
- Secondary Index Subtables
- Stored Procedures
- User Defined Functions (UDFs)
- Permanent Journals


Views and **Macros** do **NOT** take up any Perm Space!

A Series of Quizzes on Adding and Subtracting Space

Marketing
10 GB Perm
10 GB Spool

Sales
5 GB Perm
5 GB Spool

- 1 Marketing has 10 GB of Perm and Spool. Sales has 5 GB Perm and Spool.
- 2 Marketing then Creates Stan and gives him 1 GB Perm and 10 GB Spool.
- 3 Sales then Creates Mary and gives her 1 GB Perm and 5 GB Spool.




After creating users how much Perm / Spool is in Marketing and how much is in Sales?

Answer 1 to Quiz on Space

Marketing
10 GB Perm
10 GB Spool

Sales
5 GB Perm
5 GB Spool

- 1 Marketing has 10 GB of Perm and Spool. Sales has 5 GB Perm and Spool.
- 2 Marketing then Creates Stan and gives him 1 GB Perm and 10 GB Spool.
- 3 Sales then Creates Mary and gives her 1 GB Perm and 5 GB Spool.



After creating users how much Perm / Spool is in Marketing and how much is in Sales?


9 GB Perm 10 GB Spool

4 GB Perm

5 GB Spool

Space Transfer Quiz

- 1 If a USER is dropped their PERM Space goes up to their immediate parent.
- 2 If a USER is transferred (GIVE Statement) they take their space with them.



Stan has just been transferred to Sales.

After the transfer how much Perm / Spool is in:


Marketing _____

Sales _____

Stan _____

Answer to Space Transfer Quiz

- 1 If a USER is dropped their PERM Space goes up to their immediate parent.
- 2 If a USER is transferred (GIVE Statement) they take their space with them.




Stan has just been transferred to Sales.

After the transfer how much Perm / Spool is in:

Marketing	9 GB Perm	10 GB Spool
Sales	4 GB Perm	5 GB Spool
Stan	1 GB Perm	10 GB Spool

Drop Space Quiz

- 1 If a USER is dropped their PERM Space goes up to their immediate parent.
- 2 If a USER is transferred (GIVE Statement) they take their space with them.



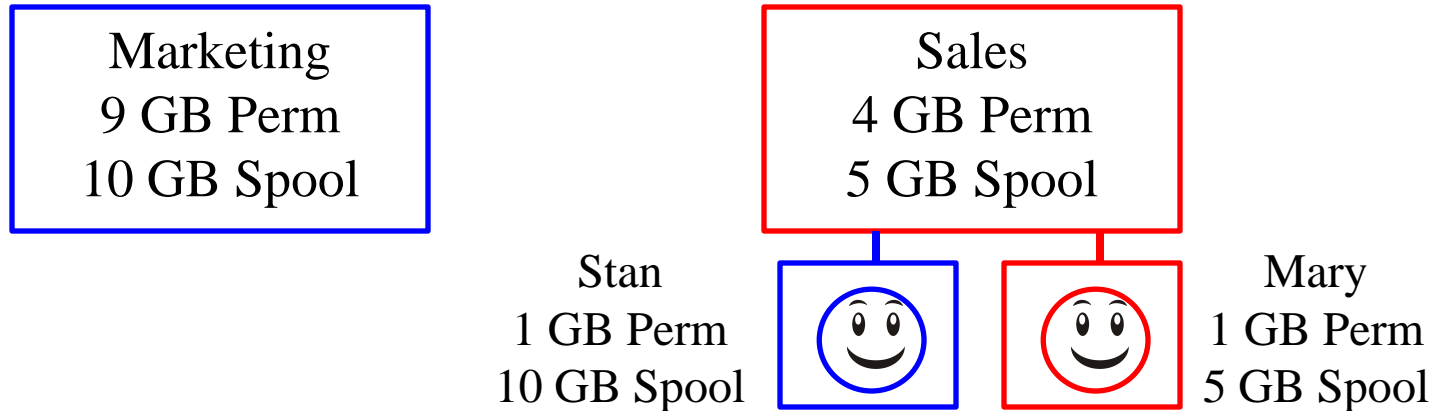
What happens **NOW** if Stan is Dropped?

After the drop how much Perm / Spool is in:

Marketing	_____	_____
Sales	_____	_____
Stan	_____	_____

Answers to Drop Space Quiz

- 1 If a USER is dropped their PERM Space goes up to their immediate parent.
- 2 If a USER is transferred (GIVE Statement) they take their space with them.



What happens **NOW** if Stan is Dropped?

After the drop how much Perm / Spool is in:

Marketing	9 GB Perm	10 GB Spool
Sales	5 GB Perm	5 GB Spool
Stan	dropped (0)	dropped (0)

Date Functions

“An inch of time cannot be bought with an inch of gold.”

- Chinese Proverb

Dates are stored Internally as INTEGERS from a Formula

$$\text{INTEGERDATE} = ((\text{Year} - 1900) * 10000) + (\text{Month} * 100) + \text{Day}$$

/* Example – Tom's Birthday January 10, 1959 */

$$\begin{aligned}\text{INTEGERDATE} &= ((1959 - 1900) = 59 && \leftarrow \text{Year Portion} \\ &\quad * 10000) = 590000 \\ &\quad + (\text{Month} * 100) = 590100 && \leftarrow \text{Month Portion} \\ &\quad + \text{Day} = 590110 && \leftarrow \text{Day Portion}\end{aligned}$$

/* Example – Tom's Birthday January 10, 1999 */

990110

/* Example – Tom's Birthday January 10, 2000 */

1000110

The way the Smart Calendar works so well is that it stores EVERY date in Teradata as something known as an **INTEGERDATE**.

Date, Time, and Timestamp Keywords

SELECT Date	AS "Date"
,Current_Date	AS ANSI_Date
,Time	AS "Time"
,Current_Time	AS ANSI_Time
,Current_Timestamp(6)	AS ANSI_Timestamp

Answer Set

<u>Date</u>	<u>ANSI Date</u>	<u>Time</u>	<u>ANSI Time</u>	<u>ANSI Timestamp</u>
2011/03/22	2011/03/22	10:34:44	10:34:44	2011/03/22 10:34:44.123456 -04:00

There's **no** keyword **Timestamp**, but only ANSI's **Current_Timestamp**

Above are the keywords you can utilize to get the date, time, or timestamp. These are reserved words that the system will deliver to you when requested.

INTEGER Date Vs ANSIDATE is how the Date is Displayed

SELECT Date ,Current_Date	AS "Date" AS ANSI_Date
------------------------------	---------------------------

INTEGERDATE (YY/MM/DD)

June 30, 2012

<u>Date</u>	<u>ANSI_Date</u>
12/06/30	12/06/30

ANSIDATE (YYYY-MM-DD)

June 30, 2012

<u>Date</u>	<u>ANSI_Date</u>
2012-06-30	2012-06-30

NEXUS Query Chameleon MM-DD-YYYY

<u>Date</u>	<u>ANSI_Date</u>
06-30-2012	06-30-2012

Teradata in release V2R3 defaulted to a display of YY/MM/DD. This is called the INTEGERDATE. This can be changed to ANSIDATE, which is YYYY-MM-DD for a specific session or by Default if the DBA changes the DATEFORM in DBS Control. This has nothing to do with how the date is stored internally. It has to do with the display of dates when using any ODBC tool or load utility. Above are some examples.

DATEFORM

DATEFORM Controls the default display of dates.

DATEFORM display choices are either **INTEGERDATE** or **ANSIDATE**.

INTEGERDATE is (YY/MM/DD) and **ANSIDATE** is (YYYY-MM-DD).

DATEFORM is the expected format for import and export of dates in Load Utilities.

Can be over-ridden by **USER** or within a Session at any time.

The Default can be changed by the DBA by changing the **DATEFORM** in DBSControl.

INTEGERDATE (YY/MM/DD)

June 30, 2012

Date	ANSI_Date
12/06/30	12/06/30

ANSIDATE (YYYY-MM-DD)

June 30, 2012

Date	ANSI_Date
2012-06-30	2012-06-30

Teradata in release V2R3 defaulted to a display of YY/MM/DD. This is called the **INTEGERDATE**. This can be changed to **ANSIDATE**, which is YYYY-MM-DD for a specific session or by Default if the DBA changes the **DATEFORM** in DBS Control. This has nothing to do with how the date is stored internally. It has to do with the display of dates when using any ODBC tool or load utility.

Changing the DATEFORM in Client Utilities such as BTEQ

Enter your logon or BTEQ Command:

.logon localtd/dbc

Password: *****

Logon successfully completed

BTEQ – Enter your DBC/SQL request or BTEQ command:

SELECT DATE;

Date

12/06/30

← INTEGERDATE is the Default

BTEQ – Enter your DBC/SQL request or BTEQ command:

SET Session DATEFORM = ANSIDATE;

← Changing the DATEFORM
for this BTEQ session.

SELECT DATE;

Current Date

2012-06-30

← ANSIDATE is the Display Form

Notice the
Word **Date**

Notice the
Word
Current_Date

Date, Time, and Timestamp Recap

SELECT Date	AS "Date"
,Current_Date	AS ANSI_Date

INTEGERDATE (YY/MM/DD)

June 30, 2012

Date	ANSI_Date
12/06/30	12/06/30

ANSIDATE (YYYY-MM-DD)

June 30, 2012

Date	ANSI_Date
2012-06-30	2012-06-30

Dates are converted to an **integer** through a formula before being **stored**.

Dates are **displayed** by **default** as INTEGERDATE **YY-MM-DD**.

The DBA can set up the system to display as **ANSIDATE YYYY-MM-DD**.

Keywords **Date** or **Current_Date** will return the date automatically.





Time, **Current_Time** and **Current_Timestamp** are keywords.

The **Nexus** Query Chameleon displays dates as **MM-DD-YYYY**.

Timestamp Differences

```
SELECT Current_Timestamp(0) AS Col1  
       ,Current_Timestamp(6) AS Col2
```

Answer Set

Col1	Col2
2011/03/22 10:34:44	2011/03/22 10:34:44.123456
 Date  Space  Time	 Milliseconds

A timestamp has the date separated by a space and the time. In our second example we have asked for 6 milliseconds.

Troubleshooting Timestamp

```
SELECT Timestamp(0) AS Col1  
       , Timestamp(6) AS Col2
```

Error

There is Date and Current_Date (both work).

There is Time and Current_Time (both work).

There is **NO** Timestamp, but only **Current_Timestamp**!

There is **NO** Timestamp command, but only ANSI's **Current_Timestamp**!

Add or Subtract Days from a date

```
SELECT Order_Date
       ,Order_Date + 60 as "Due Date"
       ,Order_Total
       ,"Due date" -10  as Discount
       ,Order_Total *.98 (FORMAT '$$$,$$$$.99', Title 'Discounted')
FROM   Order_Table
ORDER BY 1 ;
```

<u>Order_Date</u>	<u>Due Date</u>	<u>Order_Total</u>	<u>Discount</u>	<u>Discounted</u>
05/04/1998	07/03/1998	12347.53	06/23/1998	12100.57
01/01/1999	03/02/1999	8005.91	02/20/1999	7845.79
09/09/1999	11/08/1999	23454.84	10/29/1999	22985.74
10/01/1999	11/30/1999	5111.47	11/20/1999	5009.24
10/10/1999	12/09/1999	15231.62	11/29/1999	14926.98

When you **add** or **subtract** from a Date you are adding/subtracting **Days**

Because Dates are stored internally on disk as integers it makes it easy to add days to the calendar. In the query above we are adding 60 days to the Order_Date.

A Summary of Math Operations on Dates

- 1 DATE – DATE = Interval (days between dates)
- 2 DATE + or - Integer = Date

Let's find the number of days Tera-Tom has been alive since his last birthday.

```
SELECT (1120110(date)) - (590110 (date)) (Title 'Tera-Tom's Age In Days');  
  
Tera-Tom's Age In Days  
19358
```

Below is the same exact query, but with a clearer example of the dates.

```
SELECT ('1959-01-10'(date)) - ('2012-01-10' (date)) (Title 'Tera-Tom's Age In Days');  
  
Tera-Tom's Age In Days  
19358
```

A DATE – DATE is an interval of days between dates. A DATE + or – Integer = Date. Both queries above perform the same function, but the top query uses the internal date functions and the query on the bottom does dates the traditional way.

Using a Math Operation to find your Age in Years

- 1 DATE – DATE = Interval (days between dates)
- 2 DATE + or - Integer = Date

Let's find the number of **days** Tera-Tom has been alive since his last birthday.

```
SELECT (1120110(date)) - (590110 (date)) (Title 'Tera-Tom's Age In Days');
```

Tera-Tom's Age In Days

19358

Let's find the number of **years** Tera-Tom has been alive since his last birthday.

```
SELECT ((1120110(date)) - (590110 (date))) / 365 (Title 'Tera-Tom's Age In Years');
```

Tera-Tom's Age In Years

53

A DATE – DATE is an interval of days between dates. A DATE + or – Integer = Date. Both queries above perform the same function, but the top query uses the internal date functions and the query on the bottom does dates the traditional way.

Find What Day of the week you were Born

Let's find the actual day of the week Tera-Tom was born

```
SEL 'Tera-Tom was born on day ' || ((590110(date)) - (101(date))) MOD 7 (TITLE ' ');
```

Tera-Tom was born on day 5

This will produce
No Title

Result	Day of the Week
0	Monday
1	Tuesday
2	Wednesday
3	Thursday
4	Friday
5	Saturday
6	Sunday

This chart can be used
In conjunction with the
above SQL

The above subtraction results in the number of days between the two dates. Then, the MOD 7 divides by 7 to get rid of the number of weeks and results in the remainder. A MOD 7 can only result in values 0 thru 6 (always 1 less than the MOD operator). Since January 1, 1900 (101(date)) is a Monday, Tom was born on a Saturday.

The ADD_MONTHS Command

Order_Table

<u>Order_Number</u>	<u>Customer_Number</u>	<u>Order_Date</u>	<u>Order_Total</u>
123456	11111111	12347.53	1998/05/04
123512	11111111	8005.91	1999/01/01
123552	31323134	5111.47	1999/10/01
123585	87323456	15231.62	1999/10/10
123777	57896883	23454.84	1999/09/09

```
SELECT Order_Date
       ,Add_Months (Order_Date,2) as "Due Date"
       ,Order_Total
FROM   Order_Table ORDER BY 1 ;
```


<u>Order_Date</u>	<u>Due Date</u>	<u>Order_Total</u>
05/04/1998	07/04/1998	12347.53
01/01/1999	03/01/1999	8005.91
09/09/1999	11/09/1999	23454.84
10/01/1999	12/01/1999	5111.47
10/10/1999	12/10/1999	15231.62

This is the Add_Months Command. What you can do with it is add a month or many months your columns date. Can you convert this to one year?

Using the ADD_MONTHS Command to Add 1-Year

Order_Table			
<u>Order_Number</u>	<u>Customer_Number</u>	<u>Order_Date</u>	<u>Order_Total</u>
123456	11111111	12347.53	1998/05/04
123512	11111111	8005.91	1999/01/01
123552	31323134	5111.47	1999/10/01
123585	87323456	15231.62	1999/10/10
123777	57896883	23454.84	1999/09/09

```
SELECT Order_Date
       ,Add_Months (Order_Date,12) as "Due Date"
       ,Order_Total
FROM   Order_Table
ORDER BY 1 ;
```




There is **no** Add_Year command, so put in **12** months for 1-year

The Add_Months command adds months to any date. Above we used a great technique that would give us 1-year. Can you give me **5-years**?

Using the ADD_MONTHS Command to Add 5-Years

Order_Table			
<u>Order_Number</u>	<u>Customer_Number</u>	<u>Order_Date</u>	<u>Order_Total</u>
123456	11111111	12347.53	1998/05/04
123512	11111111	8005.91	1999/01/01
123552	31323134	5111.47	1999/10/01
123585	87323456	15231.62	1999/10/10
123777	57896883	23454.84	1999/09/09

```
SELECT Order_Date
       ,Add_Months (Order_Date,12 * 5) as "Due Date"
       ,Order_Total
FROM   Order_Table
ORDER BY 1 ;
```



In this example we multiplied **12 months** times **5** for a total of **5 years!**

Above you see a great technique for adding multiple years to a date. Can you now **SELECT** only the orders in **September?**

The EXTRACT Command

Order_Table			
<u>Order_Number</u>	<u>Customer_Number</u>	<u>Order_Date</u>	<u>Order_Total</u>
123456	11111111	12347.53	1998/05/04
123512	11111111	8005.91	1999/01/01
123552	31323134	5111.47	1999/10/01
123585	87323456	15231.62	1999/10/10
123777	57896883	23454.84	1999/09/09

```
SELECT Order_Date
       ,Add_Months (Order_Date,12 * 5) as "Due Date"
       ,Order_Total
FROM   Order_Table
WHERE  EXTRACT(Month from Order_Date) = 09
ORDER BY 1 ;
```

The EXTRACT command **extracts** portions of Date, Time, and Timestamp.

This is the Extract command. It extracts a portion of the date and it can be used in the SELECT list or the WHERE Clause, or the ORDER BY Clause!

EXTRACT from DATES and TIME

```
SELECT Current_Date
      ,EXTRACT(Year from Current_Date) as Yr
      ,EXTRACT(Month from Current_Date) as Mo
      ,EXTRACT(Day from Current_Date) as Da
      ,Current_Time
      ,EXTRACT(Hour from Current_Time) as Hr
      ,EXTRACT(Minute from Current_Time) as Mn
      ,EXTRACT(Second from Current_Time) as Sc
      ,EXTRACT(TIMEZONE_HOUR from Current_Time) as Th
      ,EXTRACT(TimeZONE_MINUTE from Current_Time) as Tm
```

Answer Set

<u>Order</u>	<u>Date</u>	<u>Yr</u>	<u>Mo</u>	<u>Day</u>	<u>Current Time (0)</u>	<u>Hr</u>	<u>Mn</u>	<u>Sc</u>	<u>Th</u>	<u>Tm</u>
	2011/03/22	2011	03	22	20:01:14 20	1	14	0	0	0

Just like the Add_Months, the EXTRACT Command is a Temporal Function or a Time-Based Function.

CURRENT_DATE and Math to get Temporal Functions

```
SELECT Current_Date  
      ,EXTRACT(Year from Current_Date) as Yr  
      ,EXTRACT(Month from Current_Date) as Mo  
      ,EXTRACT(Day from Current_Date) as Da  
      ,Current_Date / 10000 +1900 as YrMath  
      ,(Current_Date / 100) Mod 100 as MoMath  
      ,Current_Date Mod 100 as DayMath ;
```

Math can be used to **extract portions** of a **Date!**

Answer Set

<u>Order Date</u>	<u>Yr</u>	<u>Mo</u>	<u>Day</u>	<u>YrMath</u>	<u>MoMath</u>	<u>DayMath</u>
2011/03/22	2011	03	22	2011	03	22

The Extract Temporal Function can be used to extract a portion of a date. As you can see, Basic Arithmetic accomplish the same thing.

CAST the Date of January 1, 2011 and the Year 1800

```
SELECT
  cast('2011-01-01' as date)   as ANSI_Literal
,cast(1110101 as date)         as INTEGER_Literal
,cast('11-01-01' as date)     as YY_Literal
,cast(Date '2011-01-01' as Integer) as Dates_Stored
,cast(Date '1800-01-01' as Integer) as Dates_1800s
```

Answer Set

<u>ANSI_Literal</u>	<u>INTEGER_Literal</u>	<u>YY_Literal</u>	<u>Dates_Stored</u>	<u>Dates_1800s</u>
01/01/2011	01/01/2011	01/01/1911	111010	-999899

The Convert And Store (CAST) command is used to give columns a different data type temporarily for the life of the query. Notice our dates and how their stored.

The System Calendar

Teradata systems have a **table** called **Caldates**.

Caldates has only one column in it called **Cdates**.

Cdates is a date column that contains a row for each date starting from January 1, **1900** to December 31, **2100**.

No user can access the table Caldates directly.

Views in the **Sys_Calendar** database accesses Caldates.

A **view** called Calendar is how USER's work with the calendar.

Users use **Sys_Calendar.Calendar** for advanced dates.

In every Teradata system, they have something known as a **System Calendar** (or as Teradata calls it **Sys_Calendar.Calendar**). Get ready for AWESOME!

```
SELECT * FROM Sys_Calendar.Calendar  
WHERE Calendar_Date = '1959-01-10';
```

Birthday of
Tera-Tom



Calendar_Date = 01/10/1959'

day_of_week = 7 (Sunday = 1)

day_of_month = 10

day_of_year = 10

day_of_Calendar = 21559 (since Jan 1, 1900)

weekday_of_month = 2

week_of_month = 1 (0 for partial week for any month not starting with Sunday)

week_of_year = 1

week_of_calendar = 3079 (since Jan 1, 1900)

month_of_quarter = 1

month_of_year = 1

month_of_calendar = 709 (since Jan 1, 1900)

quarter_of_year = 1

quarter_of_calendar = 237 (since Jan 1, 1900)

year_of_calendar = 1959

Tera-Tom was born on a Saturday! It was the first full week of the month, the first full week of the year and it was the first quarter of the year!

How to really use the Sys_Calendar.Calendar

```
SELECT O.*  
FROM Order_Table as O  
INNER JOIN  
    Sys_Calendar.Calendar  
ON Order_Date = Calendar_Date  
AND Quarter_Of_Year = 4  
AND Day_of_Week = 6  
AND Week_of_Month = 0;
```

Join a date column
with the
Calendar_Date

<u>Order_Number</u>	<u>Customer_Number</u>	<u>Order_Date</u>	<u>Order_Total</u>
123552	31323134	10/01/1999	5111.47

We just brought back all Orders from the Order_Table that were purchased on a **Friday** in the 4th **Quarter**, during the **1st partial week**. This means **no Sunday seen yet** for that month.

Above is the perfect example of how you can utilize the Sys_Calendar.Calendar to join to any date field and then expand your search options.

Storing Dates Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL ,  
    CHECKSUM = DEFAULT  
    (Date_col          Date,  
     TIME_col          TIME(6),  
     TIMETIMEZONE_col TIME(6) WITH TIME ZONE,  
     TIMESTAMP_col     TIMESTAMP(6),  
     TIMEZONE_col      TIMESTAMP(6) WITH TIME ZONE)  
    UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

DATE '1999-01-10' is stored as 990110

DATE '2000-01-10' is stored as 1000110

4-bytes store Date_col internally because dates are considered a 4-byte integer.

Storing Time Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL,  
CHECKSUM = DEFAULT  
    (Date_col          Date,  
     TIME_col          TIME(6),  
     TIMETIMEZONE_col  TIME(6) WITH TIME ZONE,  
     TIMESTAMP_col     TIMESTAMP(6),  
     TIMEZONE_col       TIMESTAMP(6) WITH TIME ZONE)  
UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

Time(n) stored as HHMMSS.nnnnnn

It takes 6 bytes to store Time_col internally.

Storing TIME With TIME ZONE Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL ,  
CHECKSUM = DEFAULT  
    (Date_col          Date,  
    TIME_col          TIME(6),  
    TIMETIMEZONE_col  TIME(6) WITH TIME ZONE,  
    TIMESTAMP_col     TIMESTAMP(6),  
    TIMEZONE_col      TIMESTAMP(6) WITH TIME ZONE)  
UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

Time(n) WITH ZONE stored as **HHMMSS.nnnnnn+HHMM**

It takes 8 bytes to store TimeTimezone_col internally.

Storing Timestamp Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL ,  
    CHECKSUM = DEFAULT  
    (Date_col                Date,  
     TIME_col                TIME(6),  
     TIMETIMEZONE_col        TIME(6) WITH TIME ZONE,  
     TIMESTAMP_col           TIMESTAMP(6),  
     TIMEZONE_col            TIMESTAMP(6) WITH TIME ZONE)  
    UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

TimeStamp(n) stored as **YYMMDDHHMMSS.nnnnnn**

It takes **10 bytes** to store TimeStamp_col internally.

Storing Timestamp with TIME_ZONE Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL ,  
CHECKSUM = DEFAULT  
    (Date_col                Date,  
    TIME_col                 TIME(6),  
    TIMETIMEZONE_col         TIME(6) WITH TIME_ZONE,  
    TIMESTAMP_col            TIMESTAMP(6),  
    TIMEZONE_col              TIMESTAMP(6) WITH TIME_ZONE)  
UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

TimeStamp(n) With Zone stored as
YYMMDDHHMMSS.nnnnnn+HHMM

It will take 12 bytes to store Timezone_col internally.

Storing Date, Time, Timestamp with Zone Internally

```
CREATE SET TABLE TIMEZONE_table ,FALLBACK ,  
    NO BEFORE JOURNAL,  
    NO AFTER JOURNAL ,  
    CHECKSUM = DEFAULT  
    (Date_col                Date,  
     TIME_col                TIME(6),  
     TIMETIMEZONE_col        TIME(6) WITH TIME ZONE,  
     TIMESTAMP_col           TIMESTAMP(6),  
     TIMEZONE_col            TIMESTAMP(6) WITH TIME ZONE)  
    UNIQUE PRIMARY INDEX ( TIMEZONE_col );
```

Date	Stored Internally	4 Bytes
Time(n)	Stored Internally	6 Bytes
Time(n) With Zone	Stored Internally	8 Bytes
Timestamp(n)	Stored Internally	10 Bytes
Timestamp(n) with zone	Stored Internally	12 Bytes

Each data type increase their internal storage by 2 bytes.

Time Zones

A time zone relative to **London** (UTC) might be:

LA-----Miami-----Frankfurt-----Hong Kong			
+8:00	+05:00	00:00	-08:00

A time zone relative to **New York** (EST) might be:

LA-----Miami-----Frankfurt-----Hong Kong			
+3:00	00:00	-05:00	-13:00

Time zones are set either at the system level (DBS Control), the user level (when user is created or modified), or at the session level as an override.

Teradata has the ability to access and store both the hours and the minutes reflecting the difference between the user's time zone and the system time zone. From a World perspective, this difference is normally the number of hours between a specific location on Earth and the United Kingdom location that was historically called Greenwich Mean Time (GMT). Since the Greenwich observatory has been "decommissioned," the new reference to this same time zone is called Universal Time Coordinate (UTC).

Setting Time Zones

A Time Zone should be established for the system and every user in each different time zone.

Setting the **system default** time zone is done by the DBA in the **DBSControl** record:

```
MODIFY GENERAL 16 = x    /* Hours,  n= -12 to 13    */  
MODIFY GENERAL 17 = x    /* Minutes, n = -59 to 59    */
```

Setting a **User's** time zone requires choosing either **LOCAL**, **NULL**, or an **explicit value**:

```
CREATE USER  Tera-Tom  
TIME ZONE    = LOCAL    /* use system level      */  
              = NULL     /* no default, set to system or session level at logon */  
              = '16:00'  /* explicit setting      */  
              = '-06:30' /* explicit setting      */
```


Setting a **Session's** time zone:

```
SET TIME ZONE LOCAL ; /* use system level    */  
SET TIME ZONE USER ; /* use user level      */  
SET TIME ZONE INTERVAL '08:00' HOUR TO MINUTE ; /* explicit setting    */
```

A Teradata session can modify the time zone without requiring a logoff and logon

Seeing your Time Zone

Help Session ;



<u>User Name</u>	<u>Account Name</u>	<u>Logon Date</u>	<u>Logon Time</u>	<u>Current Database</u>	<u>Collation</u>	<u>Char Set</u>	<u>Transaction Semantics</u>	<u>Current Dateform</u>	<u>Session Time Zone</u>
DBC	DBC	12/06/17	15:55:39	SQL_CLASS	ASCII	ASCII	Teradata	IntegerDate	00:00

Not all output
is displayed
above from the
HELP Session

A user's time zone is now part of the information maintained by Teradata. The settings can be seen in the extended information available in the HELP SESSION request. Teradata converts all TIME and TIMESTAMP values to Universal Time Coordinate (UTC) prior to storing them. All operations, including hashing, collation, and comparisons that act on TIME and TIMESTAMP values are performed using their UTC forms. This will allow users to CAST the information to their local times.

Creating a Sample Table for Time Zone Examples

```
CREATE TABLE Tstamp_Test
(
  TS_Zone CHAR(3)
,TS_with_Zone TIMESTAMP(6) WITH TIME ZONE
,TS_Without_Zone TIMESTAMP(6)
)
UNIQUE PRIMARY INDEX ( TS_Zone );
```

Not all output
is displayed
above from the
HELP Session

A user's time zone is now part of the information maintained by Teradata. The settings can be seen in the extended information available in the **HELP SESSION** request.

Selecting the Data from our Time Zone Table

```
SELECT * FROM Tstamp_Test ;
```

<u>TS_Zone</u>	<u>TS_with_Zone</u>	<u>TS_Without_Zone</u>
UTC	2000-10-01 08:12:00.000000+05:00	2000-10-01 08:12:00.000000
EST	2000-10-01 08:12:00.000000+00:00	2000-10-01 08:12:00.000000
PST	2000-10-01 08:12:00.000000-03:00	2000-10-01 08:12:00.000000
HKT	2000-10-01 08:12:00.000000-11:00	2000-10-01 08:12:00.000000



Notice the
Accompanying
Time Zone Offsets

Our Insert statements were done at 08:12:00 exactly. Notice the Time Zone offsets in the column TS_with_Zone and how their not there for the column TS_Without_Zone. Teradata converts all TIME and TIMESTAMP values to Universal Time Coordinate (UTC) prior to storing them. All operations, including hashing, collation, and comparisons that act on TIME and TIMESTAMP values are performed using their UTC forms. This will allow users to CAST the information to their local times.

Normalizing our Time Zone Table with a CAST

```
SELECT TS_Zone, TS_with_Zone
       ,CAST(TS_with_Zone AS TIMESTAMP(6)) AS T_Normal
FROM Tstamp_Test ORDER BY 3 ;
```

TS_Zone	TS_with_Zone	T_Normal
UTC	2000-10-01 08:12:00.000000+05:00	2000-10-01 03:12:00.000000
EST	2000-10-01 08:12:00.000000+00:00	2000-10-01 08:12:00.000000
PST	2000-10-01 08:12:00.000000-03:00	2000-10-01 11:12:00.000000
HKT	2000-10-01 08:12:00.000000-11:00	2000-10-01 19:12:00.000000



The System is on EST Time. The New Times are Normalized to the time zone of the System!

Notice that the Time Zone value was added to or subtracted from the time portion of the time stamp to adjust them to a perspective of the same time zone. As a result, at that moment, it has normalized the different Times Zones in respect to the system time.

As an illustration, when the transaction occurred at 8:12 AM locally in the PST Time Zone, it was already 11:12 AM in EST, the location of the system. The times in the columns have been **normalized** in respect to the time zone of the **system**.

Intervals for Date, Time and Timestamp

Interval Chart

Simple Intervals	More involved Intervals
YEAR	DAY TO HOUR
MONTH	DAY TO MINUTE
DAY	DAY TO SECOND
HOUR	HOUR TO MINUTE
MINUTE	HOUR TO SECOND
SECOND	MINUTE TO SECOND

To make Teradata SQL more ANSI compliant and compatible with other RDBMS SQL, Teradata has added INTERVAL processing. Intervals are used to perform DATE, TIME and TIMESTAMP arithmetic and conversion.

Although Teradata allowed arithmetic on DATE and TIME, it was not performed in accordance to ANSI standards and therefore, an extension instead of a standard. With INTERVAL being a standard instead of an extension, more SQL can be ported directly from an ANSIcompliant database to Teradata without conversion.

Interval Data Types and the Bytes to Store Them

Interval Chart

Bytes	Data Type	Comments
2	INTERVAL YEAR	10 for 32-bit systems; 12 for 64-bit
4	INTERVAL YEAR TO MONTH	
2	INTERVAL MONTH	
2	INTERVAL MONTH TO DAY	
2	INTERVAL DAY	
8	INTERVAL DAY TO MINUTE	
10/12	INTERVAL DAY TO SECOND	
2	INTERVAL HOUR 2	
4	INTERVAL HOUR TO MINUTE 4	
8	INTERVAL HOUR TO SECOND 8	
2	INTERVAL MINUTE 2	6 for 32-bit systems; 8 for 64-bit 6 for 32-bit systems; 8 for 64-bit
6/8	INTERVAL MINUTE TO SECOND	
6/8	INTERVAL SECOND	

The Basics of a Simple Interval

```
SELECT Current_Date as Our_Date  
       ,Current_Date + Interval '1' Day      as Plus_1_Day  
       ,Current_Date + Interval '3' Month    as Plus_3_Months  
       ,Current_Date + Interval '5' Year     as Plus_5_Years
```

<u>Our_Date</u>	<u>Plus_1_Day</u>	<u>Plus_3_Months</u>	<u>Plus_5_Years</u>
06/18/2012	06/19/2012	09/18/2012	06/18/2017

In the example SQL above we take a simple date and add 1 day, 3 months and 5 years. Notice that our current_date is 06/18/2012 and that our intervals come out perfectly.

Troubleshooting The Basics of a Simple Interval

```
SELECT Date '2012-01-29' as Our_Date  
      ,Date '2012-01-29' + INTERVAL '1' Month as Leap_Year
```

<u>Our_Date</u>	<u>Leap_Year</u>
01/29/2012	02/29/2012

```
SELECT Date '2011-01-29' as Our_Date  
      ,Date '2011-01-29' + INTERVAL '1' Month as Leap_Year
```

Error – Invalid Date

The first example works because we added 1 month to the date '2012-01-29' and we got '2012-02-29'. Because this was leap year there actually is a date of February 29, 2012. The next example is the real point. We have a date of '2011-01-29' and we add 1-month to that, but there is no February 29th in 2011 so the query fails.

Interval Arithmetic Results

DATE and TIME arithmetic Results using intervals:

DATE	- DATE	= Interval
TIME	- TIME	= Interval
TIMESTAMP	- TIMESTAMP	= Interval
DATE	- or + Interval	= DATE
TIME	- or + Interval	= TIME
TIMESTAMP	- or + Interval	= TIMESTAMP
Interval	- or + Interval	= Interval

To use DATE and TIME arithmetic, it is important to keep in mind the results of various operations. The above chart is your Interval guide.

A Date Interval Example

```
SELECT (DATE '1999-10-01' - DATE '1988-10-01') DAY AS Actual_Days ;
```

ERROR – Interval Field Overflow

The **Error** occurred because the default for all intervals is 2 digits.

```
SELECT (DATE '1999-10-01' - DATE '1988-10-01') DAY(4) AS Actual_Days ;
```

Makes the output 4 digits

<u>Actual_Days</u>
4017

The default for all intervals is 2 digits. We received an overflow error because the Actual_Days is 4017. The second example works because we demanded the output to be 4 digits (the maximum for intervals).

A Time Interval Example

Makes the output 3 digits

```
SELECT (TIME '12:45:01' - TIME '10:10:01') HOUR AS Actual_Hours
      ,(TIME '12:45:01' - TIME '10:10:01') MINUTE(3) AS Actual_Minutes
      ,(TIME '12:45:01' - TIME '10:10:01') SECOND(4) AS Actual_Seconds
      ,(TIME '12:45:01' - TIME '10:10:01') SECOND(4,4) AS Actual_Seconds4
```

<u>Actual_Hours</u>	<u>Actual_Minutes</u>	<u>Actual_Seconds</u>	<u>Actual_Seconds4</u>
2	155	9300.000000	9300.0000

```
SELECT (TIME '12:45:01' - TIME '10:10:01') HOUR AS Actual_Hours
      ,(TIME '12:45:01' - TIME '10:10:01') MINUTE AS Actual_Minutes
      ,(TIME '12:45:01' - TIME '10:10:01') SECOND(4) AS Actual_Seconds
      ,(TIME '12:45:01' - TIME '10:10:01') SECOND(4,4) AS Actual_Seconds4
```

ERROR – Interval Field Overflow

The default for all intervals is 2 digits, but notice in the top example we put in 3 digits for Minute, 4 digits for Second and 4,4 digits for the Actual_Seconds4. If we had not we would have received an overflow error as in the bottom example.

A - DATE Interval Example

```
SELECT Current_Date,  
       INTERVAL '-2' YEAR + CURRENT_DATE as Two_years_Ago;
```

<u>Date</u>	<u>Two_Year_Ago</u>
06/18/2012	06/18/2010

The above Interval example uses a –'2' to go back in time.

A Complex Time Interval Example using CAST

Below is the syntax for using the CAST with a date:

```
SELECT CAST (<interval> AS INTERVAL <interval> )  
FROM <table-name> ;
```

The following converts an INTERVAL of 6 years and 2 months to an INTERVAL number of months:

```
SELECT CAST( (INTERVAL '6-02' YEAR TO MONTH) AS INTERVAL MONTH );
```

$$\frac{6-02}{74}$$

The CAST function (Convert And Store) is the ANSI method for converting data from one type to another. It can also be used to convert one INTERVAL to another INTERVAL representation. Although the CAST is normally used in the SELECT list, it works in the WHERE clause for comparison reasons.

A Complex Time Interval Example using CAST

This request attempts to convert 1300 months to show the number of years and months.
Why does it fail?

```
SELECT CAST(INTERVAL '1300' MONTH AS INTERVAL YEAR TO MONTH)  
(Title 'Years & Months') ;
```

ERROR

```
SELECT CAST(INTERVAL '1300' MONTH as interval YEAR(3) TO MONTH) ;
```

Years & Month
108-04

The top query failed because the INTERVAL result defaults to 2-digits and we have a 3-digit answer for the year portion (108). The bottom query fixes that specifying 3-digits. The biggest advantage in using the INTERVAL processing is that SQL written on another system is now compatible with Teradata.

The OVERLAPS Command

Compatibility: Teradata Extension

The syntax of the OVERLAPS is:

```
SELECT <literal>  
      WHERE (<start-date-time>, <end-date-time>) OVERLAPS  
(<start-date-time>, <end-date-time>);
```

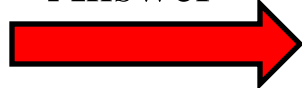
```
SELECT 'The Dates Overlap' (TITLE ' ')  
WHERE (DATE '2001-01-01', DATE '2001-11-30') OVERLAPS  
      (DATE '2001-10-15', DATE '2001-12-31');
```

Answer
 The Dates Overlap

When working with dates and times, sometimes it is necessary to determine whether two different ranges have common points in time. Teradata provides a Boolean function to make this test for you. It is called OVERLAPS; it evaluates true, if multiple points are in common, otherwise it returns a false. The literal is returned because both date ranges have from October 15 through November 30 in common.

An OVERLAPS Example that Returns No Rows

```
SELECT 'The dates overlap' (TITLE ' ')  
WHERE (DATE '2001-01-01', DATE '2001-11-30') OVERLAPS  
      (DATE '2001-11-30', DATE '2001-12-31') ;
```

Answer  No rows found

The above SELECT example tests two literal dates and uses the OVERLAPS to determine whether or not to display the character literal.

The literal was not selected because the ranges do not overlap. So, the common single date of November 30 does not constitute an overlap. When dates are used, 2 days must be involved and when time is used, 2 seconds must be contained in both ranges.

The OVERLAPS Command using TIME

```
SELECT 'The Times Overlap' (TITLE ' ' )  
WHERE (TIME '08:00:00', TIME '02:00:00') OVERLAPS  
      (TIME '02:01:00', TIME '04:15:00') ;
```

Answer  The Times Overlap

The above SELECT example tests two literal times and uses the OVERLAPS to determine whether or not to display the character literal.

This is a tricky example and it is shown to prove a point. At first glance, it appears as if this answer is incorrect because 02:01:00 looks like it starts 1 second after the first range ends. However, the system works on a 24-hour clock when a date and time (timestamp) is not used together. Therefore, the system considers the earlier time of 2AM time as the start and the later time of 8 AM as the end of the range. Therefore, not only do they overlap, the second range is entirely contained in the first range.

The OVERLAPS Command using a NULL Value

```
SELECT 'The Times Overlap' (TITLE ' ')  
WHERE (TIME '10:00:00', NULL) OVERLAPS (TIME '01:01:00', TIME '04:15:00')
```

Answer  No Rows Found

The above SELECT example tests two literal dates and uses the OVERLAPS to determine whether or not to display the character literal:

When using the OVERLAPS function, there are a couple of situations to keep in mind:

1. A single point in time, i.e. the same date, does not constitute an overlap. There must be at least one second of time in common for TIME or one day when using DATE.
2. Using a NULL as one of the parameters, the other DATE or TIME constitutes a single point in time versus a range.

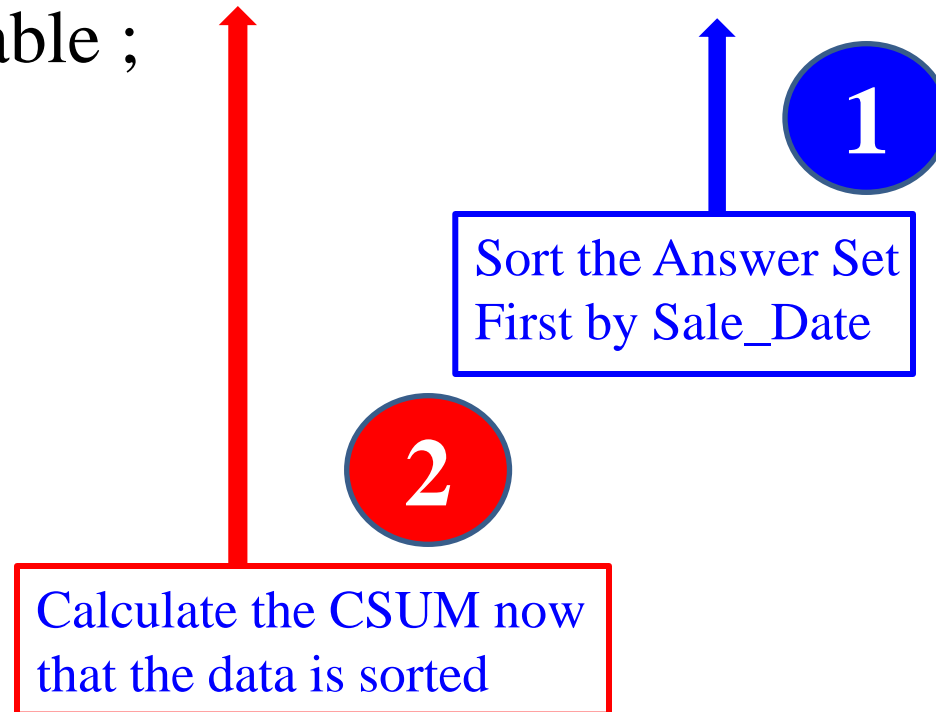
OLAP Functions

“The best we can do is size up the chances, calculate the risks involved, estimate our ability to deal with them, and then make our plans with confidence.”

- Henry Ford

On-Line Analytical Processing (OLAP) or Ordered Analytics

```
SELECT    Product_ID
          , Sale_Date
          , Daily_Sales
          , CSUM(Daily_Sales, Sale_Date) AS "CSum"
FROM Sales_Table ;
```



OLAP is often called Ordered Analytics because the first thing every OLAP does before any calculating is SORT all the rows. The query above sorts by Sale_Date!

Cumulative Sum (CSUM) Command and how OLAP Works

```
SELECT    Product_ID , Sale_Date, Daily_Sales
          ,CSUM(Daily_Sales, Sale_Date)  AS “CSum”
FROM Sales_Table ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>CSUM</u>
1000	2000-09-28	48850.40	
2000	2000-09-28	41888.88	
3000	2000-09-28	61301.77	
1000	2000-09-29	54500.22	
2000	2000-09-29	48000.00	
3000	2000-09-29	34509.13	
1000	2000-09-30	36000.07	
2000	2000-09-30	49850.03	
3000	2000-09-30	43868.86	
1000	2000-10-01	40200.43	
2000	2000-10-01	54850.29	
3000	2000-10-01	28000.00	
1000	2000-10-02	32800.50	
2000	2000-10-02	36021.93	
3000	2000-10-02	19678.94	

Not all rows
are displayed in
this answer set

Sort the Answer Set
first by Sale_Date, but
Don't do any CSUM
Calculations yet!

1

OLAP always sorts first and then is in a position to calculate starting with the first sorted row and continuing to the last sorted row, thus calculating all Daily_Sales.

OLAP Commands always Sort (ORDER BY) in the Command

```
SELECT    Product_ID , Sale_Date, Daily_Sales
          ,CSUM(Daily_Sales, Sale_Date)  AS “CSum”
FROM Sales_Table ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>CSUM</u>
1000	2000-09-28	48850.40	48850.40
2000	2000-09-28	41888.88	90739.28
3000	2000-09-28	61301.77	152041.05
1000	2000-09-29	54500.22	206541.27
2000	2000-09-29	48000.00	254541.27
3000	2000-09-29	34509.13	289050.40
1000	2000-09-30	36000.07	325050.47
2000	2000-09-30	49850.03	374900.50
3000	2000-09-30	43868.86	418769.36
1000	2000-10-01	40200.43	458969.79
2000	2000-10-01	54850.29	513820.08
3000	2000-10-01	28000.00	541820.08
1000	2000-10-02	32800.50	574620.58

Not all rows
are displayed in
this answer set

2

Calculate the
CSUM starting
with the first
sorted row and
go to the last.

Once the data is first sorted by Sale_Date then phase 2 is ready and the OLAP calculation can be performed on the sorted data. Day 1 we made 48850.40! Add the next row’s Daily_Sales to get a Cumulative Sum (CSUM) to get 90739.28!

Calculate the Cumulative Sum (CSUM) after Sorting the Data

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Sale_Date) AS "CSUM"  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>CSUM</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75

This is our first OLAP known as a CSUM. Right now, the syntax wants to see the cumulative sum of the Daily_Sales sorted by Sale_Date. The first thing the above query does before calculating is SORT all the rows on Sale_Date.

The OLAP Major Sort Key

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Sale_Date)  AS “CSum”  
FROM  Sales_Table WHERE Product_ID BETWEEN 1000 and 2000
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>CSum</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75

Not all rows
are displayed in
this answer set

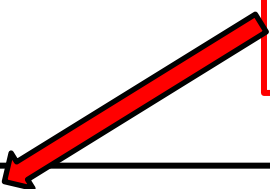
In a CSUM, the **second column** listed is always the **major SORT KEY**. The SORT KEY in the above query is Sale_Date. Notice again the answer set is sorted by this. After the sort has finished the CSUM is calculated starting with the first sorted row till the end.

The OLAP Major Sort Key and the Minor Sort Key(s)

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
           CSUM(Daily_Sales, Product_ID, Sale_Date)  AS “CSum”  
FROM Sales_Table;
```

Major
Sort Key

Minor
Sort Key




<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>CSUM</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	103350.62
1000	2000-09-30	36000.07	139350.69
1000	2000-10-01	40200.43	179551.12
1000	2000-10-02	32800.50	212351.62
1000	2000-10-03	64300.00	276651.62
1000	2000-10-04	54553.10	331204.72
2000	2000-09-28	41888.88	373093.60
2000	2000-09-29	48000.00	421093.60
2000	2000-09-30	49850.03	470943.63
2000	2000-10-01	54850.29	525793.92

Not all rows
are displayed in
this answer set

Product_ID is the MAJOR sort key and Sale_Date is the MINOR Sort key above.

Troubleshooting OLAP – My Data isn't coming back Correct

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            CSUM(Daily_Sales, Product_ID, Sale_Date)  AS "CSum"  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000  
ORDER BY Daily_Sales;
```




Don't place an
ORDER BY
at the end!

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>CSUM</u>
1000	2000-10-02	32800.50	212351.62
2000	2000-10-04	32800.50	637816.53
1000	2000-09-30	36000.07	139350.69
2000	2000-10-02	36021.93	561815.85
1000	2000-10-01	40200.43	179551.12
2000	2000-09-28	41888.88	373093.60
2000	2000-10-03	43200.18	605016.03
2000	2000-09-29	48000.00	421093.60
1000	2000-09-28	48850.40	48850.40
2000	2000-09-30	49850.03	470943.63
1000	2000-09-29	54500.22	103350.62
1000	2000-10-04	54553.10	331204.72
2000	2000-10-01	54850.29	525793.92
1000	2000-10-03	64300.00	276651.62

The first thing every OLAP does is **SORT**. That means you should **NEVER** put an **ORDER BY** at the end. It will mess up the **ENTIRE** result set.

GROUP BY in Teradata OLAP Syntax Resets on the Group

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            CSUM(Daily_Sales, Product_ID, Sale_Date)  AS "CSum"  
FROM Sales_Table  
GROUP BY Product_ID ;
```



<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>CSUM</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	103350.62
1000	2000-09-30	36000.07	139350.69
1000	2000-10-01	40200.43	179551.12
1000	2000-10-02	32800.50	212351.62
1000	2000-10-03	64300.00	276651.62
1000	2000-10-04	54553.10	331204.72
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	89888.88
2000	2000-09-30	49850.03	139738.91
2000	2000-10-01	54850.29	194589.20

Reset
now!




Not all rows
displayed In
answer set

The **GROUP BY** Statement cause the **CSUM** to start over (reset) on its calculating the cumulative sum of the **Daily_Sales** each time it runs into a **NEW Product_ID**.

CSUM the Number 1 to get a Sequential Number

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Product_ID, Sale_Date) AS "CSum",  
       CSUM(1, Product_ID, Sale_Date) as "Seq_Number"  
FROM Sales_Table ;
```



<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>CSUM</u>	<u>Seq_Number</u>
1000	2000-09-28	48850.40	48850.40	1
1000	2000-09-29	54500.22	103350.62	2
1000	2000-09-30	36000.07	139350.69	3
1000	2000-10-01	40200.43	179551.12	4
1000	2000-10-02	32800.50	212351.62	5
1000	2000-10-03	64300.00	276651.62	6
1000	2000-10-04	54553.10	331204.72	7
2000	2000-09-28	41888.88	373093.60	8
2000	2000-09-29	48000.00	421093.60	9
2000	2000-09-30	49850.03	470943.63	10

Not all rows
are displayed in
this answer set

With "Seq_Number", because you placed the number 1 in the area where it calculates, it will continuously add 1 to the answer for each row.

A Single GROUP BY Resets each OLAP with Teradata Syntax

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Product_ID, Sale_Date) AS "CSum",  
       CSUM(1, Product_ID, Sale_Date) as "Seq_Number"  
FROM Sales_Table GROUP BY Product_ID ;
```



Not all rows
are displayed in
this answer set


<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>CSUM</u>	<u>Seq_Number</u>
1000	2000-09-28	48850.40	48850.40	1
1000	2000-09-29	54500.22	103350.62	2
1000	2000-09-30	36000.07	139350.69	3
1000	2000-10-01	40200.43	179551.12	4
1000	2000-10-02	32800.50	212351.62	5
1000	2000-10-03	64300.00	276651.62	6
1000	2000-10-04	54553.10	331204.72	7
2000	2000-09-28	41888.88	41888.88	1
2000	2000-09-29	48000.00	89888.88	2
2000	2000-09-30	49850.03	139738.91	3
2000	2000-10-01	54850.29	194589.20	4

What does the GROUP BY Statement cause? Both OLAP Commands to reset!

A Better Choice – The ANSI Version of CSUM

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
                               ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM   Sales_Table  
WHERE  Product_ID BETWEEN 1000 and 2000 ;
```

Start on 1st row
and continue
till the end



<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SUMOVER</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75

Not all rows
are displayed in
this answer set

This ANSI version of CSUM is **SUM() Over**. Right now, the syntax wants to see the **sum** of the **Daily_Sales** after it is first sorted by **Sale_Date**. Rows Unbounded Preceding makes this a CSUM. The ANSI Syntax seems difficult, but only at first.

The ANSI Version of CSUM – The Sort Explained

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
       ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SUMOVER</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75
2000	2000-10-03	43200.18	550462.93

Not all rows
are displayed in
this answer set

The first thing the above query does before calculating is **SORT** all the rows by **Sale_Date**. The Sort is located right after the **ORDER BY**.

The ANSI CSUM – Rows Unbounded Preceding Explained

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
                             ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SUMOVER</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75

The keywords **Rows Unbounded Preceding** determines that this is a **CSUM**. There are only a few different statements and Rows Unbounded Preceding is the main one. It means start calculating at the beginning row and continue calculating until the last row..

The ANSI CSUM – Making Sense of the Data

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
                               ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>SUMOVER</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75

Not all rows
are displayed in
this answer set

The second “SUMOVER” row is **90739.28**. That is derived by the first row’s Daily_Sales (**41888.88**) added to the SECOND row’s Daily_Sales (**48850.40**).

The ANSI CSUM – Making Even More Sense of the Data

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
                               ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SUMOVER</u>
2000	2000-09-28	41888.88	41888.88
1000	2000-09-28	48850.40	90739.28
2000	2000-09-29	48000.00	138739.28
1000	2000-09-29	54500.22	193239.50
1000	2000-09-30	36000.07	229239.57
2000	2000-09-30	49850.03	279089.60
1000	2000-10-01	40200.43	319290.03
2000	2000-10-01	54850.29	374140.32
1000	2000-10-02	32800.50	406940.82
2000	2000-10-02	36021.93	442962.75
1000	2000-10-03	64300.00	507262.75

The third “SUMOVER” row is **138739.28**. That is derived by taking the first row’s Daily_Sales (**41888.88**) and adding it to the SECOND row’s Daily_Sales (**48850.40**). Then you add that total to the THIRD row’s Daily_Sales (**48000.00**).

The ANSI CSUM – The Major and Minor Sort Key(s)

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
       ROWS UNBOUNDED PRECEDING) AS SumOVER  
FROM Sales_Table ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SUMOVER</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	103350.62
1000	2000-09-30	36000.07	139350.69
1000	2000-10-01	40200.43	179551.12
1000	2000-10-02	32800.50	212351.62
1000	2000-10-03	64300.00	276651.62
1000	2000-10-04	54553.10	331204.72
2000	2000-09-28	41888.88	373093.60
2000	2000-09-29	48000.00	421093.60
2000	2000-09-30	49850.03	470943.63
2000	2000-10-01	54850.29	525793.92

Not all rows
are displayed in
this answer set

You can have more than one SORT KEY. In the top query, Product_ID is the MAJOR Sort and Sale_Date is the MINOR Sort.

The ANSI CSUM – Getting a Sequential Number

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS UNBOUNDED PRECEDING) as SUMOVER,  
       SUM(1) OVER (ORDER BY Product_ID, Sale_Date  
                   ROWS UNBOUNDED PRECEDING) AS Seq_Number  
FROM Sales_Table ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>SUM OVER</u>	<u>Seq_Number</u>
1000	2000-09-28	48850.40	48850.40	1
1000	2000-09-29	54500.22	103350.62	2
1000	2000-09-30	36000.07	139350.69	3
1000	2000-10-01	40200.43	179551.12	4
1000	2000-10-02	32800.50	212351.62	5
1000	2000-10-03	64300.00	276651.62	6
1000	2000-10-04	54553.10	331204.72	7
2000	2000-09-28	41888.88	373093.60	8
2000	2000-09-29	48000.00	421093.60	9
2000	2000-09-30	49850.03	470943.63	10

Not all rows
are displayed in
this answer set

With “Seq_Number”, because you placed the number 1 in the area which calculates the cumulative sum, it'll continuously add 1 to the answer for each row.

Troubleshooting The ANSI OLAP on a GROUP BY

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (ORDER BY Sale_Date  
                             ROWS UNBOUNDED PRECEDING) AS SUMOVER  
FROM Sales_Table  
GROUP BY Product_ID ;
```

Error! Why?

Never **GROUP BY** in a **SUM()Over** or with any ANSI Syntax OLAP command. If you want to reset you use a **PARTITION BY** Statement, but never a **GROUP BY**.

The ANSI OLAP – Reset with a PARTITION BY Statement

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (PARTITION BY Product_ID  
                              ORDER BY Product_ID, Sale_Date  
                              ROWS UNBOUNDED PRECEDING) AS SumANSI  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>SumANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	103350.62
1000	2000-09-30	36000.07	139350.69
1000	2000-10-01	40200.43	179551.12
1000	2000-10-02	32800.50	212351.62
1000	2000-10-03	64300.00	276651.62
1000	2000-10-04	54553.10	331204.72
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	89888.88
2000	2000-09-30	49850.03	139738.91

The **PARTITION** Statement is how you reset in **ANSI**. This will cause the **SUMANSI** to start over (reset) on its calculating for each **NEW Product_ID**.

PARTITION BY only Resets a Single OLAP not ALL of them

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (PARTITION BY Product_ID  
ORDER BY Product_ID, Sale_Date  
ROWS UNBOUNDED PRECEDING) AS Subtotal,  
       SUM(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
ROWS UNBOUNDED PRECEDING) AS GRANDTotal  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>		<u>SubTotal</u>	<u>GRANDTotal</u>
1000	2000-09-28	48850.40		48850.40	48850.40
1000	2000-09-29	54500.22		103350.62	103350.62
1000	2000-09-30	36000.07		139350.69	139350.69
1000	2000-10-01	40200.43		179551.12	179551.12
1000	2000-10-02	32800.50		212351.62	212351.62
1000	2000-10-03	64300.00		276651.62	276651.62
1000	2000-10-04	54553.10		331204.72	331204.72
2000	2000-09-28	41888.88	↔	41888.88	373093.60
2000	2000-09-29	48000.00		89888.88	421093.60
2000	2000-09-30	49850.03		139738.91	470943.63

Above are two OLAP statements. Only one has PARTITION BY so only it resets.

The Moving Average (MAVG) and Moving Window

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table
```

↑
Moving
Average

↑
Moving
Window

↑ ↑
Major and Minor
Sort keys

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3 Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11

This is the **Moving Average (MAVG)**. It will calculate the average of 3 rows because that is the Moving Window. It will read the current row and TWO preceding to find the MAVG of those 3 rows. It will be sorted by Product_ID and Sale_Date first.

How the Moving Average is Calculated

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table
```

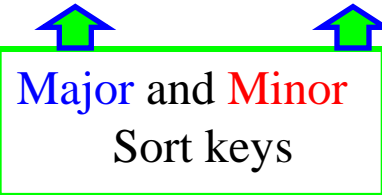
Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3 Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

With a Moving Window of 3, how is the 46450.23 amount derived in the **AVG3_Rows** column in the third row? It is the AVG of 48850.40, 54500.22 and 36000.07! The fourth row has AVG3_Rows equal to 43566.91. That was the average of 54500.22, 36000.07 and 40200.43. The calculation is on the current row and the two before.

How the Sort works for Moving Average (MAVG)

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table
```


Major and Minor
Sort keys


Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3 Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

The sorting is show above.

GROUP BY in the Moving Average does a Reset

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table  
GROUP BY Product_ID;
```




Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3 Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	44944.44
2000	2000-09-30	49850.03	46579.64

What does the GROUP BY Product_ID do? It causes a reset on all Product_ID breaks.

Quiz – Can you make the Advanced Calculation in your mind?

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table  
GROUP BY Product_ID;
```



Not all rows
are displayed in
this answer set


<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>AVG3_Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	44944.44
2000	2000-09-30	49850.03	46579.64
2000	2000-10-01	54850.29	50900.11



How is the 44944.44 derived in the 9th row of the AVG_for_3_Rows?

Answer to Quiz for the Advanced Calculation in your mind?

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table  
GROUP BY Product_ID;
```



Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3 Rows</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	44944.44
2000	2000-09-30	49850.03	46579.64

AVG of **41888.88** and **48000.00**

Notice there are only two calculations although this has a moving window of 3. That is because the GROUP BY caused the MAVG to reset when Product_ID 2000 came.

Quiz – Write that Teradata Moving Average in ANSI Syntax

```
SELECT      Product_ID , Sale_Date, Daily_Sales,  
            MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG3_Rows  
FROM Sales_Table ;
```

Challenge

Can you place
another **equivalent**
Moving Average
in the SQL above
using **ANSI Syntax**?

Here is a challenge that almost everyone fails. Can you do it perfectly?

Both the Teradata Moving Average and ANSI Version

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG_3,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID,  
                               Sale_Date ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG 3</u>	<u>AVG 3 ANSI</u>
1000	2000-09-28	48850.40	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31	51675.31
1000	2000-09-30	36000.07	46450.23	46450.23
1000	2000-10-01	40200.43	43566.91	43566.91
1000	2000-10-02	32800.50	36333.67	36333.67
1000	2000-10-03	64300.00	45788.98	45788.98
1000	2000-10-04	54553.10	50551.20	50551.20
2000	2000-09-28	41888.88	53580.66	53580.66
2000	2000-09-29	48000.00	48147.33	48147.33
2000	2000-09-30	49850.03	46579.11	46579.11

Not all rows
are displayed in
this answer set

The MAVG and AVG(Over) commands above are equivalent. Notice the Moving Window of 3 in the Teradata syntax and that it is a 2 in the ANSI version. That is because in ANSI it is considered the Current Row and 2 preceding.

The ANSI Moving Window is Current Row and Preceding

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding)AS AVG_3_ANSI  
FROM Sales_Table ;
```

Moving
Window
of 3 rows



Calculate the
Current Row
and 2 rows
preceding



Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>AVG_3_ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11

The AVG () Over allows you to do is to get the moving average of a certain column.

How ANSI Moving Average Handles the Sort

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                               ROWS 2 Preceding)AS AVG_3_ANSI  
FROM   Sales_Table ;
```

 Major and  Minor
Sort keys

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG 3 ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

Much like the SUM OVER Command, the Average OVER places the sort after the ORDER BY.

Quiz – How is that Total Calculated?

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>AVG_3_ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

With a Moving Window of 3, how is the 46450.23 amount derived in the AVG_3_ANSI column in the third row?

Answer to Quiz – How is that Total Calculated?

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG 3 ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

AVG of **48850.40**, **54500.22**, and **36000.07**

Quiz – How is that 4th Row Calculated?

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG 3 ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

With a Moving Window of 3, how is the 43566.91 amount derived in the AVG_3_ANSI column in the fourth row?

Answer to Quiz – How is that 4th Row Calculated?

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table;
```

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>AVG_3_ANSI</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31
1000	2000-09-30	36000.07	46450.23
1000	2000-10-01	40200.43	43566.91
1000	2000-10-02	32800.50	36333.67
1000	2000-10-03	64300.00	45788.98
1000	2000-10-04	54553.10	50551.20
2000	2000-09-28	41888.88	53580.66
2000	2000-09-29	48000.00	48147.33
2000	2000-09-30	49850.03	46579.11
2000	2000-10-01	54850.29	50900.11
2000	2000-10-02	36021.93	46907.42

AVG of **54500.22**, **36000.07**, and **40200.43**

Moving Average every 3-rows Vs a Continuous Average

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
    ROWS 2 Preceding) AS AVG3,  
AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
    ROWS UNBOUNDED Preceding) AS Continuous  
FROM Sales_Table;
```


Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>AVG3</u>	<u>Continuous</u>
1000	2000-09-28	48850.40	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31	51675.31
1000	2000-09-30	36000.07	46450.23	46450.23
1000	2000-10-01	40200.43	43566.91	44887.78
1000	2000-10-02	32800.50	36333.67	42470.32
1000	2000-10-03	64300.00	45788.98	46108.60
1000	2000-10-04	54553.10	50551.20	47314.96
2000	2000-09-28	41888.88	53580.66	46636.70
2000	2000-09-29	48000.00	48147.33	46788.18

The ROWS 2 Preceding gives the MAVG for every 3 rows. The ROWS UNBOUNDED Preceding gives the continuous MAVG.

Partition By Resets an ANSI OLAP

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS 2 Preceding) AS AVG3,  
       AVG(Daily_Sales) OVER (PARTITION BY Product_ID  
                             ORDER BY Product_ID, Sale_Date  
                             ROWS UNBOUNDED Preceding) AS Continuous  
FROM Sales_Table;
```



ANSI RESET
much Like a
GROUP BY


<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>AVG3</u>	<u>Continuous</u>
1000	2000-09-28	48850.40	48850.40	48850.40
1000	2000-09-29	54500.22	51675.31	51675.31
1000	2000-09-30	36000.07	46450.23	46450.23
1000	2000-10-01	40200.43	43566.91	44887.78
1000	2000-10-02	32800.50	36333.67	42470.32
1000	2000-10-03	64300.00	45788.98	46108.60
1000	2000-10-04	54553.10	50551.20	47314.96
2000	2000-09-28	41888.88	53580.66	41888.88
2000	2000-09-29	48000.00	48147.33	44944.44

Not all rows
are displayed in
this answer set

Use a **PARTITION BY** Statement to Reset the ANSI OLAP.

The Moving Difference (MDIFF)

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MDIFF(Daily_Sales, 4, Product_ID, Sale_Date) as “MDiff”  
FROM Sales_Table ;
```



Moving
Difference

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MDiff</u>
1000	2000-09-28	48850.40	?
1000	2000-09-29	54500.22	?
1000	2000-09-30	36000.07	?
1000	2000-10-01	40200.43	?
1000	2000-10-02	32800.50	-16049.90
1000	2000-10-03	64300.00	9799.78
1000	2000-10-04	54553.10	18553.03
2000	2000-09-28	41888.88	1688.45
2000	2000-09-29	48000.00	15199.50
2000	2000-09-30	49850.03	-14449.97
2000	2000-10-01	54850.29	297.19
2000	2000-10-02	36021.93	-5866.95
2000	2000-10-03	43200.18	-4799.82
2000	2000-10-04	32800.50	-17049.53

This is the Moving Difference (MDIFF). What this does is calculate the difference between the current row and only the 4th row preceding.

Moving Difference (MDIFF) Visual

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MDIFF(Daily_Sales, 4, Product_ID, Sale_Date) as “MDiff”  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MDiff</u>
1000	2000-09-28	48850.40	?
1000	2000-09-29	54500.22	?
1000	2000-09-30	36000.07	?
1000	2000-10-01	40200.43	?
1000	2000-10-02	32800.50	-16049.90
1000	2000-10-03	64300.00	9799.78
1000	2000-10-04	54553.10	18553.03
2000	2000-09-28	41888.88	1688.45
2000	2000-09-29	48000.00	15199.50
2000	2000-09-30	49850.03	-14449.97
2000	2000-10-01	54850.29	297.19
2000	2000-10-02	36021.93	-5866.95
2000	2000-10-03	43200.18	-4799.82
2000	2000-10-04	32800.50	-17049.53

How much more did we make for Product_ID 1000 on 2000-10-03 versus Product_ID 1000 which was 4 rows earlier on 2000-09-29?

Moving Difference using ANSI Syntax

```
SELECT Product_ID, Sale_Date, Daily_Sales,  
       Daily_Sales - SUM(Daily_Sales)  
       OVER ( ORDER BY   Product_ID ASC, Sale_Date ASC  
             ROWS BETWEEN 4 PRECEDING AND 4 PRECEDING)  
           AS "MDiff_ANSI"  
FROM Sales_Table ;
```

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>MDiff_ANSI</u>
1000	2000-09-28	48850.40	?
1000	2000-09-29	54500.22	?
1000	2000-09-30	36000.07	?
1000	2000-10-01	40200.43	?
1000	2000-10-02	32800.50	-16049.90
1000	2000-10-03	64300.00	9799.78
1000	2000-10-04	54553.10	18553.03
2000	2000-09-28	41888.88	1688.45
2000	2000-09-29	48000.00	15199.50
2000	2000-09-30	49850.03	-14449.97
2000	2000-10-01	54850.29	297.19
2000	2000-10-02	36021.93	-5866.95

This is how you do a **MDiff** using the **ANSI** Syntax with a **moving window of 4**.

Moving Difference using ANSI Syntax with Partition By

SELECT Product_ID, Sale_Date (Format 'yyyy-mm-dd'), Daily_Sales,
Daily_Sales - SUM(Daily_Sales) OVER (**PARTITION BY** Product_ID
ORDER BY Product_ID ASC, Sale_Date ASC
ROWS BETWEEN 4 PRECEDING AND 4 PRECEDING)
AS "MDiff_ANSI"

FROM Sales_Table;

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>MDiff_ANSI</u>
1000	2000-09-28	48850.40	?
1000	2000-09-29	54500.22	?
1000	2000-09-30	36000.07	?
1000	2000-10-01	40200.43	?
1000	2000-10-02	32800.50	-16049.90
1000	2000-10-03	64300.00	9799.78
1000	2000-10-04	54553.10	18553.03
2000	2000-09-28	41888.88	?
2000	2000-09-29	48000.00	?
2000	2000-09-30	49850.03	?
2000	2000-10-01	54850.29	?
2000	2000-10-02	36021.93	-5866.95
2000	2000-10-03	43200.18	-4799.82
2000	2000-10-04	32800.50	-17049.53

Wow! This is how you do a MDiff using the ANSI Syntax with a PARTITION BY.

Trouble Shooting the Moving Difference (MDIFF)

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MDIFF(Daily_Sales, 7, Product_ID, Sale_Date) as Compare2Rows  
FROM Sales_Table  
GROUP BY Product_ID ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Compare2Rows</u>
1000	2000-09-28	48850.40	?
1000	2000-09-29	54500.22	?
1000	2000-09-30	36000.07	?
1000	2000-10-01	40200.43	?
1000	2000-10-02	32800.50	?
1000	2000-10-03	64300.00	?
1000	2000-10-04	54553.10	?
2000	2000-09-28	41888.88	?
2000	2000-09-29	48000.00	?
2000	2000-09-30	49850.03	?
2000	2000-10-01	54850.29	?
2000	2000-10-02	36021.93	?
2000	2000-10-03	43200.18	?
2000	2000-10-04	32800.50	?

Do you notice that column Compare2Rows did not produce any data? That is because the **GROUP BY** Reset before it could get 7 records to find the MDIFF.

The RANK Command

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales) AS “Rank”  
FROM Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank</u>
1000	2000-10-03	64300.00	1
2000	2000-10-01	54850.29	2
1000	2000-10-04	54553.10	3
1000	2000-09-29	54500.22	4
2000	2000-09-30	49850.03	5
1000	2000-09-28	48850.40	6
2000	2000-09-29	48000.00	7
2000	2000-10-03	43200.18	8
2000	2000-09-28	41888.88	9
1000	2000-10-01	40200.43	10
2000	2000-10-02	36021.93	11
1000	2000-09-30	36000.07	12
2000	2000-10-04	32800.50	13
1000	2000-10-02	32800.50	13

This is the RANK. In this example, it will rank your Daily_Sales from greatest to least. The default for this type of RANK is to sort DESC.

How to get Rank to Sort in Ascending Order

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales ASC) AS “Rank”  
FROM  Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank</u>
1000	2000-10-02	32800.50	1
2000	2000-10-04	32800.50	1
1000	2000-09-30	36000.07	3
2000	2000-10-02	36021.93	4
1000	2000-10-01	40200.43	5
2000	2000-09-28	41888.88	6
2000	2000-10-03	43200.18	7
2000	2000-09-29	48000.00	8
1000	2000-09-28	48859.40	9
2000	2000-09-30	49850.03	10
1000	2000-09-29	54500.22	11
1000	2000-10-04	54553.10	12
2000	2000-10-01	54850.29	13
1000	2000-10-03	64300.00	14

This RANK query sorts in **Ascending** mode.

Two ways to get Rank to Sort in Ascending Order

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales ASC) AS Rank1,  
            RANK(-Daily_Sales)      AS Rank2  
FROM  Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>	<u>Rank2</u>
1000	2000-10-02	32800.50	1	1
2000	2000-10-04	32800.50	1	1
1000	2000-09-30	36000.07	3	3
2000	2000-10-02	36021.93	4	4
1000	2000-10-01	40200.43	5	5
2000	2000-09-28	41888.88	6	6
2000	2000-10-03	43200.18	7	7
2000	2000-09-29	48000.00	8	8
1000	2000-09-28	48859.40	9	9
2000	2000-09-30	49850.03	10	10
1000	2000-09-29	54500.22	11	11
1000	2000-10-04	54553.10	12	12
2000	2000-10-01	54850.29	13	13
1000	2000-10-03	64300.00	14	14

A minus sign or keyword ASC will sort Both RANK in **Ascending mode**.

RANK using ANSI Syntax Defaults to Ascending Order

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK() OVER (ORDER BY Daily_Sales) AS Rank1  
FROM    Sales_Table  
WHERE Product_ID IN (1000, 2000)
```

Not all rows
are displayed in
this answer set

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>Rank1</u>
2000	2000-10-04	32800.50	1
1000	2000-10-04	32800.50	1
1000	2000-09-30	65000.07	3
2000	2000-10-02	36021.93	4
1000	2000-10-01	40200.43	5
2000	2000-09-28	41888.88	6
2000	2000-10-03	43200.18	7
2000	2000-09-29	48000.00	8

This is the RANK() OVER. It provides a rank for your queries. Notice how you do not place anything within the () after the word RANK. Default Sort is ASC.

Getting RANK using ANSI Syntax to Sort in DESC Order

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK() OVER (ORDER BY Daily_Sales DESC) AS Rank1  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000) ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
2000	2000-09-29	48000.00	1
2000	2000-10-03	43200.00	2
2000	2000-09-28	41888.88	3
1000	2000-10-01	40200.43	4
2000	2000-10-02	36032.93	5
1000	2000-09-30	65000.07	6
1000	2000-10-04	32800.50	8
2000	2000-10-04	32800.50	8

Is the query above in ASC mode or DESC mode for sorting?

RANK() OVER and PARTITION BY with a QUALIFY

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK() OVER (PARTITION BY Product_ID  
                    ORDER BY Daily_Sales DESC) AS Rank1  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000)  
QUALIFY Rank1 < 4
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-03	65000.07	1
1000	2000-10-04	54553.10	2
1000	2000-09-29	54500.22	3
2000	2000-10-01	54850.29	1
2000	2000-09-30	49850.03	2
2000	2000-09-29	48000.00	3

What does the PARTITION Statement in the RANK() OVER do? It resets the rank.
The QUALIFY statement limits rows once the Rank's been calculated.

QUALIFY and WHERE

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales ASC) AS Rank1,  
            RANK(-Daily_Sales)      AS Rank2  
FROM  Sales_Table  
WHERE Product_ID IN (1000, 2000)  
QUALIFY Rank(-Daily_Sales) < 6 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>	<u>Rank2</u>
1000	2000-10-02	32800.50	1	1
2000	2000-10-04	32800.50	1	1
1000	2000-09-30	36000.07	3	3
2000	2000-10-02	36021.93	4	4
1000	2000-10-01	40200.43	5	5

The **WHERE** statement is performed first. It limits the rows being calculated. Then the **QUALIFY** takes the calculated rows and limits the returning rows. **QUALIFY** is to OLAP what **HAVING** is to Aggregates. Both limit after the calculations.

Quiz – How can you simplify the QUALIFY Statement

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK(Daily_Sales ASC) AS Rank1,  
       RANK(-Daily_Sales)    AS Rank2  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000)  
QUALIFY Rank(-Daily_Sales) < 6 ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>Rank1</u>	<u>Rank2</u>
1000	2000-10-02	32800.50	1	1
2000	2000-10-04	32800.50	1	1
1000	2000-09-30	36000.07	3	3
2000	2000-10-02	36021.93	4	4
1000	2000-10-01	40200.43	5	5

How can you improve the QUALIFY Statement above for simplicity?

Answer to Quiz –Can you simplify the QUALIFY Statement

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK(Daily_Sales ASC) AS Rank1,  
       RANK(-Daily_Sales)      AS Rank2  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000)  
QUALIFY Rank2 < 6 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>	<u>Rank2</u>
1000	2000-10-02	32800.50	1	1
2000	2000-10-04	32800.50	1	1
1000	2000-09-30	36000.07	3	3
2000	2000-10-02	36021.93	4	4
1000	2000-10-01	40200.43	5	5

QUALIFY Rank2 < 6
(Use the Alias)

The QUALIFY Statement without Ties

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales) AS Rank1  
FROM Sales_Table  
WHERE Product_ID IN (1000, 2000)  
QUALIFY Rank1 < 6 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-03	64300.00	1
2000	2000-10-01	54850.29	2
1000	2000-10-04	54553.10	3
1000	2000-09-29	54500.22	4
2000	2000-09-30	49850.03	5

A QUALIFY < 6 will provide a result that is 5 rows. Notice there are NO ties, yet!

The QUALIFY Statement with Ties

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales ASC) AS Rank1  
FROM Sales_Table  
WHERE Product_ID IN (1000, 2000)  
QUALIFY Rank1 < 6 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-02	32800.50	1
2000	2000-10-04	32800.50	1
1000	2000-09-30	36000.07	3
2000	2000-10-02	36021.93	4
1000	2000-10-01	40200.43	5

A **QUALIFY** < 6 will provide a result that is 5 rows. Notice there are **Ties**!

The QUALIFY Statement with Ties Brings back Extra Rows

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales ASC) AS Rank1  
FROM Sales_Table  
WHERE Product_ID IN (1000, 2000)  
QUALIFY Rank1 < 2 ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>Rank1</u>
1000	2000-10-02	32800.50	1
2000	2000-10-04	32800.50	1

A QUALIFY < 2 will provide more rows than 1 because of the **Ties!**

Mixing Sort Order for QUALIFY Statement

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales) AS Rank1  
FROM  Sales_Table  
WHERE Product_ID IN (1000, 2000)  
QUALIFY RANK (Daily_Sales ASC) < 6 ;
```

← **DESC Mode**

← **ASC Mode**

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-02	32800.50	13
2000	2000-10-04	32800.50	13
1000	2000-09-30	36000.07	12
2000	2000-10-02	36021.93	11
1000	2000-10-01	40200.43	10

Look at the Rankings and the Daily_Sales. This data come out odd because Rank1 is **DESC** by default (using this Syntax) and the QUALIFY specifies ASC mode.

Quiz – What Caused the RANK to Reset?

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales) AS Rank1  
FROM Sales_Table  
WHERE Product_ID IN (1000, 2000)  
GROUP BY Product_ID  
QUALIFY Rank1 < 4 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-03	64300.00	1
1000	2000-10-04	54553.10	2
1000	2000-09-29	54500.22	3
2000	2000-10-01	54850.29	1
2000	2000-09-30	49850.03	2
2000	2000-09-29	48000.00	3

What caused the data to reset the column Rank1?

Answer to Quiz – What Caused the RANK to Reset?

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
            RANK(Daily_Sales) AS Rank1  
FROM Sales_Table  
WHERE Product_ID IN (1000, 2000)  
GROUP BY Product_ID  
QUALIFY Rank1 < 4 ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Rank1</u>
1000	2000-10-03	64300.00	1
1000	2000-10-04	54553.10	2
1000	2000-09-29	54500.22	3
2000	2000-10-01	54850.29	1
2000	2000-09-30	49850.03	2
2000	2000-09-29	48000.00	3

GROUP BY

Quiz – Name those Sort Orders

RANK() OVER (ORDER BY Daily_Sales) AS ANSI_Rank

Is the default above ASC or DESC?

RANK(Daily_Sales) AS NON_ANSI_Rank

Is the default above ASC or DESC?

Answer the questions above.

Answer to Quiz – Name those Sort Orders

RANK() OVER (ORDER BY Daily_Sales) AS ANSI_Rank

Defaults to ASC

RANK(Daily_Sales) AS NON_ANSI_Rank

Defaults to DESC

Please note that by default these different syntaxes sort completely opposite.

PERCENT_RANK() OVER

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       PERCENT_RANK() OVER (PARTITION BY PRODUCT_ID  
       ORDER BY Daily_Sales DESC) AS PercentRank1  
FROM   Sales_Table WHERE Product_ID in (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>PercentRank1</u>
1000	2000-10-03	64300.00	0.000000
1000	2000-10-04	54553.10	0.166667
1000	2000-09-29	54500.22	0.333333
1000	2000-09-28	48850.40	0.500000
1000	2000-10-01	40200.43	0.666667
1000	2000-09-30	36000.07	0.833333
1000	2000-10-02	32800.50	1.000000
2000	2000-10-01	54850.29	0.000000
2000	2000-09-30	49850.03	0.166667
2000	2000-09-29	48000.00	0.333333
2000	2000-10-03	43200.18	0.500000
2000	2000-09-28	41888.88	0.666667
2000	2000-10-02	36021.93	0.833333
2000	2000-10-04	32800.50	1.000000

7 Rows in
Calculation
for **1000**
Product_ID

7 Rows in
Calculation
for **2000**
Product_ID

We now have added a Partition statement which produces 7 rows per Product_ID.

PERCENT_RANK() OVER with 14 rows in Calculation

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       PERCENT_RANK() OVER ( ORDER BY Daily_Sales DESC)  
                                AS PercentRank1  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>PercentRank1</u>
1000	2000-10-03	64300.00	0.000000
2000	2000-10-01	54850.29	0.076923
1000	2000-10-04	54553.10	0.153846
1000	2000-09-29	54500.22	0.230769
2000	2000-09-30	49850.03	0.307692
1000	2000-09-28	48850.40	0.384615
2000	2000-09-29	48000.00	0.461538
2000	2000-10-03	43200.18	0.538462
2000	2000-09-28	41888.88	0.615385
1000	2000-10-01	40200.43	0.692308
2000	2000-10-02	36021.93	0.769231
1000	2000-09-30	36000.07	0.846154
2000	2000-10-04	32800.50	0.923077
1000	2000-10-02	32800.50	0.923077

14 Rows in
Calculation
for both the
1000 and 2000
Product_IDs

Percentage_Rank is just like RANK however, it gives you the Rank as a percent, but only a percent of all the other rows up to 100%.

PERCENT_RANK() OVER with 21 rows in Calculation

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       PERCENT_RANK() OVER ( ORDER BY Daily_Sales DESC)  
                                AS PercentRank1  
FROM   Sales_Table ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>PercentRank1</u>
1000	2000-10-03	64300.00	0.000000
3000	2000-09-28	61301.77	0.050000
2000	2000-10-01	54850.29	0.100000
1000	2000-10-04	54553.10	0.150000
1000	2000-09-29	54500.22	0.200000
2000	2000-09-30	49850.03	0.250000
1000	2000-09-28	48850.40	0.300000
2000	2000-09-29	48000.00	0.350000
3000	2000-09-30	43868.86	0.400000
2000	2000-10-03	43200.18	0.450000
2000	2000-09-28	41888.88	0.500000
1000	2000-10-01	40200.43	0.550000
2000	2000-10-02	36021.93	0.600000
1000	2000-09-30	36000.07	0.650000

Not all rows
are displayed in
this answer set

21 Rows in
Calculation
for all of the
Product_IDs

Percentage_Rank is just like RANK however, it gives you the Rank as a percent, but only a percent of all the other rows up to 100%.

Quiz – What Cause the Product_ID to Reset

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       PERCENT_RANK() OVER (PARTITION BY PRODUCT_ID  
       ORDER BY Daily_Sales DESC) AS PercentRank1  
FROM   Sales_Table WHERE Product_ID in (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>PercentRank1</u>
1000	2000-10-03	64300.00	0.000000
1000	2000-10-04	54553.10	0.166667
1000	2000-09-29	54500.22	0.333333
1000	2000-09-28	48850.40	0.500000
1000	2000-10-01	40200.43	0.666667
1000	2000-09-30	36000.07	0.833333
1000	2000-10-02	32800.50	1.000000
2000	2000-10-01	54850.29	0.000000
2000	2000-09-30	49850.03	0.166667
2000	2000-09-29	48000.00	0.333333
2000	2000-10-03	43200.18	0.500000
2000	2000-09-28	41888.88	0.666667
2000	2000-10-02	36021.93	0.833333
2000	2000-10-04	32800.50	1.000000

What caused the Product_IDs to be sorted?

Answer to Quiz – What Cause the Product_ID to Reset

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       PERCENT_RANK() OVER (PARTITION BY PRODUCT_ID  
       ORDER BY Daily_Sales DESC) AS PercentRank1  
FROM   Sales_Table WHERE Product_ID in (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>PercentRank1</u>
1000	2000-10-03	64300.00	0.000000
1000	2000-10-04	54553.10	0.166667
1000	2000-09-29	54500.22	0.333333
1000	2000-09-28	48850.40	0.500000
1000	2000-10-01	40200.43	0.666667
1000	2000-09-30	36000.07	0.833333
1000	2000-10-02	32800.50	1.000000
2000	2000-10-01	54850.29	0.000000
2000	2000-09-30	49850.03	0.166667
2000	2000-09-29	48000.00	0.333333
2000	2000-10-03	43200.18	0.500000
2000	2000-09-28	41888.88	0.666667
2000	2000-10-02	36021.93	0.833333
2000	2000-10-04	32800.50	1.000000

PARTITION BY caused the data to be sorted!

COUNT OVER for a Sequential Number

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       COUNT(*) OVER (ORDER BY Product_ID, Sale_Date  
       ROWS UNBOUNDED PRECEDING) AS Seq_Number  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Seq Number</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	8
2000	2000-09-29	48000.00	9
2000	2000-09-30	49850.03	10
2000	2000-10-01	54850.29	11

This is the **COUNT OVER**. It will provide a sequential number starting at 1. The Keyword(s) **ROWS UNBOUNDED PRECEDING** causes Seq_Number to start at the beginning and increase sequentially to the end.

Troubleshooting COUNT OVER

```
SELECT      Product_ID ,Sale_Date , Daily_Sales,  
COUNT(*) OVER (ORDER BY Product_ID, Sale_Date) AS No_Seq  
FROM    Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

Rows Unbounded Preceding is missing in this statement.

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>Seq Number</u>
1000	2000-09-28	48850.40	14
1000	2000-09-29	54500.22	14
1000	2000-09-30	36000.07	14
1000	2000-10-01	40200.43	14
1000	2000-10-02	32800.50	14
1000	2000-10-03	64300.00	14
1000	2000-10-04	54553.10	14
2000	2000-09-28	41888.88	14
2000	2000-09-29	48000.00	14
2000	2000-09-30	49850.03	14
2000	2000-10-01	54850.29	14
2000	2000-10-02	36021.93	14
2000	2000-10-03	43200.18	14
2000	2000-10-04	32800.50	14

14 rows
came back

When you don't have a **ROWS UNBOUNDED PRECEDING**, No_Seq get a value of 14 on every row. Why? Because 14 is the **FINAL COUNT NUMBER**.

Quiz – What caused the COUNT OVER to Reset?

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       COUNT(*) OVER (PARTITION BY Product_ID  
                      ORDER BY Product_ID, Sale_Date  
                      ROWS UNBOUNDED PRECEDING) AS StartOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>StartOver</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	1
2000	2000-09-29	48000.00	2
2000	2000-09-30	49850.03	3
2000	2000-10-01	54850.29	4
2000	2000-10-02	36021.93	5
2000	2000-10-03	43200.18	6
2000	2000-10-04	32800.50	7

What Keyword(s) caused StartOver to reset?

Answer to Quiz – What caused the COUNT OVER to Reset?

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       COUNT(*) OVER (PARTITION BY Product_ID  
                      ORDER BY Product_ID, Sale_Date  
                      ROWS UNBOUNDED PRECEDING) AS StartOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>StartOver</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	1
2000	2000-09-29	48000.00	2
2000	2000-09-30	49850.03	3
2000	2000-10-01	54850.29	4
2000	2000-10-02	36021.93	5
2000	2000-10-03	43200.18	6
2000	2000-10-04	32800.50	7

PARTITION BY

The MAX OVER Command

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       MAX(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                               ROWS UNBOUNDED PRECEDING) AS MaxOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>MaxOver</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	54500.22
1000	2000-09-30	36000.07	54500.22
1000	2000-10-01	40200.43	54500.22
1000	2000-10-02	32800.50	54500.22
1000	2000-10-03	64300.00	64300.00
1000	2000-10-04	54553.10	64300.00
2000	2000-09-28	41888.88	64300.00
2000	2000-09-29	48000.00	64300.00
2000	2000-09-30	49850.03	64300.00
2000	2000-10-01	54850.29	64300.00
2000	2000-10-02	36021.93	64300.00
2000	2000-10-03	43200.18	64300.00
2000	2000-10-04	32800.50	64300.00

After the sort the **Max() Over** shows the **Max Value** up to that point.

MAX OVER with PARTITION BY Reset

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       MAX(Daily_Sales) OVER (PARTITION BY Product_ID  
                              ORDER BY Product_ID, Sale_Date  
                              ROWS UNBOUNDED PRECEDING) AS MaxOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MaxOver</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	54500.22
1000	2000-09-30	36000.07	54500.22
1000	2000-10-01	40200.43	54500.22
1000	2000-10-02	32800.50	54500.22
1000	2000-10-03	64300.00	64300.00
1000	2000-10-04	54553.10	64300.00
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	48000.00
2000	2000-09-30	49850.03	49850.03
2000	2000-10-01	54850.29	54850.29

Not all rows
are displayed in
this answer set

The largest value is 64300.00 in the column **MaxOver**. Once it was evaluated it did not continue until the end because of the PARTITION BY reset.

Troubleshooting MAX OVER

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       MAX(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date )  
                                AS MaxOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

Rows Unbounded Preceding is missing in this statement.

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>MaxOver</u>
1000	2000-09-28	48850.40	64300.00
1000	2000-09-29	54500.22	64300.00
1000	2000-09-30	36000.07	64300.00
1000	2000-10-01	40200.43	64300.00
1000	2000-10-02	32800.50	64300.00
1000	2000-10-03	64300.00	64300.00
1000	2000-10-04	54553.10	64300.00
2000	2000-09-28	41888.88	64300.00
2000	2000-09-29	48000.00	64300.00
2000	2000-09-30	49850.03	64300.00

Not all rows
are displayed in
this answer set

You can also use MAX as a OLAP. 64300.00 came back in MaxOver because that was the MAX value for Daily_Sales in this Answer Set. Notice that it **doesn't** have a **ROWS UNBOUNDED PRECEDING**.

The MIN OVER Command

```
SELECT Product_ID, Sale_Date ,Daily_Sales  
      ,Min(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                             ROWS UNBOUNDED PRECEDING) AS MinOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MinOver</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	48850.40
1000	2000-09-30	36000.07	36000.07
1000	2000-10-01	40200.43	36000.07
1000	2000-10-02	32800.50	32800.50
1000	2000-10-03	64300.00	32800.50
1000	2000-10-04	54553.10	32800.50
2000	2000-09-28	41888.88	32800.50
2000	2000-09-29	48000.00	32800.50
2000	2000-09-30	49850.03	32800.50
2000	2000-10-01	54850.29	32800.50
2000	2000-10-02	36021.93	32800.50
2000	2000-10-03	43200.18	32800.50
2000	2000-10-04	32800.50	32800.50

After the sort the MIN () Over shows the Max Value up to that point.

Troubleshooting MIN OVER

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       MIN(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date )  
                                AS MinOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

Rows Unbounded Preceding is missing in this statement.

Not all rows
are displayed in
this answer set

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MinOver</u>
1000	2000-09-28	48850.40	32800.50
1000	2000-09-29	54500.22	32800.50
1000	2000-09-30	36000.07	32800.50
1000	2000-10-01	40200.43	32800.50
1000	2000-10-02	32800.50	32800.50
1000	2000-10-03	64300.00	32800.50
1000	2000-10-04	54553.10	32800.50
2000	2000-09-28	41888.88	32800.50
2000	2000-09-29	48000.00	32800.50
2000	2000-09-30	49850.03	32800.50
2000	2000-10-01	54850.29	32800.50

Min only displayed 32800.50 because there is **NOT** a **ROWS UNBOUNDED PRECEDING** statement so it found the **lowest Daily_Sales** and repeated it.

Quiz – Fill in the Blank

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       Min(Daily_Sales) OVER (PARTITION BY Product_ID  
                               ORDER BY Product_ID, Sale_Date  
                               ROWS UNBOUNDED PRECEDING) AS MinOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MinOver</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	48850.40
1000	2000-09-30	36000.07	36000.07
1000	2000-10-01	40200.43	36000.07
1000	2000-10-02	32800.50	32800.50
1000	2000-10-03	64300.00	32800.50
1000	2000-10-04	54553.10	32800.50
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	41888.88
2000	2000-09-30	49850.03	41888.88
2000	2000-10-01	54850.29	41888.88
2000	2000-10-02	36021.93	36021.93
2000	2000-10-03	43200.18	
2000	2000-10-04	32800.50	

The last two answers (MinOver) are blank so you can fill in the blank.

Answer to Quiz – Fill in the Blank

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       Min(Daily_Sales) OVER (PARTITION BY Product_ID  
                              ORDER BY Product_ID, Sale_Date  
                              ROWS UNBOUNDED PRECEDING) AS MinOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>MinOver</u>
1000	2000-09-28	48850.40	48850.40
1000	2000-09-29	54500.22	48850.40
1000	2000-09-30	36000.07	36000.07
1000	2000-10-01	40200.43	36000.07
1000	2000-10-02	32800.50	32800.50
1000	2000-10-03	64300.00	32800.50
1000	2000-10-04	54553.10	32800.50
2000	2000-09-28	41888.88	41888.88
2000	2000-09-29	48000.00	41888.88
2000	2000-09-30	49850.03	41888.88
2000	2000-10-01	54850.29	41888.88
2000	2000-10-02	36021.93	36021.93
2000	2000-10-03	43200.18	36021.93
2000	2000-10-04	32800.50	32800.50

The Row_Number Command

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       ROW_NUMBER() OVER (ORDER BY Product_ID, Sale_Date)  
       AS Seq_Number  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>Seq_Number</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	8
2000	2000-09-29	48000.00	9
2000	2000-09-30	49850.03	10
2000	2000-10-01	54850.29	11

Not all rows
are displayed in
this answer set

The **ROW_NUMBER()** Keyword(s) caused Seq_Number to increase sequentially.
Notice that this does **NOT** have a **Rows Unbounded Preceding** and it still works!

Quiz – How did the Row_Number Reset?

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       ROW_NUMBER() OVER (PARTITION BY Product_ID  
                           ORDER BY Product_ID, Sale_Date ) AS StartOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product ID</u>	<u>Sale Date</u>	<u>Daily Sales</u>	<u>StartOver</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	1
2000	2000-09-29	48000.00	2
2000	2000-09-30	49850.03	3
2000	2000-10-01	54850.29	4
2000	2000-10-02	36021.93	5
2000	2000-10-03	43200.18	6
2000	2000-10-04	32800.50	7

What Keyword(s) caused StartOver to reset?

Quiz – How did the Row_Number Reset?

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       ROW_NUMBER() OVER (PARTITION BY Product_ID  
                          ORDER BY Product_ID, Sale_Date ) AS StartOver  
FROM   Sales_Table WHERE Product_ID IN (1000, 2000) ;
```

<u>Product_ID</u>	<u>Sale_Date</u>	<u>Daily_Sales</u>	<u>StartOver</u>
1000	2000-09-28	48850.40	1
1000	2000-09-29	54500.22	2
1000	2000-09-30	36000.07	3
1000	2000-10-01	40200.43	4
1000	2000-10-02	32800.50	5
1000	2000-10-03	64300.00	6
1000	2000-10-04	54553.10	7
2000	2000-09-28	41888.88	1
2000	2000-09-29	48000.00	2
2000	2000-09-30	49850.03	3
2000	2000-10-01	54850.29	4
2000	2000-10-02	36021.93	5
2000	2000-10-03	43200.18	6
2000	2000-10-04	32800.50	7

What Keyword(s) caused StartOver to **reset**?

PARTITION BY

Testing Your Knowledge

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Product_ID, Sale_Date)  AS “CSum”  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000  
GROUP BY Product_ID ;
```

This is the **CSUM**. However, what we want to see is the **Sum()Over ANSI** version. Use the information in the **CSUM** and **convert** this to the equivalent **Sum()Over**.

Testing Your Knowledge

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       CSUM(Daily_Sales, Product_ID, Sale_Date)  AS “CSum”  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000  
GROUP BY Product_ID ;
```

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       SUM(Daily_Sales) OVER (PARTITION BY Product_ID  
                              ORDER BY Product_ID, Sale_Date  
                              ROWS UNBOUNDED PRECEDING)  AS SumANSI  
FROM Sales_Table  
WHERE Product_ID BETWEEN 1000 and 2000 ;
```

Both statements are exactly the same except the bottom example uses ANSI syntax.

Testing Your Knowledge

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
       MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG_for_3_Rows  
FROM   Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

Write the equivalent to the SQL above using ANSI Syntax such as AVG () Over.

Testing Your Knowledge

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
MAVG( Daily_Sales, 3, Product_ID, Sale_Date) AS AVG_for_3_Rows  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

```
SELECT Product_ID , Sale_Date, Daily_Sales,  
    AVG(Daily_Sales) OVER (ORDER BY Product_ID, Sale_Date  
                           ROWS 2 Preceding) AS AVG_3_ANSI  
FROM Sales_Table WHERE Product_ID BETWEEN 1000 and 2000 ;
```

The SQL above is equivalent except the bottom example uses ANSI Syntax.

Testing Your Knowledge

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK(Daily_Sales) AS “Rank”  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000) ;
```

This is the **Rank**. However, what we want to see is the `RANK()Over`. Use the information in the Rank to make it the `Rank()Over`.

Testing Your Knowledge

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK(Daily_Sales) AS "Rank"  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000) ;
```

```
SELECT Product_ID ,Sale_Date , Daily_Sales,  
       RANK() OVER (ORDER BY Daily_Sales DESC) AS Rank1  
FROM   Sales_Table  
WHERE  Product_ID IN (1000, 2000)
```

The SQL above is equivalent except the bottom example uses ANSI Syntax. Also notice the sort key. DESC is the default in the top example.